A FULL-SEQUENCE QUANTITATIVE GAP BETWEEN THE CHROMATIC AND COCHROMATIC NUMBERS OF A RANDOM GRAPH

SAMUIL PETKOV

Abstract. A proper coloring partitions the vertices of a graph into independent sets. In a cocoloring, each class may be an independent set or a clique. The minimum numbers of classes are the chromatic number $\chi(G)$ and cochromatic number $\zeta(G)$, respectively. Thus $\zeta(G) \le \chi(G)$. For $G_n \sim G(n, 1/2)$, Erdős and Gimbel asked whether $\chi(G_n) - \zeta(G_n)$ tends to infinity with high probability. We prove

$$\mathbb{P}\left(\chi(G_n) - \zeta(G_n) \ge \frac{(\log 2)^2}{4} \log\left(\frac{200}{153}\right) \frac{n}{(\log n)^3}\right) \longrightarrow 1.$$

All logarithms are natural, and the limit holds as $n \to \infty$ through all integers. This lower bound has the conjectured scale $n/(\log n)^3$.

We choose four consecutive class sizes below the real-valued first-moment center for independent sets. The fractional part of this center is the rounding phase. Optimizing the proportions at each phase gives a sharper explicit coefficient, with a deterministic error tending to zero uniformly over the phase. For the second moment, we count partitions with each class declared independent or complete, and bound the contribution from pairs of overlapping partitions. This gives a cocoloring with a possibly small probability. Concentration and a coloring of the leftover vertices then give high probability at a cost of $o(n/(\log n)^3)$ extra classes.

Introduction

A *cocoloring* of a graph G is a partition of $V(G)$ into nonempty classes, each inducing either an edgeless graph or a complete graph. The least number of classes in such a partition is the *cochromatic number* $\zeta(G)$. Since an ordinary proper coloring is a cocoloring using only independent-set classes, one always has $\zeta(G) \le \chi(G)$.

Write $[n] := \{1, \dots, n\}$, and let $G_n \sim G(n, 1/2)$ be the labeled random graph on $[n]$. All logarithms are natural unless a base is displayed. The first-order scales of $\chi(G_n)$ and $\zeta(G_n)$ coincide, so this asymptotic alone does not tell us whether their difference grows. Erdős and Gimbel asked whether

$$\chi(G_n) - \zeta(G_n) \longrightarrow \infty$$

with high probability (Erdős and Gimbel, 1993, p. 263). The question was restated by Gimbel (Gimbel, 2016, Section 7.4) and is cataloged as Erdős Problem 625 (Bloom, 2026). We prove a quantitative bound along the full sequence of integers.

The expected number of independent sets of size k is

$$\mu_k = \binom{n}{k} 2^{-\binom{k}{2}}.$$

The usual asymptotic solution of $\mu_k = 1$ is the real-valued center $\alpha_0(n)$ defined below (Bollobás and Erdős, 1976, p. 420). Its fractional part records the position of this center relative to the neighboring integer class sizes and will be the phase δ_n. We include both quantities in the theorem statement.

Date: 30 August 2026.

2020 *Mathematics Subject Classification.* Primary 05C80; Secondary 05C15, 60C05.

Key words and phrases. random graph, chromatic number, cochromatic number, second moment method, configuration model.

Main theorem (Uniform lower bound and phase refinement)**.** *For every integer $n \geq 2$, let $G_n \sim G(n, 1/2)$ be the random simple graph on $[n]$, with edges present independently with probability $1/2$. All logarithms without a displayed base are natural, and all limits are as $n \to \infty$ through the integers. Then*

$$\mathbb{P}\left(\chi(G_n) - \zeta(G_n) \geq \frac{(\log 2)^2}{4} \log\left(\frac{200}{153}\right) \frac{n}{(\log n)^3}\right) \longrightarrow 1.$$

For the sharper phase-dependent form, set

$$\alpha_0(n) := 2\log_2 n - 2\log_2 \log_2 n + 2\log_2(\mathrm{e}/2) + 1, \qquad \delta_n := \{\alpha_0(n)\} = \alpha_0(n) - \lfloor \alpha_0(n) \rfloor.$$

Thus $\delta_n \in [0, 1)$ is deterministic. For $0 \leq \delta \leq 1$, define

$$S_+ := \{i \in \mathbb{Z} : i \geq -1\}, \qquad S_4 := \{2, 3, 4, 5\}, \qquad T_\delta := 1 + \frac{2}{\log 2} - \delta.$$

Then $T_\delta \in (2, 5)$. For $S \in \{S_+, S_4\}$ and $T \in (2, 5)$, put

$$\mathcal{F}_S(T) := \inf_{\lambda \in \mathbb{R}} \left\{ \log \sum_{i \in S} \exp\left(\lambda i - \frac{\log 2}{2} i^2\right) - \lambda T \right\}, \qquad A_4(\delta) := \log 2 - \mathcal{F}_{S_+}(T_\delta) + \mathcal{F}_{S_4}(T_\delta).$$

The sum inside the logarithm defining $\mathcal{F}_{S_+}(T)$ converges for every $\lambda \in \mathbb{R}$; for $T \in (2, 5)$, each displayed infimum is a finite real and is uniquely attained.

There is a deterministic sequence of nonnegative reals $(\varepsilon_n)_{n \geq 2}$ with $\varepsilon_n \to 0$ such that

$$\mathbb{P}\left(\chi(G_n) - \zeta(G_n) \geq \left\lceil \frac{(\log 2)^2}{4} A_4(\delta_n) - \varepsilon_n \right\rceil \frac{n}{(\log n)^3}\right) \longrightarrow 1.$$

Moreover, A_4 is continuous and nonconstant on $[0, 1]$, and for every $\delta \in [0, 1]$,

$$A_4(\delta) > \log\left(\frac{200}{153}\right).$$

Continuity and compactness therefore give a positive uniform margin above this bound.

The uniform coefficient displayed above is

$$\frac{(\log 2)^2}{4} \log\left(\frac{200}{153}\right) = 0.032175871697936\ldots.$$

The uniform bound uses a single constant. The refinement keeps $A_4(\delta_n)$, so its leading coefficient depends on where the independent-set center lies between neighbouring integers. The positive uniform margin absorbs the vanishing error for all sufficiently large integers. Since $n/(\log n)^3 \to \infty$, both forms resolve the question of Erdős and Gimbel along the full sequence.

The class sizes are integers, but their center is real. A profile records how many classes have each size. When $\alpha_0(n)$ crosses an integer, the cutoff and admissible sizes change, so the optimization problem determining the profile and its first-moment root changes too. A bound for most integers can still miss a phase window. To cover the full sequence, we keep the phase in the estimates and control them uniformly up to both endpoints, including integers on either side of every cutoff change.

Relation to previous work. The first-order asymptotic for the chromatic number of a dense random graph was proved by Bollobás (Bollobás, 1988), following the early work of Grimmett and McDiarmid (Grimmett and McDiarmid, 1975); later refinements include McDiarmid (1990), Panagiotou and Steger (2009), and Heckel (2018). The cochromatic number belongs to the theory of generalized chromatic numbers associated with hereditary graph properties (Scheinerman, 1992; Bollobás and Thomason, 1995).

For the difference $\chi(G_n) - \zeta(G_n)$, Heckel and, independently, Steiner related divergence to the nonconcentration of the chromatic number (Heckel, 2024; Steiner, 2025). Heckel subsequently proved a near-linear lower bound for a phase-dependent set containing approximately 95% of the integers and conjectured that the natural full-sequence scale is $n/(\log n)^3$ (Heckel, 2025). We prove this lower bound along the full sequence $n \to \infty$, with an explicit phase-resolved coefficient. We start from Heckel's gain obtained by declaring each class independent or complete, and from her probability-amplification argument. Four consecutive class sizes make our optimization uniform in the phase. The second moment counts the compatible class declarations exactly and controls all overlaps between two partitions, including whole classes they share.

Proof strategy. We construct deterministic integers k_χ^- and k_{co}, and a deterministic $a_n = o(n/(\log n)^3)$, for which

(A)
$$k_\chi^- - k_{\mathrm{co}} - a_n = \left[\frac{(\log 2)^2}{4} A_4(\delta_n) + o(1)\right] \frac{n}{(\log n)^3};$$

(B)
$$\mathbb{P}\big(\chi(G_n) > k_\chi^-\big) \longrightarrow 1;$$

(C)
$$\mathbb{P}\big(\zeta(G_n) \le k_{\mathrm{co}} + a_n\big) \longrightarrow 1.$$

The intersection of the events in (B) and (C), together with (A), gives the phase-resolved assertion of the main theorem; no independence between the two events is required.

Sections 1–5 compare the first-moment roots and prove (B). Sections 6–9 bound the second moment. Section 10 constructs a_n and proves (C). The final section completes (A) and combines the bounds.

The two counting difficulties. Two partitions may share whole classes. They may also have large intersections between classes that are not equal. Section 7 controls the whole-class overlaps by exact recurrences near the empty and full choices, and by a negative exponential rate between them. The estimate is uniform over the phase, including the range omitted in the earlier argument.

For the remaining overlaps, we form a table whose entries are the sizes of intersections between a class of each partition. Entries larger than half the maximum class size cannot share a row or column, so they form a matching. Section 8 bounds their contribution; Section 9 handles the remaining entries and the compatibility of the class declarations. The chosen vertex pairs are retained throughout, so every labelled configuration is counted once. If Z counts the marked cocolorings of the chosen profile, the resulting second-moment bound is

$$\frac{\mathbb{E}Z^2}{(\mathbb{E}Z)^2} \le \exp\left\{o\left(\frac{n}{(\log n)^4}\right)\right\}.$$

Appendix H records the companion Lean theorem and its reproducibility materials.

Notation and counting conventions

Notation specific to an estimate is introduced where it is used.

Asymptotic and counting conventions. An event holds *with high probability* if its probability tends to one as $n \to \infty$. Whenever a phase parameter is present, each occurrence of $o(1)$ denotes a deterministic error sequence that is uniform over the full phase interval, including sequences approaching either endpoint. For integers $x, r \ge 0$, we use the falling factorial

$$(x)_r = x(x-1)\cdots(x-r+1), \qquad (x)_0 = 1,$$

and set $(x)_r = 0$ for $r > x$. A quotient involving falling factorials is used only after the denominator has been shown to be nonzero.

Class marks and profiles. For counting, mark each class of a partition by I or K, requiring it to be independent or complete, respectively. All classes in the selected profile have size at least two for sufficiently large n, so these two requirements are disjoint. Forgetting the marks recovers the cocoloring without multiplicity.

A *profile* is a finitely supported sequence $(k_s)_{s\ge1}$ of nonnegative integers, where k_s is the number of classes of size s (Heckel, 2025, Section 2.2). It is feasible on n vertices with k parts when

$$\sum_s k_s = k, \qquad \sum_s s k_s = n.$$

The ordinary first moment counts profile partitions whose classes are all required to be independent. The marked first moment counts profile partitions together with one I/K mark per class, subject to the corresponding independence or completeness requirement.

Overlap table and support graph. Let $(A_a)_a$ and $(B_b)_b$ be two ordered profile partitions. Their overlap table is

$$r_{ab} = |A_a \cap B_b|.$$

Its row and column sums are the class sizes of the two partitions. In the bipartite configuration model, we represent each vertex by one labelled placeholder, or stub, in its row class and one in its column class. Pairing the two stubs records which classes contain that vertex.

The *support graph* $H(r)$ is the simple bipartite graph whose vertices are the row and column classes incident to an edge and whose edge set is

$$E(H(r)) = \{(a, b) : r_{ab} \geq 2\}.$$

A subset of its edges is *even* if every vertex has even degree. Writing $c(H)$ for the number of connected components, with $c(\varnothing) = 0$, we put

$$\beta(H) = |E(H)| - |V(H)| + c(H).$$

Then $\beta(H)$ is the dimension of the binary cycle space, so the number of even edge sets is $2^{\beta(H)}$. The threshold two is exact: a cell of size zero or one contains no edge internal to both partition classes, whereas a cell of size at least two forces the marks on its row and column classes to agree.

The large-cell threshold, deficit notation, endpoint tables, and the precise separation between exposed matching and residual factors are introduced where they are first used in Sections 8–9.

1. Independent-set counts and probability inequalities

For integers $v, s \geq 0$, with $\binom{v}{s} = 0$ when $s > v$, let

$$\mu_s(v) = \binom{v}{s} 2^{-\binom{s}{2}}, \qquad \mu_s = \mu_s(n). \tag{1.1}$$

The proof needs the following standard inequalities. We state the forms used below to fix their hypotheses and constants.

1. For integers $m \geq 1$,
$$\log(m!) = m \log m - m + \frac{1}{2} \log(2\pi m) + O(1/m). \tag{1.2}$$
Consequently, with $0 \log 0 = 0$, one has $\log(m!) = m \log m - m + O(\log(m + 1))$ uniformly for $m \in \mathbb{N}_0$, with an absolute implied constant.
2. If $r \geq 1$, $t \geq 0$, and a random variable Y is a function of r independent blocks and changing one block changes Y by at most one, then
$$\mathbb{P}(|Y - \mathbb{E}[Y]| \geq t) \leq 2 \exp(-2t^2/r). \tag{1.3}$$
We will use the corresponding one-sided bounds as well.
3. If $Z \geq 0$, $0 < \mathbb{E}[Z] < \infty$, and $\mathbb{E}\left[Z^2\right] < \infty$, then
$$\mathbb{P}(Z > 0) \geq \frac{\mathbb{E}[Z]^2}{\mathbb{E}\left[Z^2\right]}. \tag{1.4}$$
4. If $X \sim \mathrm{Bin}(m, 1/2)$, then
$$\mathbb{P}(X \leq m/4) \leq e^{-m/16}. \tag{1.5}$$
Indeed, exponential Markov with $t = \log 3$ bounds the probability by $[3^{1/4}(2/3)]^m \leq e^{-m/16}$.
5. For every nonnegative random variable X and $a > 0$, $\mathbb{P}(X \geq a) \leq \mathbb{E}[X]\,/a$.

The first is Stirling's estimate, the second is McDiarmid's bounded-differences inequality, and the third is the zero-threshold case of Paley–Zygmund. The fourth is the only binomial-tail estimate used below; the fifth is Markov's inequality. The bounded-differences method is due to McDiarmid (1989); the one-sided formulation used here is McDiarmid (1998, Theorem 3.1).

2. Uniform first-moment estimates across the rounding phase

Put

$$q := \log 2, \qquad L := \log n, \qquad \ell := \log\log n, \qquad C := 1 + \log q - q.$$

We shall use the elementary bounds

$$\frac{2}{3} < q < \frac{7}{10}. \tag{2.0}$$

The location parameter α_0 defined in the introduction has the following equivalent form in natural logarithms (Heckel, 2025, Section 2.1):

$$\alpha_0 = 2\log_2 n - 2\log_2\log_2 n + 2\log_2(\mathrm{e}/2) + 1 = \frac{2(L - \ell + C)}{q} + 1. \tag{2.1}$$

Let

$$\alpha := \lfloor \alpha_0 \rfloor, \qquad \delta := \alpha_0 - \alpha, \qquad b := 1 - \delta.$$

Thus $0 \le \delta < 1$, $0 < b \le 1$, and

$$\alpha = \frac{2(L - \ell + C)}{q} + b.$$

When emphasizing dependence on n, we write $\delta_n := \delta$. We prove the estimates on the closed parameter range $0 \le \delta \le 1$. Uniformity on this closure also covers integer sequences approaching either endpoint of the actual half-open phase interval.

Lemma 2.1 (Uniform phase expansion and adjacent-size control)**.** *There exist a bounded continuous function $K\colon [0,1] \to \mathbb{R}$, absolute constants $C_{\mathrm{ph}}, C_2, c > 0$, and a real sequence $(E_n^{\mathrm{ph}})_{n\ge 2}$ such that, with*

$$\varepsilon_n^{\mathrm{ph}} := C_{\mathrm{ph}} \frac{1 + \ell^2}{L},$$

one has $\varepsilon_n^{\mathrm{ph}} \to 0$ and, for all sufficiently large n, the following expansion holds at the induced phase $\delta_n = \alpha_0 - \lfloor \alpha_0 \rfloor$:

$$\log \mu_\alpha = \delta_n L + \left(\frac{2}{q} - \frac{1}{2} - \delta_n\right)\ell + K(\delta_n) + E_n^{\mathrm{ph}}, \qquad |E_n^{\mathrm{ph}}| \le \varepsilon_n^{\mathrm{ph}}. \tag{2.2}$$

The error bound is uniform over all induced phase values, and therefore also along sequences approaching either endpoint of $[0,1]$. Moreover, uniformly over those phase values,

$$\begin{aligned} \log \mu_{\alpha+2} &= (\delta_n - 2)L + \left(\frac{2}{q} + \frac{3}{2} - \delta_n\right)\ell + O(1), \\ \mu_{\alpha+2} &\le \exp\{-L + C_2\ell\} = o(1) \end{aligned} \tag{2.3}$$

while

$$\mu_{\alpha-2} \ge c\, n^2 (\log n)^{2/q - 5/2} \tag{2.4}$$

for all sufficiently large n.

Stirling's formula gives the phase expansion with a uniform remainder. The exact ratios for adjacent class sizes then give both the upper cap and the lower first-moment bound. Appendix D contains the remainder estimates and the complete ratio calculation.

Let $\alpha(G)$ denote the independence number of a graph G, and let $X_{\alpha+2}$ denote the number of independent sets of size $\alpha + 2$. Then $\mathbb{E}X_{\alpha+2} = \mu_{\alpha+2}$, and the event $\alpha(G_n) > \alpha + 1$ implies $X_{\alpha+2} \ge 1$. Consequently Markov's inequality and (2.3) give, uniformly along the full sequence of integers, the deterministic cap error

$$\mathbb{P}\big(\alpha(G_n) > \alpha + 1\big) \le \mu_{\alpha+2} =: \varepsilon_n^{\mathrm{cap}} \longrightarrow 0 \tag{2.9}$$

3. Real profile optimization and first-moment roots

We first allow real type counts and optimize the counting exponent under the two conservation constraints. We call its zeros first-moment roots. The ordinary first-moment root will give a lower location for the chromatic number. The four-size first-moment root for marked partitions will give the number of parts in our cocoloring. Section 5 compares these roots before rounding the type counts; Proposition 5.3 collects the properties of the resulting integer profile. Throughout this section, n is sufficiently large that $\alpha \geq 6$. For an integer class size $u \geq 1$, put

$$d_u := 2^{\binom{u}{2}} u!.$$

Write $u = \alpha - i$, so i is the deficit from the phase center α. We use

$$S_+ := \{-1, 0, 1, 2, \ldots\}, \qquad S_4 := \{2, 3, 4, 5\}. \tag{3.1}$$

At finite n, the unrestricted support is $S_+^{(n)} := \{-1, 0, \ldots, \alpha - 1\}$: the lower endpoint comes from the cap event (2.9), and the upper endpoint from positivity of class sizes. For $S \in \{S_+, S_4\}$, set

$$S^{(n)} := \begin{cases} S_+^{(n)}, & S = S_+, \\ S_4, & S = S_4. \end{cases}$$

Every finite-n profile, maximization, and sum below uses $S^{(n)}$; limiting quantities use the displayed support S itself.

The class-size contribution will be split into an affine part and a remainder. For $i \in S_+^{(n)}$, define it by

$$h_n(i) := \begin{cases} -\dfrac{\log 2}{2} i^2 + \displaystyle\sum_{r=0}^{i-1} \log\Big(1 - \frac{r}{\alpha}\Big), & i \geq 0, \\ -\dfrac{\log 2}{2} + \log\Big(\dfrac{\alpha}{\alpha+1}\Big), & i = -1. \end{cases} \tag{3.1a}$$

The affine part is fixed by the mean constraint:

$$-\log d_{\alpha - i} = A_n + B_n i + h_n(i), \tag{3.1b}$$

where

$$A_n := -\log d_\alpha, \qquad B_n := q\alpha - \frac{q}{2} + \log \alpha, \tag{3.1c}$$

Direct subtraction gives this identity also at $i = -1$.

For $S \in \{S_+, S_4\}$ and a nonnegative real profile $\mathbf{k} = (k_i)_{i \in S^{(n)}}$, put $k := \sum_{i \in S^{(n)}} k_i$. When $k > 0$, define

$$L_{\mathbf{k}} := n \log n - n + k - \sum_{i \in S^{(n)}} k_i \log(k_i d_{\alpha - i}), \tag{3.2}$$

with the convention $0 \log 0 = 0$, and let $L_S(n, k)$ be its maximum over real $k_i \geq 0$ satisfying

$$\sum_{i \in S^{(n)}} k_i = k, \qquad \sum_{i \in S^{(n)}} (\alpha - i) k_i = n. \tag{3.3}$$

For k feasible for the support S under consideration, set

$$s := \frac{n}{k}, \qquad T := \alpha - s = \alpha - \frac{n}{k}. \tag{3.4}$$

With $p_i := k_i / k$, the constraints give $T = \sum_{i \in S^{(n)}} i p_i$; thus T is the mean deficit of the normalized profile.

For $S \in \{S_+, S_4\}$, the positive weights below are normalized by their sum, called the partition function. The mean map gives the mean deficit under the resulting distribution. Define the finite and limiting versions by

$$Z_{n,S}(\lambda) := \sum_{i \in S^{(n)}} \exp\{\lambda i + h_n(i)\}, \qquad M_{n,S}(\lambda) := \frac{d}{d\lambda} \log Z_{n,S}(\lambda),$$

$$Z_S(\lambda) := \sum_{i \in S} \exp\Big(\lambda i - \frac{q}{2} i^2\Big), \qquad M_S(\lambda) := \frac{d}{d\lambda} \log Z_S(\lambda).$$

Whenever the equations $M_{n,S}(\lambda) = T$ and $M_S(\lambda) = T$ have unique solutions, denote them by $\lambda_{n,S}(T)$ and $\lambda_S(T)$, respectively.

Lemma 3.1 (Uniform root, slope, and finite-dual estimates)**.** *Let*

$$s_0 := \alpha_0 - 1 - \frac{2}{\log 2}, \qquad T_0 := \alpha - s_0 = 1 + \frac{2}{\log 2} - \delta. \tag{3.5}$$

There exist a fixed $A_* > 0$ and an integer $n_{\mathrm{root}} \geq 2$ such that, for every $n \geq n_{\mathrm{root}}$, every $S \in \{S_+, S_4\}$, and every $0 \leq c \leq \log 2$, the equation

$$L_S(n,k) + ck = 0$$

has a unique feasible zero $r_{S,c}$ in the corridor

$$\left| \frac{n}{r_{S,c}} - s_0 \right| \leq A_* \frac{\log\log n}{\log n}.$$

No uniqueness outside this corridor is asserted. Uniformly in S, c, and the phase,

$$\frac{n}{r_{S,c}} = s_0 + O\left(\frac{\log\log n}{\log n}\right), \qquad \alpha - \frac{n}{r_{S,c}} = T_0 + O\left(\frac{\log\log n}{\log n}\right). \tag{3.6}$$

For every fixed $A > 0$, uniformly over $S \in \{S_+, S_4\}$, $0 \leq c \leq \log 2$, and k feasible for S with $|n/k - s_0| \leq A \log\log n/\log n$,

$$\frac{\partial}{\partial k}\{L_S(n,k) + ck\} = \frac{2}{\log 2}(\log n)^2 + O_A((\log n)(\log\log n)). \tag{3.7}$$

In particular, the deterministic normalized slope error

$$\varepsilon_{n,A}^{\mathrm{slope}} := \max_{\substack{S \in \{S_+,S_4\} \\ 0 \leq c \leq \log 2}} \sup_{\substack{k \text{ feasible for } S \\ |n/k - s_0| \leq A \log\log n/\log n}} \left| \frac{1}{(\log n)^2} \frac{\partial}{\partial k}\{L_S(n,k) + ck\} - \frac{2}{\log 2} \right|$$

satisfies $\varepsilon_{n,A}^{\mathrm{slope}} \to 0$.

For two supports $R, S \in \{S_+, S_4\}$ at a k feasible for both finite supports $R^{(n)}$ and $S^{(n)}$,

$$\frac{L_R(n,k) - L_S(n,k)}{k} = \mathcal{F}_{n,R}(T) - \mathcal{F}_{n,S}(T), \tag{3.8}$$

where

$$\mathcal{F}_{n,S}(T) := \max_{\substack{p_i \geq 0 \ (i \in S^{(n)}) \\ \sum_{i \in S^{(n)}} p_i = 1 \\ \sum_{i \in S^{(n)}} i p_i = T}} \left[-\sum_{i \in S^{(n)}} p_i \log p_i + \sum_{i \in S^{(n)}} p_i h_n(i) \right]. \tag{3.8a}$$

Fix the compact target interval

$$K_* := \left[\frac{2}{\log 2} - \frac{1}{10}, 1 + \frac{2}{\log 2} + \frac{1}{10} \right] \subset (2,5). \tag{3.9}$$

Then, uniformly for $T \in K_*$,

$$\mathcal{F}_{n,S}(T) \longrightarrow \mathcal{F}_S(T),$$

where $\mathcal{F}_S$ is the function defined in the main theorem; equivalently,

$$\mathcal{F}_S(T) = \max_{\substack{p_i \geq 0 \ (i \in S) \\ \sum_{i \in S} p_i = 1 \\ \sum_{i \in S} i p_i = T \\ \sum_{i \in S} i^2 p_i < \infty}} \left[-\sum_{i \in S} p_i \log p_i - \frac{\log 2}{2} \sum_{i \in S} i^2 p_i \right]. \tag{3.9a}$$

For every $S \in \{S_+, S_4\}$ and $T \in K_*$, the limiting mean equation $M_S(\lambda) = T$ has the unique solution $\lambda_S(T)$. There exist a fixed $\Lambda > 0$ and an integer n_{dual} such that, for every $n \geq n_{\mathrm{dual}}$, every such S, and every $T \in K_*$, the finite equation $M_{n,S}(\lambda) = T$ has the unique solution $\lambda_{n,S}(T)$ and

$$|\lambda_{n,S}(T)| + |\lambda_S(T)| \leq 2\Lambda.$$

For $n \geq n_{\mathrm{dual}}$, define the deterministic dual error

$$\varepsilon_n^{\mathrm{dual}} := \max_{S \in \{S_+,S_4\}} \sup_{T \in K_*} \big(|\mathcal{F}_{n,S}(T) - \mathcal{F}_S(T)| + |\lambda_{n,S}(T) - \lambda_S(T)| \big)$$

and set it arbitrarily, say equal to zero, for $n < n_{\mathrm{dual}}$. Then

$$\varepsilon_n^{\mathrm{dual}} = O(1/\alpha) = O(1/\log n),$$

with a phase-independent implied constant. In particular, $\varepsilon_n^{\mathrm{dual}} \to 0$.

For S_4, both optimizers are unique; denote them by $p_{n,i}(T)$ and $p_i(T)$. Then

$$\sup_{T\in K_*} \max_{2\le i\le 5} |p_{n,i}(T) - p_i(T)| = O(1/\alpha). \tag{3.9b}$$

Moreover, there are $\eta_* > 0$ and $n_{\mathrm{int}} \in \mathbb{N}$ such that

$$p_{n,i}(T) \ge \eta_* \qquad (n \ge n_{\mathrm{int}},\ T \in K_*,\ 2 \le i \le 5).$$

All displayed bounds and eventuality thresholds are uniform over the induced phases δ_n, including sequences approaching either endpoint.

The proof is in Appendix A. It first controls the finite exponential tilts uniformly on K_*, then applies Stirling's formula to locate the zeros and bound their slopes.

4. A uniform lower bound for the chromatic number χ

Let $r_+(n) = r_{S_+,0}$ be the unrestricted-support corridor zero from Lemma 3.1, write $L_+(n,k) := L_{S_+}(n,k)$, and put

$$k_\chi^- := \lfloor r_+(n) \rfloor - \lceil \log n \rceil. \tag{4.1}$$

This deterministic integer is chosen a logarithmic distance below the ordinary first-moment root. The slope estimate will make the resulting decrease in the counting exponent dominate the error in the profile sum. Since $r_+(n) = \Theta(n/\log n)$ uniformly in the phase, one has $1 \le k_\chi^- < n$ for all sufficiently large n.

For an integer profile $\mathbf{k} = (k_i)_{-1\le i\le \alpha-1}$ feasible on n vertices, let $X_{\mathbf{k}}$ be the number of unordered proper colorings having exactly k_i classes of size $\alpha - i$. Direct enumeration, as in Heckel (2025, Section 2.3), gives

$$\mathbb{E}[X_{\mathbf{k}}] = \frac{n!}{\prod_i ((\alpha-i)!)^{k_i} k_i!} 2^{-\sum_i k_i \binom{\alpha-i}{2}}. \tag{4.2}$$

For an integer k, let $\mathcal{K}_{n,k,\alpha+1}$ be the set of these profiles satisfying

$$\sum_i k_i = k, \qquad \sum_i (\alpha - i) k_i = n,$$

and define

$$E_{n,k,\alpha+1} := \sum_{\mathbf{k}\in\mathcal{K}_{n,k,\alpha+1}} \mathbb{E}[X_{\mathbf{k}}]\,.$$

Thus $E_{n,k,\alpha+1}$ is the expected number of unordered proper colorings with exactly k nonempty parts, each of size at most $\alpha + 1$. When the feasible profile set is empty, we use the extended-real convention $\log E_{n,k,\alpha+1} = -\infty$.

Summing the capped coloring profiles. There are $\alpha + 1 = O(\log n)$ profile coordinates, and each coordinate lies between zero and n. Hence

$$|\mathcal{K}_{n,k,\alpha+1}| \le (n+1)^{\alpha+1} = \exp\{O((\log n)^2)\} \tag{4.3}$$

uniformly in k and in the phase. Keep the class-size factorials inside $d_{\alpha-i}$ exact, and put

$$R(m) := \log(m!) - (m \log m - m), \qquad R(0) := 0.$$

For every feasible profile, exact cancellation in (4.2) gives

$$\log \mathbb{E}[X_{\mathbf{k}}] - L_{\mathbf{k}} = R(n) - \sum_i R(k_i).$$

The uniform estimate (1.2) therefore gives

$$\log \mathbb{E}[X_{\mathbf{k}}] \le L_{\mathbf{k}} + O((\log n)^2).$$

The error is uniform even when some $k_i = 0$: there are $O(\log n)$ profile coordinates and $|R(k_i)| = O(\log(n+1))$. For every k feasible for $S_+^{(n)}$, maximizing over the feasible profile and then summing (4.3) yields

$$\log E_{n,k,\alpha+1} \le L_+(n,k) + O((\log n)^2). \tag{4.3a}$$

Displacement below the root. Set

$$\Delta_n := r_+(n) - k_\chi^-.$$

The floor and ceiling in (4.1) give the deterministic bounds

$$\log n \le \Delta_n < \log n + 2. \tag{4.3b}$$

We next verify that the entire interval $[k_\chi^-, r_+(n)]$ lies in the derivative corridor of Lemma 3.1. Write $s(k) = n/k$. Uniformly on this interval, $k = \Theta(n/\log n)$ and $s(k) = \Theta(\log n)$. Therefore

$$|s(k_\chi^-) - s(r_+)| \le \sup_{k\in[k_\chi^-, r_+]} \frac{n}{k^2}\,\Delta_n = O\!\left(\frac{(\log n)^3}{n}\right) = o\!\left(\frac{\log\log n}{\log n}\right).$$

Thus, after one phase-independent threshold, every point of the interval satisfies

$$\left|\frac{n}{k} - s_0\right| \le (A_* + 1)\frac{\log\log n}{\log n}.$$

Equation (3.7), with $A = A_* + 1$, $S = S_+$, and $c = 0$, consequently gives an absolute $c_* > 0$ such that

$$\frac{d}{dk}L_+(n,k) \ge c_*(\log n)^2 \qquad (k_\chi^- \le k \le r_+) \tag{4.3c}$$

for all sufficiently large n, uniformly in the phase. Since $L_+(n, r_+) = 0$, the mean-value theorem and (4.3b) imply

$$L_+(n, k_\chi^-) \le -c_*(\log n)^2\Delta_n \le -c_*(\log n)^3. \tag{4.4}$$

Combining (4.3a) and (4.4), and decreasing the constant once, gives an absolute constant $c_\chi > 0$ and a deterministic sequence

$$\varepsilon_n^{\mathrm{prof}} := \exp\{-c_\chi(\log n)^3\} \longrightarrow 0$$

such that

$$E_{n,k_\chi^-,\alpha+1} \le \varepsilon_n^{\mathrm{prof}}. \tag{4.4a}$$

Removing the size cap. Let

$$\mathcal{A}_n := \{\alpha(G_n) \le \alpha + 1\}.$$

Equation (2.9) gives a deterministic phase-uniform sequence $\varepsilon_n^{\mathrm{cap}} \to 0$ such that

$$\mathbb{P}(\mathcal{A}_n^c) \le \varepsilon_n^{\mathrm{cap}}.$$

On $\mathcal{A}_n$, every class of every proper coloring has size at most $\alpha + 1$.

Suppose that $\chi(G_n) \le k_\chi^-$ and choose a proper coloring with $h \le k_\chi^-$ nonempty classes. If $h < k_\chi^-$, then some class contains at least two vertices: otherwise $h = n$, contradicting $h < k_\chi^- < n$. Splitting such a class into two nonempty subsets preserves independence. Repeating this operation produces a proper coloring with exactly k_χ^- nonempty classes. The size cap is preserved under splitting. Therefore

$$\{\chi(G_n) \le k_\chi^-\} \cap \mathcal{A}_n \subseteq \{\text{an exactly } k_\chi^-,\ (\alpha+1)\text{-bounded coloring exists}\}.$$

Markov's inequality and (4.4a) now give

$$\begin{aligned}\mathbb{P}\Big(\chi(G_n) \le k_\chi^-\Big) &\le \mathbb{P}(\mathcal{A}_n^c) + E_{n,k_\chi^-,\alpha+1}\\ &\le \varepsilon_n^{\mathrm{cap}} + \varepsilon_n^{\mathrm{prof}} \longrightarrow 0.\end{aligned} \tag{4.5}$$

Equivalently,

$$\mathbb{P}\Big(\chi(G_n) > k_\chi^-\Big) \longrightarrow 1. \tag{4.6}$$

Finally, (4.3b) also records the location precision:

$$|k_\chi^- - r_+(n)| < \log n + 2 = o\!\left(\frac{n}{(\log n)^3}\right). \tag{4.7}$$

Thus both the probability statement and the location error hold with one set of phase-independent thresholds along the full sequence of integers.

5. Four-size entropy gain and the selected integer profile

The choice $S_4 = \{2, 3, 4, 5\}$ costs profile entropy but makes the later overlap estimates uniform. We must show that this cost is less than $\log 2$ per part, so that the 2^k sign choices leave a positive gain. Lemma 5.2 gives the required margin. The slope estimate in Lemma 3.1 then converts this gain in the counting exponent into a distance between its roots.

5.1. **The phase coefficient and entropy bound.** For $S \in \{S_+, S_4\}$, introduce the tilted weights, partition function, and mean map

$$w_i(\lambda) := \exp\left(\lambda i - \frac{\log 2}{2} i^2\right), \qquad Z_S(\lambda) := \sum_{i \in S} w_i(\lambda),$$
$$M_S(\lambda) := \frac{\sum_{i \in S} i\, w_i(\lambda)}{Z_S(\lambda)} = \frac{d}{d\lambda} \log Z_S(\lambda).$$

The Gaussian factor makes both sums converge. Moreover, $M_S'(\lambda) = \mathrm{Var}_{S,\lambda}(i) > 0$, and the endpoint limits of M_S are the endpoints of the convex hull of S. Thus, for every $T \in (2, 5)$, there is a unique tilt $\lambda_S(T)$ with $M_S(\lambda_S(T)) = T$. The Gibbs identity (A.3) proves that the optimizer in (3.9a) is

$$p_i = \frac{w_i(\lambda_S(T))}{Z_S(\lambda_S(T))} = \frac{e^{\lambda_S(T) i - \log 2 i^2/2}}{\sum_{j \in S} e^{\lambda_S(T) j - \log 2 j^2/2}}. \tag{5.1}$$

This also gives the dual representation

$$\mathcal{F}_S(T) = \inf_{\lambda \in \mathbb{R}} \{\log Z_S(\lambda) - \lambda T\} = \log Z_S(\lambda_S(T)) - \lambda_S(T) T.$$

Since $T_0 = 1 + 2/\log 2 - \delta = T_\delta$, introduce the entropy loss D_4 and the uniform bound γ_4 by

$$D_4(\delta) := \mathcal{F}_{S_+}(T_\delta) - \mathcal{F}_{S_4}(T_\delta), \qquad A_4(\delta) = \log 2 - D_4(\delta), \qquad \gamma_4 := \log \frac{200}{153}. \tag{5.2}$$

The restriction to four consecutive deficits also locates the smallest phase gain. At the same mean deficit, the four-size optimizer has smaller variance than the unrestricted optimizer. This gives the following curvature formula.

Lemma 5.1 (Curvature and minimum of the phase gain)**.** *For $S \in \{S_+, S_4\}$ and $T \in (2, 5)$, let*

$$V_S(T) := \sum_{i \in S} (i - T)^2 \frac{w_i(\lambda_S(T))}{Z_S(\lambda_S(T))}.$$

Then $0 < V_{S_4}(T) < V_{S_+}(T)$. The function A_4 extends to a twice continuously differentiable function on an open interval containing $[0, 1]$, and

$$A_4'(\delta) = \lambda_{S_4}(T_\delta) - \lambda_{S_+}(T_\delta), \qquad A_4''(\delta) = \frac{1}{V_{S_+}(T_\delta)} - \frac{1}{V_{S_4}(T_\delta)} < 0.$$

In particular, A_4 is strictly concave and nonconstant on $[0, 1]$, and

$$\min_{0 \le \delta \le 1} A_4(\delta) = \min\{A_4(0), A_4(1)\}.$$

Proof. The Gaussian factor permits termwise differentiation locally in λ. Since $M_S'(\lambda_S(T)) = V_S(T) > 0$, differentiating the inverse mean equation and the dual value gives

$$\lambda_S'(T) = \frac{1}{V_S(T)}, \qquad \mathcal{F}_S'(T) = -\lambda_S(T), \qquad \mathcal{F}_S''(T) = -\frac{1}{V_S(T)}.$$

For the variance comparison, fix T and write p_+ and p_4 for the two probability vectors in (5.1), extending p_4 by zero off S_4. On S_4,

$$\frac{p_4(i)}{p_+(i)} = \frac{Z_{S_+}(\lambda_{S_+}(T))}{Z_{S_4}(\lambda_{S_4}(T))} e^{(\lambda_{S_4}(T) - \lambda_{S_+}(T)) i}.$$

Within the four-size support, this ratio is monotone in the index. Thus the indices where $d_i := p_4(i) - p_+(i) > 0$ form a consecutive block $\{j, \dots, k\} \subseteq S_4$. This block is nonempty because p_+ has positive mass outside S_4 and both vectors have total mass one. Put $a = j - 1/2$ and $b = k + 1/2$. Then $(i - a)(i - b) d_i \le 0$ for every $i \in S_+$: the quadratic factor is negative on the block and positive at all other integer indices. The

inequality is strict at $i = j$. Both vectors have mass one and mean T, so the constant and linear terms of the quadratic cancel in the sum:

$$V_{S_4}(T) - V_{S_+}(T) = \sum_{i \in S_+} i^2 d_i = \sum_{i \in S_+} (i-a)(i-b) d_i < 0.$$

All sums converge absolutely. The nonstrict inequality is also a standard consequence of the equal-mean relative log-concavity comparison (Yu, 2010, Theorem 12); the quadratic argument above verifies strictness here.

The derivative formulas now follow from $T'_\delta = -1$. Since $T_\delta \in (2, 5)$ on a neighborhood of $[0, 1]$, they also apply at both endpoints. Strict concavity gives the endpoint minimum. □

Lemma 5.2 (Uniform entropy bound)**.** *For every* $0 \le \delta \le 1$,

$$0 \le D_4(\delta) < \log \frac{153}{100}, \qquad A_4(\delta) > \gamma_4. \tag{5.3}$$

To bound the entropy loss, evaluate the unrestricted dual at the four-size minimizing tilt. A fixed bracket for this tilt lets us control the omitted weights on either side of the four-size support. Appendix E gives the tilt comparisons, the rational bounds and the Gaussian-tail calculation.

5.2. **The gain from class marks and root separation.** The following 2^k gain is the one-partition identity isolated by Heckel (2025, Proposition 6). We repeat its short argument because it is an input to the new root calculation. For a fixed partition into k classes, assigning each class the declaration "independent" or "complete" multiplies its first moment by 2^k: every declaration prescribes all internal edge bits, and all 2^k declarations have probability $2^{-\sum \binom{u}{2}}$. The declarations are disjoint because every allowed class size is at least two for all sufficiently large n. Let $r_4^{\mathrm{co}} = r_{S_4, \log 2}$ be the unique corridor zero of

$$L_{S_4}(n, k) + \log 2k. \tag{5.9}$$

Put

$$\Phi_n(k) := L_{S_4}(n, k) + (\log 2)k.$$

Then $\Phi_n(r_4^{\mathrm{co}}) = 0$ by definition. To compare the two roots, evaluate the same function at the unrestricted root. Let

$$T_+(n) := \alpha - \frac{n}{r_+(n)}$$

and introduce the finite-n support loss

$$D_{4,n}(T) := \mathcal{F}_{n,S_+}(T) - \mathcal{F}_{n,S_4}(T).$$

By (3.6) and the fixed margin in (3.9), one has $T_+(n) \in K_* \subset (2, 5)$ for all sufficiently large n, uniformly in the phase. Thus r_+ is feasible for both supports before (3.8) is applied. Since $L_{S_+}(n, r_+) = 0$, the exact support-comparison identity (3.8) gives

$$\begin{aligned} \Phi_n(r_+) &= L_{S_4}(n, r_+) - L_{S_+}(n, r_+) + (\log 2) r_+ \\ &= r_+ \{\log 2 - D_{4,n}(T_+(n))\}. \end{aligned} \tag{5.9a}$$

For finite n, we have expressed the signed exponent at the unrestricted root in terms of the finite support loss. It remains to replace that loss by its phase-dependent limit.

There are two errors in this replacement: the finite dual values differ from their limits, and the root target $T_+(n)$ differs from the phase-center target T_0. For the latter, define

$$\varepsilon_n^{\mathrm{target}} := |T_+(n) - T_0|.$$

Equation (3.6) gives the equivalent phase-uniform statements

$$T_+(n) = T_0 + O\left(\frac{\log\log n}{\log n}\right), \qquad \varepsilon_n^{\mathrm{target}} = O\left(\frac{\log\log n}{\log n}\right) \longrightarrow 0. \tag{5.9b}$$

The common corridor lies inside the fixed compact interval K_* for all sufficiently large n.

The limiting dual value is differentiable, with

$$\mathcal{F}'_S(T) = -\lambda_S(T),$$

because the envelope theorem cancels the derivative of the optimizing tilt. Section 3 bounds the limiting tilts by one constant on K_*, so the two limiting dual values have a common Lipschitz constant C_{Lip}. The deterministic finite-dual error from Section 3 satisfies

$$\sup_{T\in K_*} |\mathcal{F}_{n,S}(T) - \mathcal{F}_S(T)| \le \varepsilon_n^{\mathrm{dual}} \qquad (S = S_+, S_4).$$

Consequently the single root error used in the final theorem,

$$\omega_n^{\mathrm{root}} := 2\varepsilon_n^{\mathrm{dual}} + 2C_{\mathrm{Lip}}\varepsilon_n^{\mathrm{target}} \longrightarrow 0 \tag{5.9c}$$

is deterministic and phase-uniform. Indeed,

$$\begin{aligned} |D_{4,n}(T_+(n)) - D_4(\delta)| \le\ & |\mathcal{F}_{n,S_+}(T_+) - \mathcal{F}_{S_+}(T_+)| \\ & + |\mathcal{F}_{n,S_4}(T_+) - \mathcal{F}_{S_4}(T_+)| \\ & + |\mathcal{F}_{S_+}(T_+) - \mathcal{F}_{S_+}(T_0)| \\ & + |\mathcal{F}_{S_4}(T_+) - \mathcal{F}_{S_4}(T_0)| \\ \le\ & \omega_n^{\mathrm{root}}. \end{aligned}$$

The same ω_n^{root} therefore bounds the support-loss error throughout the phase, including sequences approaching either endpoint. The quantitative finite-dual estimate in Section 3 and (5.9b) give

$$\omega_n^{\mathrm{root}} = O\left(\frac{\log\log n}{\log n}\right),$$

with a phase-independent implied constant.

Substitution into (5.9a) yields the estimate used in the final theorem,

$$L_{S_4}(n, r_+) + (\log 2)r_+ = r_+\{\log 2 - D_4(\delta) + O(\omega_n^{\mathrm{root}})\} = r_+\{\log 2 - D_4(\delta) + o(1)\}. \tag{5.10}$$

The $o(1)$ in (5.10) is bounded uniformly by the sum of the finite-dual and target-displacement errors.

Finally, Lemma 5.2 gives $\log 2 - D_4(\delta) \ge \gamma_4 > 0$. Hence (5.10) implies $\Phi_n(r_+) > 0$ for all sufficiently large n, uniformly in the phase. The normalized slope error in (3.7) tends to zero, so the derivative of Φ_n is positive throughout the common root corridor after one phase-independent threshold. Since $\Phi_n(r_4^{\mathrm{co}}) = 0$, this proves

$$r_4^{\mathrm{co}} < r_+.$$

The exponent gain becomes a root separation by division by the slope. Both roots lie in (3.6), so the slope estimate holds throughout the interval between them. The mean-value theorem gives a point $\xi_n \in (r_4^{\mathrm{co}}, r_+)$ such that

$$\Phi_n(r_+) - \Phi_n(r_4^{\mathrm{co}}) = \Phi_n'(\xi_n)(r_+ - r_4^{\mathrm{co}}).$$

The corridor estimate (3.6) gives

$$r_+ = \frac{\log 2}{2}\frac{n}{\log n}\left\{1 + O\left(\frac{\log\log n}{\log n}\right)\right\},$$

and (3.7), uniformly at ξ_n, gives

$$\Phi_n'(\xi_n) = \frac{2}{\log 2}(\log n)^2 + O((\log n)(\log\log n)).$$

In the exact comparison (5.9a), the support-loss error is bounded by ω_n^{root}. We have just shown that it has the same relative order as the corridor and slope errors. Dividing the exponent difference by the slope therefore yields

$$r_+ - r_4^{\mathrm{co}} = \frac{(\log 2)^2}{4}A_4(\delta_n)\frac{n}{(\log n)^3} + O\left(\frac{n\log\log n}{(\log n)^4}\right). \tag{5.11}$$

The implied constant is independent of the phase: A_4 is bounded on $[0, 1]$, and all three relative errors above are uniform. This remainder concerns the analytic roots. Rounding and amplification introduce further errors in the probabilistic bound.

For the fixed-offset selection below, we use the following coarse consequence of (5.11): for all sufficiently large n,

$$r_+ - r_4^{\mathrm{co}} \geq \frac{(\log 2)^2 \gamma_4}{8} \frac{n}{(\log n)^3}. \tag{5.12}$$

5.3. **An exact integer profile above the first-moment root for marked partitions.** Choose an integer a fixed distance above the first-moment root for marked partitions:

$$k_{\mathrm{co}} = \lceil r_4^{\mathrm{co}} \rceil + 16. \tag{5.13}$$

Thus

$$16 \leq k_{\mathrm{co}} - r_4^{\mathrm{co}} < 17.$$

By (5.12), $k_{\mathrm{co}} < r_+$ for all sufficiently large n. Consequently the common corridor in Lemma 3.1 contains the entire interval from r_4^{co} to k_{co}. Every real point of this interval therefore has target deficit mean in the fixed compact interval $K_* \subset (2, 5)$ and is feasible for the real S_4 optimization. In particular this holds at the selected integer before the type counts are rounded. The convenient choice 16 leaves a fixed explicit margin in (5.19).

We now choose the *selected integer profile*. We round the optimizing type counts while preserving both finite-n conservation constraints.

At $k = k_{\mathrm{co}}$, let $p_i^{(n)}$ be the finite-n maximizer of L_{S_4}. Let $\lambda_{n,\mathrm{def}}$ denote the Lagrange multiplier for the deficit-mean constraint, with the sign convention that its contribution is $e^{\lambda_{n,\mathrm{def}} i}$. The Lagrange equations and the decomposition (3.1b) then give

$$p_i^{(n)} \propto d_{\alpha-i}^{-1} e^{\lambda_{n,\mathrm{def}} i} \propto e^{\widehat{\lambda}_n i + h_n(i)}, \qquad \widehat{\lambda}_n := B_n + \lambda_{n,\mathrm{def}}. \tag{5.14}$$

The finite-support tilt $\widehat{\lambda}_n$ absorbs the affine term $B_n i$. With a size-coordinate multiplier τ_n, the sign convention is $\lambda_{n,\mathrm{def}} = -\tau_n$. The target deficit mean is $\alpha - n/k_{\mathrm{co}}$. By uniqueness in Lemma 3.1, $\widehat{\lambda}_n$ is the finite-support tilt $\lambda_{n,S_4}(\alpha - n/k_{\mathrm{co}})$ introduced there. The corridor estimate and $k_{\mathrm{co}} - r_4^{\mathrm{co}} = O(1)$ give, uniformly in the phase,

$$\left| \alpha - \frac{n}{k_{\mathrm{co}}} - T_0 \right| = O\left(\frac{\log\log n}{\log n} \right) + O\left(\frac{(\log n)^2}{n} \right) \longrightarrow 0.$$

Lemma 3.1 therefore makes these proportions converge uniformly to (5.1), with $\widehat{\lambda}_n - \lambda_{S_4}(T_0) \to 0$ uniformly in the phase. The tilt bracket (5.6), proved in Appendix E, also shows directly that the four limiting weights differ by at most a fixed factor. Thus

$$\min_{2 \leq i \leq 5} p_i^{(n)} \geq c_p > 0 \tag{5.15}$$

uniformly in the phase.

For each sufficiently large n, fix a deterministic choice of preliminary integers $\widetilde{k}_i$ with $|\widetilde{k}_i - k_{\mathrm{co}} p_i^{(n)}| \leq 1$. The estimates below hold for every such choice. Define the signed constraint errors

$$e_0 = \sum_i \widetilde{k}_i - k_{\mathrm{co}}, \qquad e_1 = \sum_i i\widetilde{k}_i - (\alpha k_{\mathrm{co}} - n),$$

so $e_0, e_1 = O(1)$, and add

$$\Delta k_2 = e_1 - 3e_0, \qquad \Delta k_3 = 2e_0 - e_1. \tag{5.16}$$

Set $\Delta k_4 = \Delta k_5 = 0$, $k_i := \widetilde{k}_i + \Delta k_i$ for $2 \leq i \leq 5$, and $\mathbf{k} := (k_2, k_3, k_4, k_5)$.

The correction removes both constraint errors by changing only $O(1)$ counts. Appendix B checks positivity and shows that the loss in the counting exponent is $O(1/k_{\mathrm{co}})$.

The resulting exact integer vector satisfies

$$u_i = \alpha - i, \qquad k_i = \Theta(n/\log n) \quad (2 \leq i \leq 5). \tag{5.17}$$

Let $Z_{\mathbf{k}}^{\mathrm{sgn}}$ be the number of pairs consisting of a partition with exactly k_i classes of size u_i for $2 \le i \le 5$ and a declaration of each class as independent or complete, such that every declaration is realized by G_n. Direct counting gives the exact identity

$$\mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big] = 2^{k_{\mathrm{co}}} \frac{n!}{\prod_i (u_i!)^{k_i} k_i!} 2^{-\sum_i k_i \binom{u_i}{2}}, \tag{5.18}$$

Keeping every $u_i!$ exact and applying (1.2) only to $n!$ and the four profile-multiplicity factorials $k_i!$ then gives

$$\log \mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big] = L_{S_4}(n, k_{\mathrm{co}}) + \log 2k_{\mathrm{co}} + O(\log n) \ge c_Z (\log n)^2 \tag{5.19}$$

for a phase-independent $c_Z > 0$. Appendix B verifies that the fixed displacement above the root exceeds the rounding and Stirling losses. Every u_i tends to infinity, so the independent and complete declarations of one part are disjoint. Thus $Z_{\mathbf{k}}^{\mathrm{sgn}} > 0$ implies an actual cocoloring with k_{co} parts.

The construction gives one deterministic profile for the rest of the proof. We collect its properties here so that the second-moment estimates can refer to a single statement.

Proposition 5.3 (The selected four-size profile)**.** *For each sufficiently large n, choose k_{co} by (5.13) and $\mathbf{k} = (k_2, k_3, k_4, k_5)$ by the deterministic rounding and correction above. There are phase-independent constants $c_* > 0$, $c_Z > 0$ and an integer n_0 such that, for every integer $n \ge n_0$, the following hold.*

(i) *With $u_i = \alpha - i$ for $2 \le i \le 5$, the counts k_i are positive integers and satisfy exactly*

$$\sum_{i=2}^{5} k_i = k_{\mathrm{co}}, \qquad \sum_{i=2}^{5} u_i k_i = n.$$

Moreover, $k_i/k_{\mathrm{co}} \ge c_*$ for every i, and $k_{\mathrm{co}} \asymp n/\log n$.

(ii) *For $G_n \sim G(n, 1/2)$, the marked cocoloring count defined in (5.18) obeys*

$$\log \mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big] \ge c_Z (\log n)^2.$$

A positive count gives a cocoloring with exactly k_{co} parts.

(iii) *The ordinary and signed roots satisfy*

$$r_+(n) - r_4^{\mathrm{co}}(n) = \frac{(\log 2)^2}{4} A_4(\delta_n) \frac{n}{(\log n)^3} + O\left(\frac{n \log\log n}{(\log n)^4}\right).$$

For the integer chromatic lower location k_χ^- in (4.1),

$$\begin{aligned} k_\chi^- - k_{\mathrm{co}} &= \left[\frac{(\log 2)^2}{4} A_4(\delta) + o(1)\right] \frac{n}{(\log n)^3} \\ &\ge \left[\frac{(\log 2)^2 \gamma_4}{4} - o(1)\right] \frac{n}{(\log n)^3}. \end{aligned} \tag{5.20}$$

Here $\delta = \delta_n$. The choices, constants and error sequences are deterministic. The implied constants and vanishing errors are uniform over the induced phases, including integer sequences approaching either endpoint.

Proof. Equation (5.16) removes both rounding errors exactly. Appendix B gives $k_i \ge (c_p/2) k_{\mathrm{co}}$ after one phase-independent threshold; take $c_* = c_p/2$. The root corridor and (5.13) give the order of k_{co}. Equation (5.19) proves the first-moment bound, and the paragraph following it identifies a realized marked cocoloring with a cocoloring. The root separation is (5.11). Finally, (4.1) and (5.13) change its two roots by $O(\log n)$ and $O(1)$ respectively, giving (5.20). All these bounds were proved with phase-independent constants and errors. □

The gap and amplification scales.

It remains to turn this large first moment of the marked count into a lower bound on the probability of a cocoloring. The second-moment analysis in Sections 8–9 produces a deterministic exponent Λ_n such that

$$\frac{\mathbb{E}[(Z_{\mathbf{k}}^{\mathrm{sgn}})^2]}{(\mathbb{E} Z_{\mathbf{k}}^{\mathrm{sgn}})^2} \le e^{\Lambda_n}$$

and yields the scale hierarchy

$$k_{\mathrm{co}} \asymp \frac{n}{\log n}, \qquad r_+ - r_4^{\mathrm{co}} \asymp \frac{n}{(\log n)^3}, \qquad \Lambda_n = o\left(\frac{n}{(\log n)^4}\right).$$

Once the last estimate is available, its amplification cost satisfies

$$\frac{\sqrt{n\Lambda_n}}{\log n} = o\left(\frac{n}{(\log n)^3}\right),$$

which is smaller than the root separation.

6. Second moment of the marked cocoloring count

Fix the deterministic four-size profile from Proposition 5.3. Its marked cocoloring count has a large first moment. To bound the probability that it is positive, we must control the edge constraints shared by two partitions. We begin by counting their compatible signs. For partitions sharing exactly ℓ whole classes, Heckel (2025, Proposition 6) used the bound $2^{2k-\ell}$ to transfer selected ordinary-coloring estimates to cocolorings. We keep the full overlap matrix instead. Lemma 6.1 evaluates the sign count from the connected components formed by cells of size at least two; common whole classes are included in this count.

Overlap law and compatible signs. Temporarily label all slots within each of the four types. This multiplies the marked partition count by the deterministic factor $\prod_i k_i!$, leaving its normalized second moment unchanged. Let $\mathcal{I}_{\mathrm{row}}$ and $\mathcal{I}_{\mathrm{col}}$ be the row- and column-slot index sets. Choose two independent uniform ordered partitions of the profile, write them as $(V_a)_{a\in\mathcal{I}_{\mathrm{row}}}$ and $(W_b)_{b\in\mathcal{I}_{\mathrm{col}}}$, and set $s_a := |V_a|$ and $t_b := |W_b|$. Define

$$\mathbf{s} := (s_a)_{a\in\mathcal{I}_{\mathrm{row}}}, \qquad \mathbf{t} := (t_b)_{b\in\mathcal{I}_{\mathrm{col}}},$$

and let $\mathcal{R}(\mathbf{s},\mathbf{t})$ be the finite set of nonnegative integer matrices with row margins $\mathbf{s}$ and column margins $\mathbf{t}$. Put

$$r_{ab} = |V_a \cap W_b|. \tag{6.1}$$

The overlap matrix has the exact law

$$p(r) = \frac{\prod_a s_a! \prod_b t_b!}{n! \prod_{a,b} r_{ab}!}, \tag{6.2}$$

Both degree lists are the multiset given by (5.17). Equivalently, place s_a stubs at each row slot and t_b stubs at each column slot and take a uniform perfect matching of the two n-stub sets. For a fixed ordered column partition, the $\prod_b t_b!$ orderings of its column stubs form a constant-size fibre; hence this matching model induces exactly the law (6.2).

Let $H(r)$ be the simple bipartite graph whose edges are cells with $r_{ab} \geq 2$, omitting isolated slot vertices. Write $H := H(r)$ and

$$W = \sum_{a,b} \binom{r_{ab}}{2}, \qquad \beta(H) = |E(H)| - |V(H)| + c(H). \tag{6.3}$$

We use the standard convention $c(\varnothing) = 0$, so that $\beta(\varnothing) = 0$.

Figure 1 shows a small exact example of the thresholding operation and the two factors that it creates.

Lemma 6.1 (Exact sign sum)**.** *For every overlap matrix $r \in \mathcal{R}(\mathbf{s},\mathbf{t})$, the normalized local factor of the corresponding pair of partitions is*

$$A_\zeta(r) = 2^{W+c(H)-|V(H)|} = \left(\prod_{a,b} g(r_{ab})\right) 2^{\beta(H)}, \tag{6.4}$$

where

$$g(0) = g(1) = g(2) = 1, \qquad g(x) = 2^{\binom{x}{2}-1} \quad (x \geq 3). \tag{6.5}$$

Consequently

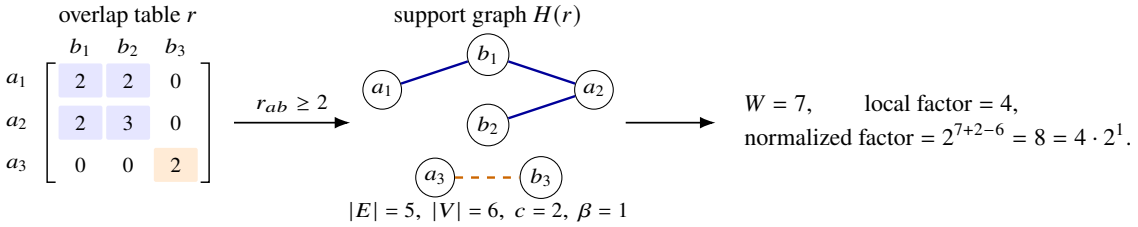


FIGURE 1. From an overlap table to the signed support factor. The shaded entries are precisely those with multiplicity at least two. Their support consists of one four-cycle (solid blue) and one isolated edge (dashed orange), so its binary cycle space has dimension one. Compatible signs are constant on each connected component; the colors and line styles mark the two components and carry no sign information. The displayed arithmetic verifies (6.4) for this example.

$$\frac{\mathbb{E}\big[(Z_{\mathbf{k}}^{\mathrm{sgn}})^2\big]}{\mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big]^2} = \sum_{r \in \mathcal{R}(\mathbf{s},\mathbf{t})} p(r) A_\zeta(r). \tag{6.6}$$

Here the sum ranges over the matrix set $\mathcal{R}(\mathbf{s},\mathbf{t})$ defined above.

Proof. Fix the two partitions. Let $B = \sum_a \binom{s_a}{2}$ be the number of internal edge bits prescribed by either one. A cell of size at least two contains an internal edge, so its row and column signs are compatible exactly when they agree. This forces a constant sign on each component of H.

The number of compatible pairs of sign assignments, including the free signs on slots isolated from H, is

$$2^{2k_{\mathrm{co}} - |V(H)| + c(H)}.$$

For each compatible pair, W edge bits are prescribed twice, so the probability of all constraints is $2^{-(2B-W)}$. Division by the square of the one-partition signed probability $(2^{k_{\mathrm{co}}} 2^{-B})^2$ gives

$$\frac{2^{2k_{\mathrm{co}} - |V(H)| + c(H)} 2^{-(2B-W)}}{(2^{k_{\mathrm{co}}} 2^{-B})^2} = 2^{W + c(H) - |V(H)|},$$

which is the first expression in (6.4).

For every support edge, split $\binom{r}{2}$ as $\big(\binom{r}{2} - 1\big) + 1$. Since $|E| - |V| + c = \beta$,

$$\prod_{a,b} g(r_{ab}) = 2^{W - |E(H)|},$$

and multiplying by $2^{\beta(H)}$ gives the second expression. Averaging over (6.2) proves (6.6). □

The identity

$$2^{\beta(H)} = \#\{F \subseteq E(H) : \deg_F(v) \text{ is even for every } v\} \tag{6.7}$$

expresses the cycle-space factor as a count of even edge sets. This description is useful in Section 9: restriction outside the exposed large-cell matching is injective, so we can bound the cycle-space contribution together with the local factors. The identity is the elementary fact that the even subgraphs form the binary cycle space of dimension $\beta(H)$.

Joint multiplicity bounds for overlap cells. The overlap cells are dependent. The next configuration-model estimate bounds a joint prescription of cell multiplicities by counting the stub pairs that realize it.

Lemma 6.2 (Joint prescribed-cell bound)**.** *Let $\mathcal{I}$ and $\mathcal{J}$ be finite row- and column-index sets, and let $(d_a)_{a \in \mathcal{I}}$ and $(d'_b)_{b \in \mathcal{J}}$ be nonnegative integer degree lists satisfying*

$$\sum_{a \in \mathcal{I}} d_a = \sum_{b \in \mathcal{J}} d'_b = m_0.$$

Choose a uniform perfect matching of the two m_0-stub sets, write $\mathbb{P}_{\mathrm{match}}$ for its law, and let r_{ab} be the number of matched pairs in cell (a,b). For a finite set $D \subseteq \mathcal{I} \times \mathcal{J}$, prescribe integer demands $x_{ab} \geq 1$ on D. Extend $x_{ab} = 0$ off D, and put

$$x = \sum_{D} x_{ab}, \quad D_a = \sum_{b} x_{ab}, \quad D'_b = \sum_{a} x_{ab}.$$

In the feasible case

$$x \leq m_0, \qquad D_a \leq d_a \ \ (\forall a), \qquad D'_b \leq d'_b \ \ (\forall b),$$

we have

$$\mathbb{P}_{\mathrm{match}}(r_{ab} \geq x_{ab} \text{ for every } (a,b) \in D) \leq \frac{\prod_a (d_a)_{D_a} \prod_b (d'_b)_{D'_b}}{(m_0)_x \prod_{(a,b)\in D} x_{ab}!}. \tag{6.8}$$

If any feasibility condition fails, the prescribed event is empty and is handled separately. In particular, in all cases with $m_0 > 0$, with $\mathrm{e} := \exp(1)$ and $\theta_{ab} = \mathrm{e}\, d_a d'_b / m_0$, its probability is at most

$$\prod_{(a,b)\in D} \frac{\theta_{ab}^{x_{ab}}}{x_{ab}!}. \tag{6.9}$$

Proof. If a feasibility condition fails, no matching realizes all demands. In the feasible case, assigning disjoint row stubs to the demanded cells gives

$$\prod_a \frac{(d_a)_{D_a}}{\prod_b x_{ab}!}$$

choices. Independently selecting the column stubs gives

$$\prod_b \frac{(d'_b)_{D'_b}}{\prod_a x_{ab}!}$$

choices. For each cell ab, the selected row and column stubs can be paired in $x_{ab}!$ ways. Hence the product over all cells cancels one of the two copies of $\prod_{(a,b)\in D} x_{ab}!$, and the total number of choices of prescribed stub pairs is

$$\frac{\prod_a (d_a)_{D_a} \prod_b (d'_b)_{D'_b}}{\prod_{(a,b)\in D} x_{ab}!}.$$

Any prescribed set of stub pairs occurs with probability $1/(m_0)_x$, because after its x prescribed pairs are exposed the remaining matching is uniform on the unused stubs. A union bound proves (6.8).

It remains to simplify the global denominator. For $1 \leq x \leq m_0$,

$$\binom{m_0}{x} = \prod_{j=0}^{x-1} \frac{m_0 - j}{x - j} \geq \left(\frac{m_0}{x}\right)^x, \qquad x! \geq \left(\frac{x}{\mathrm{e}}\right)^x.$$

The factorial inequality follows, for example, from $\sum_{j=1}^{x} \log j \geq \int_1^x \log t \, dt = x \log x - x + 1$. Consequently

$$(m_0)_x = x! \binom{m_0}{x} \geq (m_0/\mathrm{e})^x, \tag{6.10}$$

while $(m_0)_0 = 1 = (m_0/\mathrm{e})^0$ covers $x = 0$. Combining (6.10) with $(d)_r \leq d^r$ proves (6.9). Equation (6.8) retains the single global falling factorial before the coarser bound (6.9) is taken. □

7. Weights of common whole classes

Keep the profile of Proposition 5.3 fixed. First consider overlap cells that are whole classes common to both partitions. We choose equally sized row and column classes and require them to have the same vertex set. A *partial diagonal* is a specified collection of these common classes. We mark the chosen collection, so an overlap with several common classes appears once for each collection selected from them. This overcount is harmless for the upper bounds below: all marked weights are nonnegative.

For each vector of class-size counts, sum the exposed weights over all marked choices and denote the result by $D(\ell)$. These weights provide the reference case in the endpoint-table comparison of Section 8. We prove that their sum over all subprofiles is $1 + o(1)$ uniformly in the phase. The finite-n counting identities (7.1)–(7.6) let us add common blocks near the empty profile and missing blocks near the full profile. A rate estimate handles the intermediate masses.

For a selected common subprofile $\ell = (\ell_i)$, with $i \in \{2, 3, 4, 5\}$, put

$$\ell_\bullet = \sum_i \ell_i, \qquad m = \sum_i u_i \ell_i, \qquad b_\ell = \sum_i \binom{u_i}{2} \ell_i. \tag{7.1}$$

Here $0 \le \ell_i \le k_i$ coordinatewise. Since the full profile has $\sum_i u_i k_i = n$, this automatically gives $m \le n$; thus every factorial in the counting formula below has its ordinary combinatorial meaning.

Marked weights and recurrences. Summed over all choices of an m-vertex support in $[n]$, the first moment of the marked count of these partial partitions is

$$Y_\ell^{\mathrm{sgn}} = 2^{\ell_\bullet} \frac{n!}{(n-m)! \prod_i (u_i!)^{\ell_i} \ell_i!} 2^{-b_\ell}. \tag{7.2}$$

For each type i, the chosen row and column slots can be matched in $\binom{k_i}{\ell_i}^2 \ell_i!$ ways. For a fixed slot matching, labeling the common blocks contributes the corresponding factor $\ell_i!$ to Y_ℓ^{sgn}. These factors cancel type by type. In the normalized second moment a common block is then counted once, which gives the exact *marked* common-subprofile weight

$$D(\ell) = \frac{\prod_i \binom{k_i}{\ell_i}^2}{Y_\ell^{\mathrm{sgn}}}. \tag{7.3}$$

Only the incidence and local signed reward of the marked whole blocks enter $D(\ell)$. The other cell factors and the cycle-space factor remain to be estimated in Sections 8–9.

Different marked choices need not be disjoint: an overlap containing several common blocks contributes to each corresponding marked choice. The sums below are sums of these nonnegative marked weights, so no disjointness is used. In particular, $D(0) = 1$ and $D(\mathbf{k}) = 1/\mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big]$. Whenever $\ell_i < k_i$, one may add a common block of type i. Let e_i denote the ith coordinate vector. Since $n - m = \sum_j u_j (k_j - \ell_j) \ge u_i$, the denominator below is nonzero.

Moreover,

$$\frac{\binom{k_i}{\ell_i+1}^2}{\binom{k_i}{\ell_i}^2} = \frac{(k_i - \ell_i)^2}{(\ell_i + 1)^2},$$

while (7.2) gives

$$\frac{Y_{\ell+e_i}^{\mathrm{sgn}}}{Y_\ell^{\mathrm{sgn}}} = \frac{2(n-m)_{u_i}}{u_i!(\ell_i+1)2^{\binom{u_i}{2}}} = \frac{2\mu_{u_i}(n-m)}{\ell_i + 1}.$$

Dividing these two ratios gives

$$\frac{D(\ell + e_i)}{D(\ell)} = \frac{(k_i - \ell_i)^2}{2(\ell_i + 1)\mu_{u_i}(n-m)}. \tag{7.4}$$

At the opposite endpoint, put $h = \mathbf{k} - \ell$, $h_\bullet = \sum_i h_i$, $v = \sum_i u_i h_i$, and $b_h = \sum_i \binom{u_i}{2} h_i$. Define the finite first-moment factor

$$M_h(v) := 2^{h_\bullet} \frac{v!}{\prod_i (u_i!)^{h_i} h_i!} 2^{-b_h}.$$

Exact factorial cancellation gives

$$D(\mathbf{k} - h) = \frac{B(h)}{\mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big]}, \qquad B(h) = \left(\prod_i \binom{k_i}{h_i}\right) M_h(v), \tag{7.5}$$

and, whenever $h_i < k_i$,

$$\frac{\binom{k_i}{h_i+1}}{\binom{k_i}{h_i}} = \frac{k_i - h_i}{h_i + 1}, \qquad \frac{M_{h+e_i}(v + u_i)}{M_h(v)} = \frac{2\mu_{u_i}(v + u_i)}{h_i + 1}.$$

Therefore

$$\frac{B(h + e_i)}{B(h)} = \frac{2(k_i - h_i)\mu_{u_i}(v + u_i)}{(h_i + 1)^2}. \tag{7.6}$$

The finite-n identities (7.2)–(7.6) give two useful ways to grow a profile. Near the empty profile we add common blocks using (7.4). Near the full profile we add the missing blocks using (7.6). For the intervening selected masses we estimate the logarithm directly.

In Heckel's admissible-phase argument, the corresponding control of partial profiles enters through the tame-coloring estimates of Heckel and Panagiotou (2023, Lemmas 5.1–5.3, 6.3–6.5, and 7.20), as used in Heckel (2025, Proposition 5); their hypotheses are the source of the lower μ_α-window in that theorem. The next lemma supplies the uniform four-size estimate needed here, including both corners and every intermediate mass.

Lemma 7.1 (All common subprofiles)**.** *For the selected integer profile, uniformly over the complete phase,*

$$\sum_{0 \le \ell_i \le k_i} D(\ell) = 1 + o(1), \tag{7.7}$$

where the error is bounded by one deterministic sequence tending to zero.

The empty marked profile contributes one. We bound every other term by splitting the selected mass into three ranges. The forward recurrence (7.4) controls small selected mass; a negative leading rate controls the intermediate range; the reverse recurrence (7.6), together with the large complete first moment, controls nearly full selected mass.

The main point in the intermediate range is that the leading exponential rate is negative. Put $K = k_{\mathrm{co}}$, $p_i = k_i/K$, $y_i = \ell_i/K$, and $z_i = p_i - y_i$. Define

$$Y = \sum_i y_i, \qquad R = \sum_i z_i = 1 - Y, \qquad T = \sum_i i p_i, \qquad I_r = \sum_i i z_i.$$

Thus Y is the selected class fraction and R is the residual class fraction. The exact weight formula and the factorial estimates in Appendix F give, when the residual vertex fraction $\rho = (n - m)/n$ exceeds $1/32$,

$$\log D(\ell) \le K\alpha\Phi_T + CKY \log(\mathrm{e}/Y) + C \log n, \qquad \Phi_T = R \log R + \frac{\log 2}{2}(I_r - TR),$$

with a phase-independent constant C. All deficits lie in $\{2, 3, 4, 5\}$, so

$$I_r - TR \le (5 - T)R, \qquad I_r - TR \le (T - 2)(1 - R).$$

Multiplying the first bound by $1 - R$ and the second by R, then adding, gives $I_r - TR \le 3R(1 - R)$. Also $\rho > 1/32$ implies $R \ge 1/64$ eventually. For $R \le 3/4$, use $\log R \le 2(R - 1)/(R + 1)$ and $\log 2 < 7/10$:

$$\Phi_T \le -\left(\frac{2}{1 + R} - \frac{21}{20}\right) R(1 - R) \le -\frac{13}{8960} Y.$$

For $R \ge 3/4$, the selected profile has $T < 4$ eventually. The second deficit bound and $\log R \le R - 1$ then give

$$\Phi_T \le (\log 2 - R)Y \le -\frac{Y}{20}.$$

Thus $\Phi_T \le -Y/5000$ throughout this range.

To see that this dominates the error, put $\eta = \log\log n/(32\log n)$. For intermediate mass $m > \eta n$, one has $Y \ge \eta/2$ eventually. Dividing the error by $K\alpha Y$ bounds it by

$$\frac{C\log(2\mathrm{e}/\eta)}{\alpha} + \frac{2C\log n}{K\alpha\eta} = o(1),$$

uniformly in the phase. Hence, after one deterministic threshold, $\log D(\ell) \le -K\alpha\eta/20000$. There are at most $(K+1)^4$ subprofiles, and $K\alpha\eta/\log n \to \infty$, so their total intermediate contribution tends to zero. The forward recurrence covers $m \le \eta n$; the reverse recurrence covers $n - m \le n/32$. The appendix keeps the endpoint estimates and the full error accounting.

For the later second-moment assembly, write $\varepsilon_n^{\mathrm{empty}}$, $\varepsilon_n^{\mathrm{central}}$ and $\varepsilon_n^{\mathrm{full}}$ for the three deterministic, nonnegative error sequences in these estimates. The small-mass sum is bounded by the exponential of the first error, and the other two sums by their respective errors. Appendix F specifies these sequences, proves that they tend to zero uniformly in the phase, and gives the complete three-range argument.

8. Large overlap cells and endpoint comparisons

We now allow large overlap cells that stop short of whole classes. For a fixed set of such cells, sum all its labeled stub realizations and compare the resulting weight with full containment. The comparison has two costs: missing stubs inside a selected cell, and unequal sizes at its endpoints. We bound the first by summing deficits, then group the full-containment weights by endpoint table to bound the second.

The common-class weights from Section 7 have total mass $1 + o(1)$. Here we sum the weights of cells larger than half the class-size cap; Section 9 bounds the remaining local factors and the cycle-space factor.

8.1. Large cells and their labeled realizations.

Let $U := \alpha - 2$ be the largest class size in the fixed four-size profile and put

$$R_0 := \left\lfloor \frac{U}{2} \right\rfloor.$$

A cell is *large* if its multiplicity exceeds R_0. Two large cells cannot share a row or a column: otherwise that row or column would contain more than U stubs, because $2(R_0+1) > U$. Hence all large cells form a canonical bipartite matching, denoted by M.

The matching M records which class slots meet in large cells. A labeled stub realization also specifies the pairs inside each of those cells. The endpoints of M are disjoint, so selecting pairs in a cell of M does not use the stubs of another. This is why the realization counts multiply.

For $e \in M$, let s_e and t_e be its endpoint block sizes and set

$$m_e := \min\{s_e, t_e\}, \qquad d_e := |s_e - t_e|, \qquad h_e := m_e - j_e,$$

where j_e is the realized multiplicity. Thus m_e is the full-containment multiplicity, h_e is the deficit, and $j_e = m_e - h_e$. Since $j_e > U/2$ and $m_e \le U$,

$$2h_e < m_e. \tag{8.1}$$

Put $J := \sum_{e\in M} j_e$.

For fixed (M, j), let $\mathfrak{F}(M, j)$ be the finite set of labeled partial stub matchings containing exactly J pairs, all in cells of M, with exactly j_e pairs in each cell $e \in M$. Thus $\mathfrak{F}(M, j)$ records the labelled pairs, not just their cell counts.

A size-j partial matching between labeled blocks of sizes s and t can be chosen in

$$\binom{s}{j}\binom{t}{j} j! = \frac{(s)_j (t)_j}{j!}$$

ways: choose its two endpoint sets and then the bijection between them.

Proposition 8.1 (Aggregate weight of a fixed cell matching)**.** *Let M be a fixed matching support arising from two copies of the fixed feasible n-vertex profile. For a multiplicity vector $j = (j_e)_{e\in M}$ satisfying*

$0 \le j_e \le \min\{s_e, t_e\}$ *for every* $e \in M$*, one has* $J \le n$*, and the total configuration probability multiplied by the exposed local factor, summed over* $\mathfrak{F}(M, j)$*, is*

$$w(M, j) = \frac{\prod_{e \in M} (s_e)_{j_e} (t_e)_{j_e}}{(n)_J \prod_{e \in M} j_e!} \prod_{e \in M} g(j_e). \tag{8.2}$$

Proof. Because M is a matching, distinct selected cells use disjoint row and column stub sets. Therefore

$$|\mathfrak{F}(M, j)| = \prod_{e \in M} \frac{(s_e)_{j_e} (t_e)_{j_e}}{j_e!}.$$

A prescribed set of J stub pairs occurs with probability $(n)_J^{-1}$ in the uniform bipartite configuration matching. Multiplying this probability by the displayed cardinality and by the exposed reward $\prod_{e \in M} g(j_e)$ gives (8.2). □

A pair (M, j) is *feasible* when M is a matching between classes and

$$R_0 < j_e \le m_e \qquad (e \in M).$$

For a feasible pair (M, j), write

$$w_{\mathrm{hi}}(M, j) := w(M, j).$$

The same exposed weight appears in the exact conditional decomposition of Section 9.

Remark 8.2. The matching hypothesis is essential. If two selected cells shared a row or a column, their endpoint choices would compete for the same stubs and the one-cell realization counts would not multiply.

Figure 2 shows how the overlap table is split. To recover a full labeled matching, we must also retain the exposed stub pairs; this is done in Proposition 9.1. The last line of the figure separates the exposed weight w_{hi}, defined above, from the residual factor $\mathcal{A}$ of Section 9. Each factor will have its own bound.

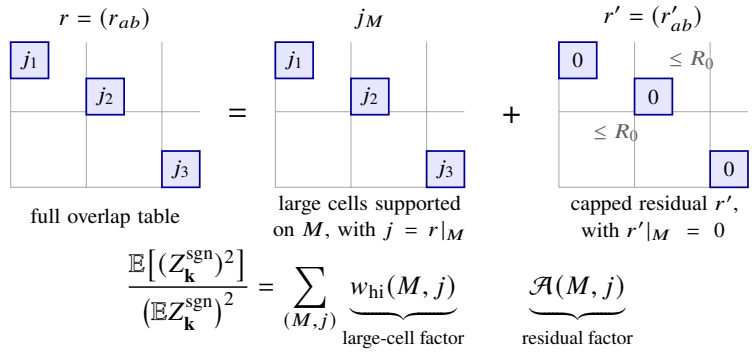


FIGURE 2. Canonical large-cell extraction. The highlighted cells $M = \{(a, b) : r_{ab} > R_0\}$ form a matching, with multiplicities $j = r|_M$. Extending j by zero gives $r = j_M + r'$, where r' vanishes on M and is capped by R_0 elsewhere.

Proposition 9.1 lifts this table decomposition to labelled configuration matchings by retaining the selected-pair set $S \in \mathfrak{F}(M, j)$ and the residual matching Π_{res}. This gives the normalized factorization shown in the figure. Blank cells have no displayed multiplicity. Outside M, the original cell equals its residual cell in r', is bounded by R_0, and has zero entry in j_M.

8.2. **Deficit ratios and the ambient denominator.** For endpoint sizes m and $m + d$, define

$$R_{m,d}(h) := \frac{\binom{m}{h}}{(d+1)(d+2)\cdots(d+h)} 2^{-hm+h(h+1)/2}, \qquad R_{m,d}(0) := 1. \tag{8.3}$$

The product in the denominator is empty when $h = 0$.

Lemma 8.3 (Exact one-cell deficit ratio)**.** *For endpoint sizes m and $m+d$, define the local numerator*

$$f_{m,d}(j) := \frac{(m)_j (m+d)_j}{j!}\, g(j).$$

For integers $m \geq 1$, $d \geq 0$, and $h \geq 0$ with $2h < m$,

$$\frac{f_{m,d}(m-h)}{f_{m,d}(m)} = R_{m,d}(h).$$

The function $f_{m,d}$ contains only the factors from one cell. The probability of all prescribed pairs contributes a single falling-factorial denominator, which we include in the aggregate comparison below.

Proof. The ratio of the partial matching counts is

$$\frac{(m)_{m-h}(m+d)_{m-h}}{(m-h)!}\,\frac{m!}{(m)_m (m+d)_m} = \frac{\binom{m}{h}}{(d+1)(d+2)\cdots(d+h)}.$$

The local signed factors satisfy

$$\frac{g(m-h)}{g(m)} = 2^{-hm+h(h+1)/2}.$$

Multiplying the two identities proves the claim. □

Write

$$\mathbf{m} := (m_e)_{e\in M}, \qquad w_{\mathrm{full}}(M) := w(M,(m_e)_{e\in M}), \qquad J_* := \sum_{e\in M} m_e, \qquad h_\bullet := J_* - J = \sum_{e\in M} h_e.$$

After the one-cell ratios have been extracted, the only factor that still couples different cells is the ambient denominator:

$$\frac{(n)_{J_*}}{(n)_J} = (n-J)_{h_\bullet} \leq n^{h_\bullet}. \tag{8.4}$$

Lemma 8.4 (Aggregate deficit comparison)**.** *For every $h = (h_e)_{e\in M} \in \mathbb{N}_0^M$ satisfying $2h_e < m_e$ for each $e \in M$,*

$$w(M, \mathbf{m} - h) \leq w_{\mathrm{full}}(M) \prod_{e\in M} n^{h_e} R_{m_e,d_e}(h_e). \tag{8.5}$$

Proof. Apply Lemma 8.3 in each selected cell, apply (8.4) once, and use $n^{h_\bullet} = \prod_{e\in M} n^{h_e}$. □

The local ratios now multiply, but the ambient matching probability still has one denominator for all exposed pairs. Summing the labeled realizations before estimating that denominator leaves exactly $(n-J)_{h_\bullet} \leq n^{h_\bullet}$, as in (8.4). Thus each missing pair is charged once in the aggregate comparison.

8.3. **Summing all positive deficits.** For a selected cell e, enlarge the exact positive-deficit range to

$$A_e := \{h \in \mathbb{N} : h \geq 1,\ 2h < m_e\}.$$

This range may include deficits that no longer give a large cell. Adding their nonnegative weights only increases the upper bound.

For any nonnegative $u_e(h)$, set $u_e(0) := 1$. Expanding the finite product gives the optional-choice identity

$$\sum_{\omega \in \prod_{e\in M}(\{0\}\cup A_e)} \prod_{e\in M} u_e(\omega_e) = \prod_{e\in M}\left(1 + \sum_{h\in A_e} u_e(h)\right),$$

where each ω_e is either the zero-deficit choice, of weight one, or one positive deficit $h \in A_e$.

For $2h < m$, integrality gives $2h+1 \leq m$. Hence

$$\frac{h+1}{2} \leq \frac{m+1}{4} \leq m - \left\lfloor \frac{3m-1}{4} \right\rfloor.$$

Multiplying by h yields

$$h\left\lfloor \frac{3m-1}{4} \right\rfloor \leq hm - \frac{h(h+1)}{2}. \tag{8.6}$$

Together with $\binom{m}{h} \le m^h$ and (8.3), this gives

$$n^h R_{m,d}(h) \le \left(\frac{nm}{2^{\lfloor (3m-1)/4 \rfloor}}\right)^h. \tag{8.7}$$

Endpoint tables have four rows and four columns. To label them from zero, we now use $i \in \{0,1,2,3\}$ for the deficit $\nu = i+2$ from Sections 5–7. Assume $\alpha > 8$, a condition included in the common eventuality threshold of Subsection 10.1, and define

$$\kappa_i := k_{i+2}, \qquad u_i := \alpha - 2 - i, \qquad 0 \le i \le 3.$$

Thus an endpoint of type i has the same size and multiplicity as a class of deficit $i+2$ in Section 5. From here through Section 9, the symbol u_i uses this shifted index; it is the earlier size u_{i+2}, with that earlier index running from 2 to 5. The counts κ_i are therefore the same selected counts, not a new profile. For $i, j \in \{0,1,2,3\}$ let $m_{ij} := \min\{u_i, u_j\}$ and put

$$\rho_{ij} := \frac{nm_{ij}}{2^{\lfloor (3m_{ij}-1)/4 \rfloor}}, \qquad \rho_\Sigma := \sum_{i,j=0}^{3} \rho_{ij}.$$

Every local base in (8.7) is at most ρ_Σ.

Lemma 8.5 (Coarse local deficit sum)**.** *Suppose $\rho_\Sigma \le 1$. Then, for every selected cell e,*

$$\sum_{h \in A_e} n^h R_{m_e,d_e}(h) \le (\alpha+1)\rho_\Sigma.$$

Proof. If e has endpoint type (i,j), then (8.7) bounds its term of deficit h by ρ_{ij}^h. Since $h \ge 1$ and $\rho_{ij} \le \rho_\Sigma \le 1$, this is at most ρ_Σ. Moreover, $|A_e| \le m_e \le \alpha + 1$. Summation proves the bound. □

Proposition 8.6 (Fixed-support all-deficit bound)**.** *Write $\mathbf{m} = (m_e)_{e \in M}$. If $\rho_\Sigma \le 1$, then*

$$\sum_{\substack{h_e \in \mathbb{N}_0 \ (e \in M) \\ m_e - h_e > R_0 \text{ for every } e}} w(M, \mathbf{m} - h) \le w_{\mathrm{full}}(M)\,(1 + (\alpha+1)\rho_\Sigma)^{|M|}. \tag{8.8}$$

Proof. Use Lemma 8.4, enlarge the deficit range to A_e, apply the optional-choice identity above, and then use Lemma 8.5. □

Define $K := \sum_{i=0}^{3} \kappa_i = k_{\mathrm{co}}$, the number of row-class slots. Projection onto the row endpoint is injective on a block matching, and therefore

$$|M| \le K. \tag{8.9}$$

8.4. **Regrouping by endpoint table.** After summing the deficits, only the full-containment reference weights remain. Their dependence on the endpoint sizes is captured by the numbers of matched type pairs. We group supports with the same endpoint table, then compare their row and column margins with the partial-diagonal weights.

For a block matching M, let

$$\ell_{ij}(M) := \#\{e \in M : e \text{ joins row type } i \text{ to column type } j\}, \qquad L(M) := (\ell_{ij}(M))_{0 \le i,j \le 3}.$$

With this deficit-indexed type convention, for $L = (\ell_{ij}) \in \mathbb{N}_0^{4\times 4}$ write

$$r = (r_i)_{i=0}^3, \qquad r_i = \sum_j \ell_{ij}, \qquad c = (c_j)_{j=0}^3, \qquad c_j = \sum_i \ell_{ij}.$$

Call L *feasible* if $r_i \le \kappa_i$ and $c_j \le \kappa_j$ for every i, j. Such a table is realized by choosing the indicated row and column block slots and pairing them cell by cell; conversely, every block matching has a feasible table. For a feasible endpoint table L, define

$$W(L) := \sum_{M:\, L(M) = L} w_{\mathrm{full}}(M).$$

Equivalently, $W(L)$ is the sum of the full-containment terms over all block-slot matching supports with table L.

Partitioning the finite set of block matchings by $L(M)$ gives

$$\sum_{M} w_{\mathrm{full}}(M) = \sum_{L \text{ feasible}} W(L). \tag{8.10}$$

Similarly, for every nonnegative function F of the endpoint table,

$$\sum_{M} w_{\mathrm{full}}(M) F(L(M)) = \sum_{L \text{ feasible}} W(L) F(L). \tag{8.11}$$

Whenever $\rho_\Sigma \le 1$, combining Proposition 8.6, (8.9), and (8.10) yields, after enlarging the nonnegative sum from canonical large-cell matchings to all block matchings,

$$\sum_{(M,j) \text{ feasible}} w_{\mathrm{hi}} \le \left(\sum_{L \text{ feasible}} W(L) \right) (1 + (\alpha+1)\rho_\Sigma)^K . \tag{8.12}$$

The finite argument through (8.12) applies to any fixed endpoint alphabet whose selected cells form a matching. The four consecutive sizes enter through the numerical deficit bound and the endpoint-table comparison below.

8.5. **Comparison with the common-class weights.** Assume throughout that $\alpha > 8$. For equal endpoint sizes, a full-containment cell is a whole class common to the two partitions. This gives the reference case estimated in Section 7. We compare a general full-containment table with the partial-diagonal weights of its two margins, including their row and column multinomial factors. Their geometric mean gives the reference value, and unequal endpoint sizes contribute a small correction.

Recall that $u_i = a - i$ and $\kappa_i = k_{i+2}$ for $0 \le i \le 3$, where $a = \alpha - 2$ and the corresponding deficit label is $i + 2$. The dictionary with Section 7 is

$$\ell^{\mathrm{pd}}_{i+2} = v_i, \qquad k_{i+2} = \kappa_i, \qquad \alpha - (i+2) = u_i.$$

Here v records a common-class margin; ℓ^{pd} denotes the same common-subprofile counts in the earlier deficit indexing. We will use this dictionary to identify the one-sided weights below with (7.3). For a full-containment table $L = (\ell_{ij})$, set

$$r = (r_i)_{i=0}^3, \qquad r_i = \sum_j \ell_{ij}, \qquad c = (c_j)_{j=0}^3, \qquad c_j = \sum_i \ell_{ij},$$

$$x_{ij} = \min\{u_i, u_j\}, \qquad J(L) = \sum_{i,j} x_{ij} \ell_{ij}.$$

The number of block matchings with table L is

$$\frac{\prod_i (\kappa_i)_{r_i} \prod_j (\kappa_j)_{c_j}}{\prod_{i,j} \ell_{ij}!}.$$

Indeed, row and column slot allocations initially produce two copies of the cell-factorial denominator, while the bijections between the selected slots inside each type cell contribute one copy back. Thus the exact reference weight is

$$W(L) = \frac{\prod_i (\kappa_i)_{r_i} \prod_j (\kappa_j)_{c_j}}{\prod_{i,j} \ell_{ij}!} \frac{1}{(n)_{J(L)}} \prod_{i,j} \left[\frac{(u_i)_{x_{ij}} (u_j)_{x_{ij}}}{x_{ij}!} g(x_{ij}) \right]^{\ell_{ij}}. \tag{8.13}$$

Let

$$\mathcal{V} := \{ v = (v_i)_{i=0}^3 \in \mathbb{N}_0^4 : 0 \le v_i \le \kappa_i \text{ for every } i \}.$$

For $v \in \mathcal{V}$, define

$$D(v) = \frac{\prod_i (\kappa_i)_{v_i}^2}{\prod_i v_i!} \frac{\prod_i [u_i! g(u_i)]^{v_i}}{(n)_{m(v)}}, \qquad m(v) = \sum_i u_i v_i. \tag{8.14}$$

With $\ell^{\mathrm{pd}}_{i+2} = v_i$, $D(v)$ is the common-subprofile weight in (7.3). This follows from $\kappa_i = k_{i+2}$, $u_i = \alpha - (i+2)$, and $g(u_i) = 2^{\binom{u_i}{2}-1}$ for $\alpha > 8$. Below, $D(r)$ and $D(c)$ evaluate (8.14) on the margins in $0 \le v_i \le \kappa_i$.

For the endpoint pair (i, j), put

$$s_{ij} := \min\{u_i, u_j\}, \qquad t_{ij} := \max\{u_i, u_j\}, \qquad d_{ij} := t_{ij} - s_{ij},$$

and define

$$Q_{ij} := (n+1)^{d_{ij}/2} \frac{\sqrt{(t_{ij})_{d_{ij}}}}{d_{ij}!} 2^{-\{d_{ij}s_{ij} + \binom{d_{ij}}{2}\}/2}.$$

Thus $Q_{ii} = 1$. Since $t_{ij} \le a$, $s_{ij} \ge a - 3$, and $d_{ij} \le 3$,

$$Q_{ij} \le \frac{(\tau_n^{\mathrm{end}})^{d_{ij}}}{d_{ij}!}, \qquad \tau_n^{\mathrm{end}} = 2^{3/2}\sqrt{\frac{(n+1)a}{2^a}} = O\left(\frac{(\log n)^{3/2}}{\sqrt{n}}\right). \tag{8.15}$$

Let

$$A_L = \frac{\prod_i r_i!}{\prod_{i,j} \ell_{ij}!}, \qquad C_L = \frac{\prod_j c_j!}{\prod_{i,j} \ell_{ij}!}, \qquad Q^L = \prod_{i,j} Q_{ij}^{\ell_{ij}}.$$

Lemma 8.7 (Endpoint-table product comparison)**.** *Assume $\alpha > 8$ and that $(\kappa_i, u_i)_{i=0}^3$ is the fixed feasible profile, so $\sum_{i=0}^3 \kappa_i u_i = n$. Then every feasible endpoint table L satisfies*

$$W(L)^2 \le \big(D(r)A_L Q^L\big)\big(D(c)C_L Q^L\big). \tag{8.16}$$

Proof. Fix the table and divide $W(L)^2$ by $(D(r)A_L)(D(c)C_L)$. The profile multiplicities and table factorials cancel. We are left with a local ratio for each endpoint pair and a ratio of ambient falling factorials; the latter will compensate the powers in the former.

Fix one endpoint pair and abbreviate $s = s_{ij}$, $t = t_{ij}$, $d = d_{ij} = t - s$, and $x_{ij} = s$. Its local atom in (8.13) is

$$\frac{(s)_s (t)_s}{s!} g(s) = \frac{t!}{d!} g(s).$$

Consequently

$$\frac{\left[\frac{(s)_s (t)_s}{s!} g(s)\right]^2}{s!g(s)\, t!g(t)} = \frac{(t)_d}{(d!)^2} 2^{-ds - \binom{d}{2}} = \frac{Q_{ij}^2}{(n+1)^d}.$$

Under the hypothesis, $s \ge \alpha - 5 > 3$, so in particular $s \ge 2$. Hence we may use $g(t)/g(s) = 2^{ds + \binom{d}{2}}$; the identity is also valid for $d = 0$.

Moreover, $m(r), m(c) \ge J(L)$ and

$$m(r) + m(c) - 2J(L) = \sum_{i,j} d_{ij}\ell_{ij}.$$

Therefore

$$\frac{(n)_{m(r)}(n)_{m(c)}}{(n)_{J(L)}^2} \le (n+1)^{m(r)+m(c)-2J(L)}.$$

Combining this inequality with the local ratio cell by cell gives

$$\begin{aligned} \frac{W(L)^2}{(D(r)A_L)(D(c)C_L)} &\le (n+1)^{m(r)+m(c)-2J(L)} \prod_{i,j} \left(\frac{Q_{ij}^2}{(n+1)^{d_{ij}}}\right)^{\ell_{ij}} \\ &= \prod_{i,j} Q_{ij}^{2\ell_{ij}} = (Q^L)^2, \end{aligned}$$

where the equality uses $m(r) + m(c) - 2J(L) = \sum_{i,j} d_{ij}\ell_{ij}$. This proves (8.16). □

Proposition 8.8 (Endpoint-table sum)**.** *There is an absolute constant $C_{\mathrm{end}} > 0$ such that, for all sufficiently large n, uniformly in the phase,*

$$\sum_{L \text{ feasible}} W(L) \le (1 + 3\tau_n^{\mathrm{end}})^K \sum_{v \in \mathcal{V}} D(v) \le \exp\left\{C_{\mathrm{end}}\sqrt{n \log n}\right\}. \tag{8.17}$$

All sums in this proposition are finite.

Proof. For all sufficiently large n, one has $\tau_n^{\mathrm{end}} \le 1$ by (8.15); hence $Q_{ij} \le \tau_n^{\mathrm{end}}$ whenever $i \ne j$. By the arithmetic–geometric mean inequality and (8.16),

$$2W(L) \le D(r)A_L Q^L + D(c)C_L Q^L.$$

Sum the row term first, keeping its margin r fixed. We can drop the column constraints and sum over all nonnegative integer tables with row margin r, since $A_L Q^L \ge 0$. The multinomial theorem then gives

$$\sum_{\substack{L \text{ feasible} \\ \mathrm{row}(L)=r}} A_L Q^L \le \sum_{\substack{L \in \mathbb{N}_0^{4\times 4} \\ \mathrm{row}(L)=r}} A_L Q^L = \prod_i \left(\sum_j Q_{ij} \right)^{r_i} \le (1 + 3\tau_n^{\mathrm{end}})^K.$$

Multiply by $D(r)$ and sum over the row margin r. The result is at most $(1 + 3\tau_n^{\mathrm{end}})^K \sum_{v \in \mathcal{V}} D(v)$, with v a dummy subprofile variable. Exchanging rows and columns gives the same bound for the second term. Thus the endpoint-table choices cost only the displayed correction factor. The remaining sum over margins is the common-subprofile sum of Section 7: $\sum_{v \in \mathcal{V}} D(v) = 1 + o(1)$ uniformly in the phase, with all common-class overlaps included. Finally, $K = \Theta(n/\log n)$ and (8.15) imply $\tau_n^{\mathrm{end}} K = O(\sqrt{n \log n})$. □

8.6. **The large-cell bound across the phase.** The phase satisfies

$$\alpha = \left(\frac{2}{\log 2} + o(1) \right) \log n.$$

In particular, for all sufficiently large n,

$$\frac{5}{2} \log n \le \alpha. \tag{8.18}$$

Every endpoint minimum satisfies $m_{ij} \ge \alpha - 5$, and elementary floor arithmetic gives

$$3\alpha - 19 \le 4 \left\lfloor \frac{3m_{ij} - 1}{4} \right\rfloor.$$

Using $\log 2 > 2/3$ and (8.18), we obtain

$$\frac{5}{4} \log n - \frac{19}{6} \le (\log 2) \left\lfloor \frac{3m_{ij} - 1}{4} \right\rfloor. \tag{8.19}$$

Hence

$$2^{\lfloor (3m_{ij}-1)/4 \rfloor} \ge e^{-19/6} n^{5/4}.$$

Since every $m_{ij} = O(\log n)$,

$$\rho_\Sigma = O\left(\frac{\log n}{n^{1/4}} \right) \longrightarrow 0. \tag{8.20}$$

Thus $\rho_\Sigma \le 1$ for all sufficiently large n.

Let Σ_n^{hi} denote the sum of exposed weights:

$$\Sigma_n^{\mathrm{hi}} := \sum_{(M,j) \text{ feasible}} w_{\mathrm{hi}}(M, j),$$

where the sum is over the feasible pairs (M, j) defined above.

Proposition 8.9 (Bound for the exposed weights)**.** *There is a phase-independent deterministic sequence $\varepsilon_n^{\mathrm{hi}} \ge 0$ with $\varepsilon_n^{\mathrm{hi}} \to 0$ such that, for all sufficiently large n,*

$$\Sigma_n^{\mathrm{hi}} \le \exp\left\{ \varepsilon_n^{\mathrm{hi}} \frac{n}{(\log n)^4} \right\}. \tag{8.21}$$

Proof. The deficit factor in (8.12) satisfies

$$\log\left(1 + (\alpha + 1)\rho_\Sigma\right)^K \le K(\alpha + 1)\rho_\Sigma = O(n^{3/4} \log n).$$

Proposition 8.8 contributes $O(\sqrt{n \log n})$ to the logarithm. Both quantities are $o(n/(\log n)^4)$. Combining these estimates with (8.12) proves the proposition. □

9. Residual factors and the second-moment bound

We keep the selected profile of Proposition 5.3 and the endpoint indexing from Section 8. The exposed matching M has pairs of class slots as its edges. The labelled stub pairs chosen inside those cells are denoted by S.

Fix one feasible pair (M, j) and one such labeled realization. Condition only on its stub pairs being present. The unused row stubs are then matched uniformly to the unused column stubs. The indicator of allowed completions is zero if a completion adds a pair in an exposed cell or exceeds the residual cap elsewhere, and one otherwise. Both the local factors and the cycle-space factor depend on this same residual matching. We multiply their product by the indicator of allowed completions before taking its expectation. A bound uniform over every feasible pair (M, j) can then be taken outside the exposed-weight sum from Section 8.

Throughout this section, write $Z := Z_{\mathbf{k}}^{\mathrm{sgn}}$ for the selected marked cocoloring count.

9.1. **Exact decomposition after exposing the large cells.** Fix a feasible pair (M, j) and recall $J = \sum_{e\in M} j_e$. Put

$$m_0 := n - J.$$

Here feasibility means that M is a block matching and $R_0 < j_e \le \min\{s_e, t_e\}$ for every $e \in M$. Each $S \in \mathfrak{F}(M, j)$ is a labelled partial matching containing exactly J pairs. Let $(d_a)_a$ and $(d'_b)_b$ be the residual row and column degrees. Their two sums equal m_0, and every degree is at most the phase cap U. Explicitly,

$$d_a = s_a - \sum_{b:(a,b)\in M} j_{ab}, \qquad d'_b = t_b - \sum_{a:(a,b)\in M} j_{ab}.$$

The unexposed pairs form a uniform bipartite configuration matching with these residual degrees; when $m_0 = 0$, this means the unique empty matching. Write $\mathbb{P}_{\mathrm{res}}$ and $\mathbb{E}_{\mathrm{res}}$ for its law and expectation. Let r'_{ab} be its cell counts and let H_{res} be the simple support graph of cells with $r'_{ab} \ge 2$. The event $\mathcal{E}(M, j)$ specifies which residual matchings complete the chosen large cells. Outside the exposed matching it imposes the cap $r'_{ab} \le \lfloor U/2 \rfloor$; in a cell of M, no further pair is allowed. The exposed data already include every pair in each large cell, so an additional pair would change its multiplicity. Thus, with $R_0 = \lfloor U/2 \rfloor$,

$$\mathcal{E}(M, j) = \{r'_{ab} = 0 \ ((a, b) \in M), \quad r'_{ab} \le R_0 \ ((a, b) \notin M)\}.$$

More explicitly, let Π be the uniform full configuration matching. For every $S \in \mathfrak{F}(M, j)$ and every event $\mathcal{B}$ of the residual matching,

$$\mathbb{P}\{S \subseteq \Pi, \ \Pi \setminus S \in \mathcal{B}\} = \frac{1}{(n)_J}\, \mathbb{P}_{\mathrm{res}}(\mathcal{B}). \tag{9.1}$$

Different labeled realizations leave the same residual degree lists. A blockwise bijection between the unused row stubs, together with one between the unused column stubs, identifies their residual matching spaces and preserves every cell count. The residual expectation therefore depends on M and the degree lists alone. Define

$$\mathcal{A}(M, j) := \mathbb{E}_{\mathrm{res}}\left[\left(\prod_{a,b} g(r'_{ab})\right) 2^{\beta(M\cup H_{\mathrm{res}})} \mathbf{1}_{\mathcal{E}(M,j)}\right]. \tag{9.2}$$

Only the residual configuration matching is random in this expectation.

Thus $w_{\mathrm{hi}}(M, j)$ sums the exposed realizations in $\mathfrak{F}(M, j)$, while $\mathbf{1}_{\mathcal{E}(M,j)}$ selects exactly their completions with capped residual cells and no further pairs in exposed cells. The next proposition puts the two counts together.

Proposition 9.1 (Canonical exact overlap decomposition)**.** *Let n be finite and let the fixed selected profile be feasible, with phase cap $U \ge 2$. Then*

$$\frac{\mathbb{E}Z^2}{(\mathbb{E}Z)^2} = \sum_{(M,j)} w_{\mathrm{hi}}(M, j)\, \mathcal{A}(M, j), \tag{9.3}$$

where the sum ranges over the finite family of feasible pairs (M, j).

Proof. Let Π be a full bipartite configuration matching and let $r_{ab}(\Pi)$ be its cell counts. Define

$$M(\Pi) := \{(a,b) : r_{ab}(\Pi) > R_0\}, \qquad j_{ab}(\Pi) := r_{ab}(\Pi) \quad ((a,b) \in M(\Pi)).$$

Since $2(R_0+1) > U$, the degree cap makes $M(\Pi)$ a matching. Moreover, $j_{ab} > \lfloor U/2 \rfloor$ and $U \geq 2$ imply $j_{ab} \geq 2$, so every cell of $M(\Pi)$ is an edge of the support graph from Section 6. Let $S(\Pi)$ be the set of all stub pairs of Π lying in cells of $M(\Pi)$, and put $\Pi_{\mathrm{res}} := \Pi \setminus S(\Pi)$. Then

$$\Pi \longmapsto \big(M(\Pi), j(\Pi), S(\Pi), \Pi_{\mathrm{res}}\big)$$

takes values in the feasible pairs (M,j), their labeled realizations, and residual matchings satisfying $\mathcal{E}(M,j)$. It is injective because $\Pi = S(\Pi) \cup \Pi_{\mathrm{res}}$.

Conversely, take a feasible (M,j), an $S \in \mathfrak{F}(M,j)$, and a perfect matching of the unused stubs satisfying $\mathcal{E}(M,j)$. Their disjoint union is a full matching. The condition $r'_{ab} = 0$ on M leaves multiplicity $j_e > R_0$ in each $e \in M$, and the residual cap leaves every cell outside M with multiplicity at most R_0. Thus its canonical large-cell data are exactly (M,j), and this construction is the inverse of the displayed map.

For each fixed S, the exposed-residual factorization (9.1) contributes $(n)_J^{-1}$ and the residual configuration law. Summing over $S \in \mathfrak{F}(M,j)$ and including the exposed local factors gives $w_{\mathrm{hi}}(M,j)$ by Proposition 8.1. By Lemma 6.1, the unexposed local factors and the remaining sign factor are exactly the integrand defining $\mathcal{A}(M,j)$, with support factor $2^{\beta(M \cup H_{\mathrm{res}})}$. The bijection therefore yields (9.3) and counts every full matching once. □

The case $m_0 = 0$ is complete without introducing cell activities. The residual matching is empty, every residual local factor is one, and $H_{\mathrm{res}} = \varnothing$. Since M is a matching, $\beta(M) = 0$, so

$$\mathcal{A}(M,j) = 1 \qquad (m_0 = 0). \tag{9.4}$$

In all subsequent subsections we assume $m_0 > 0$. In view of Proposition 8.9, it remains to bound $\mathcal{A}(M,j)$ uniformly in the positive-residual case.

9.2. **Threshold expansion and cell activities.** For every row–column pair (a,b), put

$$\theta_{ab} := \frac{\mathrm{e}\, d_a d'_b}{m_0}. \tag{9.5}$$

The intensity is defined on cells of M as well as outside M. The activities below collect the terms used to bound the expansion of the residual product; they are set to zero on M. Let

$$\Delta_x := g(x) - g(x-1), \qquad R := \left\lfloor \frac{U}{2} \right\rfloor,$$

so $\Delta_x \geq 0$: indeed $\Delta_3 = 3$, while $g(x)/g(x-1) = 2^{x-1} > 1$ for $x \geq 4$. For $(a,b) \notin M$, define

$$\lambda_{ab} := \sum_{x=3}^{R} \Delta_x \frac{\theta_{ab}^x}{x!}, \qquad q_{ab} := \frac{\theta_{ab}^2}{2} + \lambda_{ab}. \tag{9.6}$$

On M, set $\lambda_{ab} = q_{ab} = 0$.

For a finite array of nonnegative integer demands $x = (x_{ab})$ supported outside M, apply Lemma 6.2 to the residual degree lists and $\mathrm{supp}(x) := \{(a,b) : x_{ab} > 0\}$. It gives, in all cases,

$$\mathbb{P}_{\mathrm{res}}(r'_{ab} \geq x_{ab} \text{ for all } a,b) \leq \prod_{(a,b) \in \mathrm{supp}(x)} \frac{\theta_{ab}^{x_{ab}}}{x_{ab}!}. \tag{9.7}$$

Because this estimate is joint, it applies simultaneously to cells sharing a row or a column; no independence between cells is asserted.

For $0 \leq r \leq R$,

$$g(r) = 1 + \sum_{x=3}^{R} \Delta_x \mathbf{1}_{\{r \geq x\}}.$$

Using the row- and column-slot sets defined in Section 6, put $E_0 := (\mathcal{I}_{\mathrm{row}} \times \mathcal{I}_{\mathrm{col}}) \setminus M$, and let $\mathfrak{E}(M)$ be the family of even subsets of $M \cup E_0$. For an edge $e = (a, b)$, write $r'_e := r'_{ab}$, $\theta_e := \theta_{ab}$, $\lambda_e := \lambda_{ab}$, and $q_e := q_{ab}$. For $F \in \mathfrak{E}(M)$, define

$$\Phi_F(r') := \prod_{e \in F \setminus M} \left(\mathbf{1}_{\{r'_e \geq 2\}} + \sum_{x=3}^{R} \Delta_x \mathbf{1}_{\{r'_e \geq x\}} \right) \times \prod_{e \in E_0 \setminus F} \left(1 + \sum_{x=3}^{R} \Delta_x \mathbf{1}_{\{r'_e \geq x\}} \right).$$

An edge e of the even set outside the exposed matching must occur in the residual support, so the first product includes its threshold-two indicator. Cells outside both the even set and the exposed matching contribute only their local factors. On the event $\mathcal{E}(M, j)$, expansion of the cycle-space cardinality and of all local factors gives the exact identity

$$\left(\prod_{a,b} g(r'_{ab}) \right) 2^{\beta(M \cup H_{\mathrm{res}})} = \sum_{F \in \mathfrak{E}(M)} \Phi_F(r').$$

This identity uses both parts of $\mathcal{E}(M, j)$: off-M cells are capped at R, and r' vanishes on M. We next multiply by $\mathbf{1}_{\mathcal{E}(M,j)}$ and bound the expectation term by term. Every term of the truncated expansion is nonnegative, so dropping $\mathbf{1}_{\mathcal{E}(M,j)}$ gives an upper bound for its expectation. On $\mathcal{E}(M, j)^c$ we use only this expansion, without identifying it with the uncapped reward.

Lemma 9.2 (Fixed even-set expansion)**.** *For every $F \in \mathfrak{E}(M)$,*

$$\mathbb{E}_{\mathrm{res}}\left[\Phi_F(r') \mathbf{1}_{\mathcal{E}(M,j)} \right] \leq \prod_{e \in F \setminus M} q_e \prod_{e \in E_0 \setminus F} (1 + \lambda_e).$$

Proof. Expand Φ_F into its nonnegative threshold monomials. In a cell of $F \setminus M$, a monomial chooses either the threshold-two term or one higher increment; it never chooses both. In a cell outside $F \cup M$, it chooses nothing or one higher increment. Apply (9.7) once to the complete demand of each monomial. The threshold-two contribution is $\theta_e^2/2$, the higher increments sum to λ_e, and the empty choice contributes one. Every monomial is nonnegative and $\mathbf{1}_{\mathcal{E}(M,j)} \leq 1$, so omitting the indicator of $\mathcal{E}(M, j)$ can only increase the expectation. □

Summing over F and then inserting the missing factors $1 + \lambda_e \geq 1$ gives

$$\mathcal{A}(M, j) \leq \left(\prod_{e \in E_0} (1 + \lambda_e) \right) \sum_{F \in \mathfrak{E}(M)} \prod_{e \in F \setminus M} q_e. \tag{9.8}$$

9.3. **Restricting even edge sets outside the exposed matching.** The sum over even sets remains. Their edges inside the exposed matching carry no activity, so we restrict each set to its edges outside that matching. If the restriction determines the original set, we may enlarge the sum to all subsets of the remaining edges and obtain a product bound. Here is the finite inequality we need.

Lemma 9.3 (Even-set restriction bound)**.** *Let H be a finite graph, let $M \subseteq E(H)$ be a matching, and let $q_e \geq 0$ for $e \in E(H) \setminus M$. Then*

$$\sum_{\substack{F \subseteq E(H) \\ F \text{ even}}} \prod_{e \in F \setminus M} q_e \leq \prod_{e \in E(H) \setminus M} (1 + q_e).$$

Proof. Two even edge sets with the same restriction outside M have an even symmetric difference contained in M. A nonempty subset of a matching has a vertex of degree one, so the difference must be empty. Restriction outside M is therefore injective. Enlarge the resulting sum to all subsets of $E(H) \setminus M$ and expand the product. □

Apply the lemma to the complete bipartite graph on the row and column slots, with edge set $M \cup E_0$. It gives

$$\sum_{F \in \mathfrak{E}(M)} \prod_{e \in F \setminus M} q_e \leq \prod_{e \in E_0} (1 + q_e). \tag{9.9}$$

Since $0 \le \lambda_e \le q_e$, equations (9.8) and (9.9), together with $1 + x \le e^x$, imply

$$\mathcal{A}(M, j) \le \exp\left(2 \sum_{e \in E_0} q_e\right). \tag{9.10}$$

Thus both the local factors and the cycle-space cardinality are controlled by one total residual activity.

9.4. **Large residual mass: a quadratic activity bound.** It remains to sum the residual activities. We first treat residual mass large enough to make every intensity small relative to the class-size cap. In this range the higher-reward contribution to each cell activity is bounded by a constant times the square of its intensity.

Assume

$$2^U \le m_0^3. \tag{9.11}$$

Then $m_0 \ge 2^{U/3}$ and, by the degree cap,

$$\theta_{ab} \le \mathrm{e} U^2 2^{-U/3}.$$

Lemma 9.4 (Quadratic activity bound)**.** *There is an absolute constant C_0 such that, under* (9.11)*,*

$$q_{ab} \le C_0 \theta_{ab}^2 \tag{9.12}$$

for every residual cell outside M.

The quadratic activity estimate is proved in Appendix G; summing it over the residual degree lists gives the bound below.

Because θ_{ab} was defined for every pair, the cell intensities satisfy the exact identity

$$\sum_{a,b} \theta_{ab}^2 = \frac{\mathrm{e}^2}{m_0^2} \left(\sum_a d_a^2\right) \left(\sum_b (d'_b)^2\right).$$

Since every degree is at most U and both degree sums equal m_0,

$$\sum_a d_a^2 \le U m_0, \qquad \sum_b (d'_b)^2 \le U m_0.$$

On M the activities are zero, while outside M the preceding lemma applies. Therefore

$$\sum_{e \in E_0} q_e \le C_1 U^2 \qquad \text{for one absolute } C_1 > 0. \tag{9.13}$$

Together with (9.10), this proves

$$\mathcal{A}(M, j) \le \exp(2 C_1 U^2) \qquad \text{if } 2^U \le m_0^3. \tag{9.14}$$

9.5. **Small residual mass: a pointwise bound.** If the residual mass is smaller, the quadratic activity estimate need not apply. There are then few residual pairs, and a pointwise bound suffices. Suppose $2^U > m_0^3$. Then

$$m_0 < 2^{U/3} \le 2^{\lceil U/3 \rceil}. \tag{9.15}$$

The pointwise argument uses only the canonical cap and the matching property, so it also applies outside this regime. Retaining the cap gives a sharper estimate. On $\mathcal{E}(M, j)$, write

$$H := H_{\mathrm{res}}, \qquad W' := \sum_{a,b} \binom{r'_{ab}}{2}.$$

Every support edge is a cell of multiplicity at least two, so the local factor, including $g(2) = 1$, is exactly

$$\prod_{a,b} g(r'_{ab}) = 2^{W' - |E(H)|}.$$

Restriction of even edge sets outside the matching M is injective by the argument following Lemma 9.3. No residual pair lies in an exposed cell, so $E(H) \cap M = \varnothing$; hence

$$2^{\beta(M \cup H)} \le 2^{|E(H)|}.$$

The edge-cardinality factors cancel. Since $U \ge 2$, the cap $R_0 = \lfloor U/2 \rfloor$ is at least one, and for $0 \le r \le R_0$,

$$\binom{r}{2} \le \frac{R_0 - 1}{2} r.$$

Summing over cells, using that their residual multiplicities sum to m_0, gives

$$\left(\prod_{a,b} g(r'_{ab})\right) 2^{\beta(M\cup H)} \le 2^{W'} \le 2^{(R_0-1)m_0/2}.$$

The residual factor integrand vanishes outside $\mathcal{E}(M, j)$. Taking expectations proves

$$\mathcal{A}(M, j) \le 2^{(R_0-1)m_0/2} \qquad \text{for every } m_0 \ge 0. \tag{9.16}$$

We use this bound when the positive residual mass is small. It also covers zero residual mass, empty residual support, and an empty canonical event.

We combine the zero-residual case with the two positive-residual regimes $2^U \le m_0^3$ and $2^U > m_0^3$.

Proposition 9.5 (Uniform residual factor)**.** *For the fixed feasible selected profile with $U \ge 2$, there is an absolute constant $C_{\mathrm{att}} > 0$ such that every feasible pair (M, j) satisfies*

$$\mathcal{A}(M, j) \le \begin{cases} 1, & m_0 = 0, \\ \exp(C_{\mathrm{att}} U^2), & m_0 > 0 \text{ and } 2^U \le m_0^3, \\ 2^{(R_0-1)m_0/2}, & m_0 > 0 \text{ and } 2^U > m_0^3. \end{cases}$$

Proof. The three cases are exhaustive. Their bounds are (9.4), (9.14), and (9.16), with $C_{\mathrm{att}} := 2C_1$. □

9.6. **The global second-moment exponent.** We combine the exposed-weight and residual-factor bounds in the exact decomposition. Their logarithms add; we check below that the resulting exponent is small enough for amplification.

Combine the three deterministic partial-diagonal errors introduced in Section 7 and estimated in Appendix F by setting

$$\varepsilon_n^{\mathrm{pd}} := e^{\varepsilon_n^{\mathrm{empty}}} - 1 + \varepsilon_n^{\mathrm{central}} + \varepsilon_n^{\mathrm{full}}. \tag{9.17}$$

Then $\varepsilon_n^{\mathrm{pd}} \to 0$ uniformly in the phase, and the disjoint three-range assembly gives

$$\sum_{v\in\mathcal{V}} D(v) \le 1 + \varepsilon_n^{\mathrm{pd}}. \tag{9.18}$$

All sums in this subsection are finite and range over the domains defined in Sections 7–9.

Let

$$\tau_n^{\mathrm{end}} := 2^{3/2}\sqrt{\frac{(n+1)a}{2^a}}, \qquad a = \alpha - 2.$$

The notation agrees with (8.15). For all sufficiently large n, $\tau_n^{\mathrm{end}} \le 1$. Since $Q_{ii} = 1$ and each row has only three off-diagonal positions, (8.15) gives the explicit row bound

$$\sum_j Q_{ij} \le 1 + 3\tau_n^{\mathrm{end}}. \tag{9.19}$$

Indeed, for $i \ne j$ one has $Q_{ij} \le (\tau_n^{\mathrm{end}})^{d_{ij}}/d_{ij}! \le \tau_n^{\mathrm{end}}$.

Combining Proposition 8.8 with (9.18) gives

$$\sum_{L \text{ feasible}} W(L) \le (1 + 3\tau_n^{\mathrm{end}})^K (1 + \varepsilon_n^{\mathrm{pd}}). \tag{9.20}$$

Combining this with the single aggregate deficit reduction (8.12), define

$$\begin{aligned} \Gamma_n^{\mathrm{skel}} := {} & K \log(1 + (\alpha+1)\rho_\Sigma) \\ & + K \log\left(1 + 3\tau_n^{\mathrm{end}}\right) + \log\left(1 + \varepsilon_n^{\mathrm{pd}}\right). \end{aligned} \tag{9.21}$$

The sum over (M, j) in (9.3) is precisely the Σ_n^{hi} defined in Section 8. For all sufficiently large n,

$$\Sigma_n^{\mathrm{hi}} \le e^{\Gamma_n^{\mathrm{skel}}}. \tag{9.22}$$

In Γ_n^{skel}, the first term pays for the missing pairs, the second for unequal endpoint sizes, and the third for the common-class reference mass. We estimate each term on the required scale.

The phase estimates from Section 8 imply

$$K\log(1+(\alpha+1)\rho_\Sigma) \le K(\alpha+1)\rho_\Sigma = O(n^{3/4}\log n),$$
$$K\log\left(1+3\tau_n^{\mathrm{end}}\right) \le 3K\tau_n^{\mathrm{end}} = O(\sqrt{n\log n}),$$
$$\log(1+\varepsilon_n^{\mathrm{pd}}) = o(1).$$

Consequently the deterministic normalized exposed-weight error

$$\varepsilon_n^{\mathrm{skel}} := \frac{(\log n)^4}{n}\Gamma_n^{\mathrm{skel}} \longrightarrow 0. \tag{9.23}$$

Let C_{att} be the constant from Proposition 9.5. Define

$$\Gamma_n^{\mathrm{att}} := \max\left\{C_{\mathrm{att}}U^2, \frac{\log 2}{2}(R_0-1)2^{U/3}\right\}. \tag{9.24}$$

Proposition 9.5 implies the uniform bound

$$\mathcal{A}(M,j) \le e^{\Gamma_n^{\mathrm{att}}} \tag{9.25}$$

for every feasible pair (M,j). For large residual mass we use (9.14). For small positive residual mass, $m_0 < 2^{U/3}$ and

$$2^{(R_0-1)m_0/2} = \exp\left(\frac{\log 2}{2}(R_0-1)m_0\right) \le \exp\left(\frac{\log 2}{2}(R_0-1)2^{U/3}\right).$$

The zero-residual case contributes one.

Since

$$U = O(\log n), \qquad 2^U = \Theta\left(\frac{n^2}{(\log n)^2}\right),$$

one has

$$\Gamma_n^{\mathrm{att}} = O\left((\log n)^2 + n^{2/3}(\log n)^{1/3}\right) = o\left(\frac{n}{(\log n)^4}\right).$$

Thus the deterministic normalized residual factor error

$$\varepsilon_n^{\mathrm{att}} := \frac{(\log n)^4}{n}\Gamma_n^{\mathrm{att}} \longrightarrow 0. \tag{9.26}$$

Proposition 9.6 (Explicit normalized second-moment bound)**.** *For the marked cocoloring count $Z = Z_{\mathbf{k}}^{\mathrm{sgn}}$ from Proposition 5.3, define*

$$\Lambda_n := \Gamma_n^{\mathrm{skel}} + \Gamma_n^{\mathrm{att}} = \left(\varepsilon_n^{\mathrm{skel}} + \varepsilon_n^{\mathrm{att}}\right)\frac{n}{(\log n)^4}. \tag{9.27}$$

Then

$$\Lambda_n = o\left(\frac{n}{(\log n)^4}\right)$$

and for all sufficiently large n,

$$1 \le \frac{\mathbb{E}Z^2}{(\mathbb{E}Z)^2} \le e^{\Lambda_n}. \tag{9.28}$$

Proof. Insert (9.25) into the exact decomposition (9.3). Because all summands are nonnegative, the uniform residual factor may be taken outside the finite sum over (M,j). Equation (9.22) then gives

$$\frac{\mathbb{E}Z^2}{(\mathbb{E}Z)^2} \le e^{\Gamma_n^{\mathrm{att}}}\Sigma_n^{\mathrm{hi}} \le e^{\Gamma_n^{\mathrm{att}}+\Gamma_n^{\mathrm{skel}}} = e^{\Lambda_n}.$$

The lower bound follows from $\mathrm{Var}(Z) \ge 0$. Equations (9.23) and (9.26) give the stated order of Λ_n. □

The second-moment estimate is now ready for Paley–Zygmund. Γ_n^{skel} accounts for partial diagonals, deficits, and endpoint tables; Γ_n^{att} accounts for residual local factors and the cycle-space factor. Their sum is the exponent used in the amplification step.

10. Amplifying a cocoloring probability bound

Proposition 9.6 and (1.4) give the probability bound

$$\mathbb{P}\big(Z_{\mathbf{k}}^{\mathrm{sgn}} > 0\big) \geq e^{-\Lambda_n}. \tag{10.1}$$

The event $Z_{\mathbf{k}}^{\mathrm{sgn}} > 0$ yields a marked cocoloring and hence a cocoloring with k_{co} classes, so the probability bound implies

$$\mathbb{P}(\zeta(G_n) \leq k_{\mathrm{co}}) \geq e^{-\Lambda_n}. \tag{10.2}$$

The probability bound need not tend to one. Following Heckel (2025, Theorem 1), we use vertex-exposure concentration as in Scott (2017, Theorem 1) to find a large induced subgraph with the desired cocoloring. We then color the leftover vertices. Their set depends on the graph, so a bound for a fixed vertex set would not suffice. The first lemma controls all sets simultaneously; the second bounds the mean excess above the desired class count. Markov's inequality then gives an error of $o(n/(\log n)^3)$, smaller than the root separation. Appendix C gives a stronger explicit tail bound.

Coloring every leftover vertex set.

Lemma 10.1 (Simultaneous leftover coloring)**.** *There is an absolute C_0 such that, with probability $1 - o(1)$, every $S \subseteq [n]$ satisfies*

$$\chi(G_n[S]) \leq C_0 \frac{|S|}{\log n} + n^{1/3}. \tag{10.3}$$

Proof. We find independent sets in every sufficiently large vertex set by working in the complement H of G_n. Put $u_0 = \lceil n^{1/4} \rceil$. For a fixed u_0-set, (1.5) gives probability $\exp(-\Omega(u_0^2))$ that its H-edge density is below $1/4$. There are at most $\binom{n}{u_0}$ such sets, and for an absolute $c > 0$,

$$\binom{n}{u_0} e^{-cu_0^2} \leq \exp\{u_0 \log(\mathrm{e}n/u_0) - cu_0^2\} = o(1).$$

Thus a union bound shows that, with probability $1 - o(1)$, every u_0-set has density at least $1/4$. If S has size $s \geq u_0$, double-counting each edge of $H[S]$ over the u_0-subsets containing it gives

$$\frac{1}{\binom{s}{u_0}} \sum_{\substack{T \subseteq S \\ |T| = u_0}} \frac{e_H(T)}{\binom{u_0}{2}} = \frac{e_H(S)}{\binom{s}{2}}.$$

Therefore every larger set also has density at least $1/4$.

Let S_0 be any vertex set of size at least $n^{1/3}$. While $|S_t| \geq u_0$, choose $v_t \in S_t$ of maximum degree in $H[S_t]$, define $S_{t+1} := N_H(v_t) \cap S_t$, and put $s_t := |S_t|$. Density at least $1/4$ implies $|N_H(v_t) \cap S_t| \geq (s_t - 1)/4$. Hence

$$s_{t+1} \geq (s_t - 1)/4, \qquad s_t \geq 4^{-t} s_0 - 1/3. \tag{10.3a}$$

The sets are nested and $S_{t+1} \subseteq N_H(v_t)$. Thus, whenever $j > t$,

$$v_j \in S_j \subseteq S_{t+1} \subseteq N_H(v_t).$$

Every later choice is therefore adjacent in H to every earlier one, so the chosen vertices form a clique in H. Since $s_0 \geq \lceil n^{1/3} \rceil$, for every $0 \leq t \leq \lfloor \log n/(13 \log 4) \rfloor$, equation (10.3a) gives

$$s_t \geq n^{1/3 - 1/13} - \frac{1}{3} = n^{10/39} - \frac{1}{3} \geq n^{1/4}$$

for all sufficiently large n. Thus the procedure constructs at least $c \log n$ vertices for an absolute $c > 0$, and these vertices form an independent set in G_n.

On the same density event, take any S and repeatedly remove one of these independent sets. Assign it a new color, and stop when fewer than $n^{1/3}$ vertices remain. Color those vertices singly. Each removed set has at least $c \log n$ vertices, so this uses at most

$$\frac{|S|}{c \log n} + n^{1/3}.$$

After enlarging the absolute constant C_0, this is (10.3), simultaneously for every S. □

Mean excess above the class-count bound. We bound the mean number of vertices left outside a cocolorable set, then color those vertices on the simultaneous density event.

Lemma 10.2 (Mean bound for the cochromatic number)**.** *For all sufficiently large n, every deterministic integer $k \geq 0$ and every deterministic finite real $\Lambda \geq 0$ satisfying*

$$\mathbb{P}(\zeta(G_n) \leq k) \geq e^{-\Lambda}$$

obey the following bound. Put $u_0 := \lceil n^{1/4} \rceil$ and

$$\varepsilon_n^{\mathrm{dens}} := \min\left\{1, \binom{n}{u_0} \exp\left(-\frac{\binom{u_0}{2}}{16}\right)\right\}.$$

With $x_+ := \max\{x, 0\}$ and the same absolute C_0 as in Lemma 10.1,

$$\mathbb{E}[(\zeta(G_n) - k)_+] \leq \frac{C_0}{\log n}\sqrt{\frac{(n-1)\Lambda}{2}} + n^{1/3} + n\varepsilon_n^{\mathrm{dens}}.$$

The eventual threshold is independent of k, Λ, and the phase.

Proof. Let $\mathcal{D}_n$ be the density event used in the proof of Lemma 10.1: every u_0-set has edge density at least $1/4$ in the complement of G_n. By (1.5) and a union bound,

$$\mathbb{P}\big(\mathcal{D}_n^c\big) \leq \varepsilon_n^{\mathrm{dens}}.$$

On $\mathcal{D}_n$, (10.3) holds for every vertex set.

To apply concentration, we measure the largest vertex set that can be cocolored with at most the given number of classes. For $W \subseteq [n]$, let $\zeta(W) = \zeta(G_n[W])$, with $\zeta(\varnothing) = 0$, and define

$$S_k = \max\{|W| : W \subseteq [n],\ \zeta(W) \leq k\}. \tag{10.4}$$

Expose the graph in $n-1$ independent vertex blocks, placing in block v the edges from its vertex to earlier vertices. Two outcomes that differ only in block v agree on all edges away from that vertex. Take a maximizing set in the first outcome. If it contains the changed vertex, delete it and restrict the cocoloring to the remaining set; otherwise keep the set. Either way, it is feasible in the second outcome and has lost at most one vertex. Interchanging the outcomes gives the reverse comparison. Thus S_k changes by at most one, and (1.3) applies (McDiarmid, 1998, Theorem 3.1).

Since $S_k = n$ exactly when $\zeta(G_n) \leq k$, the assumed probability bound and the one-sided upper-tail form of (1.3) give

$$e^{-\Lambda} \leq \mathbb{P}(S_k - \mathbb{E}[S_k] \geq n - \mathbb{E}[S_k]) \leq \exp\left\{-\frac{2(n - \mathbb{E}[S_k])^2}{n-1}\right\}.$$

Taking logarithms bounds the mean number of uncovered vertices:

$$n - \mathbb{E}[S_k] \leq \sqrt{(n-1)\Lambda/2}. \tag{10.5}$$

Choose a maximizing set W and put $V_{\mathrm{left}} = [n] \setminus W$. Combining a cocoloring of $G_n[W]$ using at most k parts with an ordinary coloring of $G_n[V_{\mathrm{left}}]$ gives

$$\zeta(G_n) \leq k + \chi(G_n[V_{\mathrm{left}}]). \tag{10.6}$$

On $\mathcal{D}_n$, the leftover bound applies to every vertex set, including this one chosen from the graph. Combining its coloring with the cocoloring of W gives

$$(\zeta(G_n) - k)_+ \leq \frac{C_0}{\log n}(n - S_k) + n^{1/3}.$$

Outside $\mathcal{D}_n$, the positive excess is at most n. Bound the contribution of this exceptional event by its probability times the number of vertices. On the density event, removing the indicator from the nonnegative

uncovered mass $n - S_k$ can only increase its expectation. Taking expectations therefore proves the bound. No independence between the maximizing set and $\mathcal{D}_n$ is needed. Finally,

$$\log \binom{n}{u_0} = O(n^{1/4} \log n), \qquad \binom{u_0}{2} = (1 + o(1))\sqrt{n}/2,$$

so $\varepsilon_n^{\mathrm{dens}} \le e^{-c\sqrt{n}}$ eventually for one absolute $c > 0$. □

For the probability bound (10.2), the explicit second-moment bound satisfies $\Lambda_n = o(n/(\log n)^4)$. Consequently,

$$\mathbb{E}[(\zeta(G_n) - k_{\mathrm{co}})_+] = o\left(\frac{n}{(\log n)^3}\right).$$

This also gives $\mathbb{E}[\zeta(G_n)] \le k_{\mathrm{co}} + o(n/(\log n)^3)$, uniformly along the full integer sequence.

Applying the probability bound. The mean bound gives the high-probability conclusion by Markov's inequality. Put $s_n = n/(\log n)^3$ and let

$$b_n := \frac{C_0}{\log n}\sqrt{\frac{(n-1)\Lambda_n}{2}} + n^{1/3} + n\varepsilon_n^{\mathrm{dens}}, \qquad a_n := \sqrt{b_n s_n}. \tag{10.7}$$

Here $b_n > 0$ is a deterministic upper bound for $\mathbb{E}(\zeta(G_n) - k_{\mathrm{co}})_+$. Since $\Lambda_n = o(n/(\log n)^4)$ and $\varepsilon_n^{\mathrm{dens}} \le e^{-c\sqrt{n}}$,

$$b_n = o(s_n), \qquad a_n = o(s_n). \tag{10.8}$$

Markov's inequality now gives

$$\mathbb{P}\big(\zeta(G_n) > k_{\mathrm{co}} + a_n\big) \le \frac{b_n}{a_n} = \sqrt{\frac{b_n}{s_n}} \longrightarrow 0. \tag{10.9}$$

All these sequences are deterministic, and the bounds hold uniformly in the phase. The stronger tail estimate is Lemma C.1 in Appendix C.

10.1. **Uniformity across the phase.** The root estimates are uniform on K_*, and the errors above have deterministic bounds independent of the phase. The residual estimates cover all residual masses, and the leftover-coloring event holds for every vertex set simultaneously. All eventual conditions, including $\alpha > 8$, $U \ge 2$, $\rho_\Sigma \le 1$ and $\tau_n^{\mathrm{end}} \le 1$, therefore hold after one deterministic threshold. The constants in the mean bound are independent of its parameters k and Λ. Appendix C records the corresponding uniformity for the explicit tail bound.

11. The phase-resolved gap and a uniform constant

We return to (A)–(C) of the introduction. The first-moment roots give the deterministic separation; the chromatic lower bound and the amplified cocoloring upper bound hold with high probability. Intersecting their events gives the phase-resolved gap. We then use continuity over the closed phase interval to obtain the uniform constant.

Proof of the main theorem. Take n beyond the common deterministic threshold of Subsection 10.1. Let $r_+(n)$ be the ordinary first-moment root and let $r_4^{\mathrm{co}}(n)$ be the four-size first-moment root for marked partitions. Proposition 5.3 gives

$$r_+(n) - r_4^{\mathrm{co}}(n) = \left[\frac{(\log 2)^2}{4} A_4(\delta_n) + o(1)\right] \frac{n}{(\log n)^3}. \tag{11.1}$$

The selected profile is placed a fixed distance above its first-moment root:

$$k_{\mathrm{co}} = \lceil r_4^{\mathrm{co}}(n) \rceil + 16.$$

Proposition 5.3 changes the four class-size counts by $O(1)$ at this fixed k_{co}. With the same notation $i \in \{2, 3, 4, 5\}$ and $u_i = \alpha - i$, it preserves both

$$\sum_i k_i = k_{\mathrm{co}} \qquad \text{and} \qquad \sum_i u_i k_i = n.$$

Thus there is no additional correction to the total number of classes. Hence

$$r_+(n) - k_{\mathrm{co}} = \left[\frac{(\log 2)^2}{4} A_4(\delta_n) + o(1)\right] \frac{n}{(\log n)^3}. \tag{11.2}$$

Indeed, (5.13) gives $r_+(n) - k_{\mathrm{co}} = (r_+(n) - r_4^{\mathrm{co}}) - 16 + O(1)$. The fixed displacement and ceiling are absorbed by the displayed $o(1)$.

Section 4 constructs a deterministic integer k_χ^- such that

$$\mathbb{P}\big(\chi(G_n) > k_\chi^-\big) \longrightarrow 1$$

and $k_\chi^- = r_+(n) + O(\log n)$. Combining this with (11.2) gives

$$k_\chi^- - k_{\mathrm{co}} = \left[\frac{(\log 2)^2}{4} A_4(\delta_n) + o(1)\right] \frac{n}{(\log n)^3}. \tag{11.3}$$

The amplification result (10.9) gives a deterministic sequence $a_n = o(n/(\log n)^3)$ such that

$$\mathbb{P}\big(\zeta(G_n) \le k_{\mathrm{co}} + a_n\big) \longrightarrow 1.$$

On this event and the chromatic lower event,

$$\chi(G_n) - \zeta(G_n) > k_\chi^- - k_{\mathrm{co}} - a_n.$$

The deterministic locations used in this comparison are collected in Figure 3.

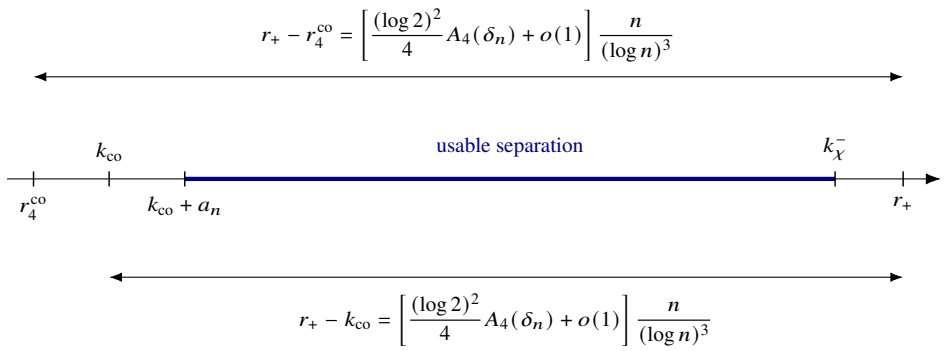


FIGURE 3. First-moment roots and the class counts used after rounding and amplification. For all sufficiently large n, $r_4^{\mathrm{co}} < k_{\mathrm{co}} < k_{\mathrm{co}} + a_n < k_\chi^- < r_+$. The thick segment is the remaining gap between the class counts. Positions are schematic and not to scale.

To make the final error explicit, put

$$s_n := \frac{n}{(\log n)^3}, \qquad \theta_n := \frac{k_\chi^- - k_{\mathrm{co}}}{s_n} - \frac{(\log 2)^2}{4} A_4(\delta_n), \qquad \eta_n := \frac{a_n}{s_n}.$$

By (11.3) and (10.9), $\theta_n \to 0$ and $\eta_n \to 0$. Hence the nonnegative deterministic error

$$\varepsilon_n^{\mathrm{gap}} := |\theta_n| + \eta_n$$

tends to zero. On the event

$$\{\chi(G_n) > k_\chi^-\} \cap \{\zeta(G_n) \le k_{\mathrm{co}} + a_n\},$$

$$\chi(G_n) - \zeta(G_n) > k_\chi^- - k_{\mathrm{co}} - a_n \ge \left[\frac{(\log 2)^2}{4} A_4(\delta_n) - \varepsilon_n^{\mathrm{gap}}\right] s_n.$$

By (4.6) and (10.9), a union bound shows that this intersection has probability $1 - o(1)$. Consequently,

$$\mathbb{P}\left(\chi(G_n) - \zeta(G_n) \ge \left[\frac{(\log 2)^2}{4} A_4(\delta_n) - \varepsilon_n^{\mathrm{gap}}\right] \frac{n}{(\log n)^3}\right) \longrightarrow 1. \tag{11.4}$$

On this tail, set $\varepsilon_n := \varepsilon_n^{\mathrm{gap}}$ and define the finitely many preceding terms arbitrarily as nonnegative reals. Thus $(\varepsilon_n)_{n\ge 2}$ is deterministic, nonnegative, and tends to zero.

The uniform coefficient.

Put $q := \log 2$ and $\gamma_4 := \log(200/153) > 0$. Lemma 5.1 locates the smallest phase gain at an endpoint, while Lemma 5.2 bounds both endpoint values. Hence

$$m_4 := \min\{A_4(0), A_4(1)\} = \min_{0\le\delta\le 1} A_4(\delta) > \gamma_4.$$

Since $\varepsilon_n^{\mathrm{gap}} \to 0$, eventually

$$\varepsilon_n^{\mathrm{gap}} < \frac{q^2}{4}(m_4 - \gamma_4).$$

Therefore the coefficient in (11.4) is at least $q^2\gamma_4/4$ for all sufficiently large n. It follows that

$$\mathbb{P}\left(\chi(G_n) - \zeta(G_n) \ge \frac{(\log 2)^2}{4}\log\left(\frac{200}{153}\right)\frac{n}{(\log n)^3}\right) \longrightarrow 1.$$

This proves the main theorem. Since $n/(\log n)^3 \to \infty$, the chromatic–cochromatic difference tends to infinity with high probability along the full sequence of integers. We have proved the lower-bound direction of the conjectured $n/(\log n)^3$ scale; a matching upper bound and the optimal constant remain open. □

AI assistance disclosure. OpenAI GPT-5.6 Sol, accessed through ChatGPT and Codex, was used to probe the written proof for counterexamples, audit symbolic consistency, and generate candidate Lean code. Codex also assisted with English-language editing, organization, and supplementary proof arguments for the finite-dual approximation and residual overlap bounds. Aristotle, Harmonic's automated theorem-proving system (Achim et al., 2025), performed automated Lean proof search and a clean-environment replay. Every formal output incorporated into the released development was reviewed and rebuilt under the pinned Lean toolchain. Machine-checked claims rest on the Lean kernel, and manuscript-only claims on the written proof. Samuil Petkov directed the work, selected the claims and arguments, reviewed the manuscript and formalization, and assumes responsibility for the accuracy, integrity, citations, and final submission.

Appendix A. Proof of the uniform root estimates

Proof of Lemma 3.1. Put $q := \log 2$, $L := \log n$, and $\ell := \log\log n$. Dividing (3.2) by $k = n/s$ and writing $p_i = k_i/k$ gives

$$\frac{L_S(n, n/s)}{n/s} = (s-1)L - s + 1 + \log s + \max_{\substack{p_i\ge 0\ (i\in S^{(n)})\\ \sum_i p_i = 1,\ \sum_i i p_i = \alpha - s}} \left[-\sum_i p_i \log p_i - \sum_i p_i \log d_{\alpha - i}\right].$$

The definition of the remainder in (3.1a) gives

$$h_n(i) = -\frac{q}{2}i^2 + O(i^2/\alpha) \quad \text{for every fixed } i, \qquad h_n(i) \le -\frac{q}{2}i^2 \quad (i \ge -1). \tag{A.1}$$

The first estimate is uniform on every fixed finite set of deficits; the second is global on the finite support.

Finite and limiting duals and optimizing tilts.

We first solve the mean constraint by exponential tilting. The target center satisfies

$$\frac{2}{q} \le T_0 \le 1 + \frac{2}{q}, \tag{A.2}$$

so it lies in the interior of K_*. Every target in K_* is also strictly inside the feasible mean range of both supports.

Choose $\Lambda > 0$ so that

$$M_S(-\Lambda) < \min K_* < \max K_* < M_S(\Lambda) \qquad (S = S_+, S_4).$$

Such a choice exists because the limiting mean maps are strictly increasing and their endpoint limits bracket K_*. We quantify the effect of the finite score before inverting the mean map. Write

$$w_i(\lambda) := \exp\left(\lambda i - \frac{q}{2}i^2\right), \qquad \rho_n(i) := -h_n(i) - \frac{q}{2}i^2$$

on the finite support. By (3.1a), the score loss is nonnegative, and

$$0 \le \rho_n(-1) = \log(1 + 1/\alpha) \le 1/\alpha.$$

For $0 \le i \le \alpha/2$, use $-\log(1-x) \le 2x$ on $0 \le x \le 1/2$ to get

$$0 \le \rho_n(i) = -\sum_{r=0}^{i-1} \log\Big(1 - \frac{r}{\alpha}\Big) \le \frac{i(i-1)}{\alpha} \le \frac{i^2}{\alpha}.$$

Since $0 \le 1 - e^{-\rho_n(i)} \le \rho_n(i)$, for $m = 0, 1, 2$ the retained low indices satisfy

$$\sum_{\substack{i \in S^{(n)} \\ i \le \alpha/2}} |i|^m \left| e^{\lambda i + h_n(i)} - w_i(\lambda) \right| \le \frac{C_m}{\alpha} \qquad (|\lambda| \le \Lambda).$$

Here $|i|^0$ means one. The constants are finite because every polynomial times $\exp(\Lambda|i| - qi^2/2)$ is summable.

For the remaining indices, (A.1) bounds the weight difference by the limiting weight. The omitted indices beyond the finite cutoff are covered by the same tail. For sufficiently large α,

$$\sum_{i > \alpha/2} i^m e^{\Lambda i - qi^2/2} \le \sum_{i > \alpha/2} i^m e^{-qi^2/4} \le C_m e^{-c\alpha^2}.$$

The last bound follows by comparison with a Gaussian integral, absorbing the fixed polynomial into a smaller Gaussian exponent. The tail for S_4 is eventually empty. Thus the zeroth, first, and second tilted moments differ by $O(1/\alpha)$ uniformly on $[-\Lambda, \Lambda]$, for both supports. Their zeroth moments have a common positive lower bound: retain the $i = 2$ term, which is present in both finite supports, and use the bounded tilt interval and bounded $h_n(2)$. Taking logarithms and ratios therefore gives, uniformly on this interval,

$$\begin{aligned} \log Z_{n,S} - \log Z_S &= O(1/\alpha), \\ M_{n,S} - M_S &= O(1/\alpha), \\ M'_{n,S} - M'_S &= O(1/\alpha). \end{aligned}$$

For all sufficiently large n, the finite support $S^{(n)}$ contains at least two indices, and hence $M'_{n,S}(\lambda) = \mathrm{Var}_{n,S,\lambda}(i) > 0$ for every $\lambda \in \mathbb{R}$. Likewise, $M'_S(\lambda) = \mathrm{Var}_{S,\lambda}(i) > 0$ for every λ. Thus both mean maps are globally strictly increasing. By continuity and compactness,

$$v_* := \frac{1}{2} \min_{\substack{S \in \{S_+, S_4\} \\ |\lambda| \le \Lambda}} M'_S(\lambda) > 0.$$

For all sufficiently large n, one has $M'_{n,S}(\lambda) \ge v_*$ on the same interval, and the finite mean maps also bracket K_*. Hence, for every $T \in K_*$, the equations

$$M_{n,S}(\lambda_{n,S}(T)) = T, \qquad M_S(\lambda_S(T)) = T$$

have unique solutions in $[-\Lambda, \Lambda]$. The inverse mean maps are uniformly v_*^{-1}-Lipschitz, and therefore

$$\sup_{T \in K_*} |\lambda_{n,S}(T) - \lambda_S(T)| \le v_*^{-1} \sup_{|\lambda| \le \Lambda} |M_{n,S}(\lambda) - M_S(\lambda)| = O(1/\alpha).$$

For the finite supports, the entropy functional in (3.8a) is strictly concave on its feasible simplex. Its unique optimizer is

$$p_{n,S,T}(i) = \frac{\exp\{\lambda_{n,S}(T) i + h_n(i)\}}{Z_{n,S}(\lambda_{n,S}(T))},$$

and its dual value is exactly

$$\mathcal{F}_{n,S}(T) = \log Z_{n,S}(\lambda_{n,S}(T)) - \lambda_{n,S}(T) T.$$

For the countable support S_+, we must also show that the entropy maximum is attained. Fix $T \in K_*$, write $\lambda = \lambda_{S_+}(T)$, and set

$$\pi_\lambda(i) := \frac{\exp\{\lambda i - qi^2/2\}}{Z_{S_+}(\lambda)}.$$

Write

$$D_{\mathrm{KL}}(p \| \pi_\lambda) := \sum_i p_i \log \frac{p_i}{\pi_\lambda(i)},$$

with the convention that a summand with $p_i = 0$ is zero. This probability distribution has mean T. For every feasible probability vector p with finite quadratic moment,

$$-\sum_i p_i \log p_i - \frac{q}{2}\sum_i i^2 p_i = \log Z_{S_+}(\lambda) - \lambda T - D_{\mathrm{KL}}(p\|\pi_\lambda). \tag{A.3}$$

The identity follows first for finitely supported p and then by truncation; finite quadratic moment also makes the entropy finite by comparison with a Gaussian law. We interpret the objective as $-\infty$ when the quadratic moment diverges. Since relative entropy is nonnegative and vanishes only at $p = \pi_\lambda$, (A.3) proves finiteness, attainment, uniqueness, and the dual formula for S_+. The same identity also recovers the finite-support case. The dual objectives for the finite and limiting supports attain their minima in the same tilt interval. Their difference at a fixed tilt is $\log Z_{n,S} - \log Z_S$, so the difference of their minimum values is at most the uniform norm of this quantity. Together with the tilt estimate, this proves (3.9a) and

$$\varepsilon_n^{\mathrm{dual}} = O(1/\alpha).$$

On the finite support S_4, the score perturbation and tilt error in the optimizer formula give the coordinatewise rate (3.9b). The limiting coordinates are positive and continuous in T on the compact set K_*; their positive minimum, together with this rate, gives the asserted uniform lower bound on the finite coordinates.

Exact cancellation between supports.

We next return to the counting exponent. Substitution of (3.1b) gives

$$\begin{aligned}\Psi_{n,S}(s) &:= \frac{L_S(n, n/s)}{n/s} \\ &= (s-1)L - s + 1 + \log s + A_n + B_n(\alpha - s) + \mathcal{F}_{n,S}(\alpha - s).\end{aligned} \tag{A.4}$$

Hold k fixed at a value feasible for both supports. The mean class size s and the mean deficit $T = \alpha - s$ are then the same, so the affine terms in (A.4) cancel. The difference is exactly the difference of the two finite dual values, as asserted in (3.8). This compares the supports at the same number of parts and vertices, before either dual value is replaced by its limit.

The exponent at the phase center and its nearby slope.

To locate the zero, we estimate the exponent at the phase center and its derivative nearby. At $s = s_0$, let $T_0 = \alpha - s_0$. Applying Stirling to A_n in (A.4) and collecting the affine terms gives

$$\begin{aligned}\Psi_{n,S}(s_0) - \mathcal{F}_{n,S}(T_0) = s_0\left(L - \log\alpha - \frac{q}{2}s_0 + \frac{q}{2}\right) - L + T_0 + 1 + \log s_0 + \frac{q}{2}T_0^2 \\ - \frac{1}{2}\log(2\pi\alpha) + O(1/L).\end{aligned} \tag{A.5}$$

Equation (2.1) gives the exact identity

$$\frac{qs_0}{2} = L - \ell + \log q - q.$$

Since $\alpha = s_0 + T_0$ and T_0 is uniformly bounded,

$$\log\alpha = \log s_0 + \frac{T_0}{s_0} + O(L^{-2}), \qquad \log s_0 = \ell + q - \log q + O(\ell/L).$$

Consequently

$$\begin{aligned}s_0\left(L - \log\alpha - \frac{q}{2}s_0 + \frac{q}{2}\right) &= -T_0 + \frac{q}{2}s_0 + O(\ell) \\ &= -T_0 + L - \ell + \log q - q + O(\ell).\end{aligned}$$

Substitution in (A.5) cancels the terms L and T_0 and leaves $O(\ell)$ uniformly in the phase. The dual term in (A.4) is $O(1)$ uniformly on K_*, by the Gaussian majorant above. Therefore

$$\Psi_{n,S}(s_0) = O(\ell) \tag{A.6}$$

uniformly for both supports.

The phase formula for α and the expansion of $\log\alpha$ also give

$$B_n = 2L - \ell + O(1) \tag{A.7}$$

uniformly in δ. Strict concavity and the optimizer formula imply

$$\frac{d}{dT}\mathcal{F}_{n,S}(T) = -\lambda_{n,S}(T).$$

Differentiating (A.4), with $T = \alpha - s$, yields

$$\Psi'_{n,S}(s) = L - 1 + s^{-1} - B_n + \lambda_{n,S}(T) = -L + O_A(\ell) \tag{A.8}$$

uniformly whenever $|s - s_0| \le A\ell/L$. Indeed, the image of this corridor under $T = \alpha - s$ lies inside K_* for all sufficiently large n, because K_* has a fixed margin around the complete range of T_0.

Roots and slopes in the phase corridor.

Set

$$\Psi_{n,S,c}(s) := \Psi_{n,S}(s) + c.$$

By (A.6), there is an absolute C_0 such that $|\Psi_{n,S,c}(s_0)| \le C_0\ell$ for both supports and every $0 \le c \le q$. By (A.8), after one phase-independent eventuality threshold,

$$\Psi'_{n,S,c}(s) \le -\frac{1}{2}L$$

throughout each fixed corridor. Choose $A_* > 2C_0$. Integration gives

$$\Psi_{n,S,c}\left(s_0 - A_*\frac{\ell}{L}\right) \ge \left(\frac{A_*}{2} - C_0\right)\ell > 0,$$
$$\Psi_{n,S,c}\left(s_0 + A_*\frac{\ell}{L}\right) \le \left(C_0 - \frac{A_*}{2}\right)\ell < 0.$$

The endpoint signs give a zero, and strict decrease in s gives uniqueness on this interval. Since $k = n/s > 0$, this is the *corridor* root $r_{S,c}$ of $L_S(n,k) + ck = 0$. We have proved (3.6); the argument makes no claim about zeros outside the corridor.

Finally,

$$L_S(n,k) + ck = k\Psi_{n,S,c}(s), \qquad \frac{ds}{dk} = -\frac{s}{k}.$$

Hence

$$\begin{aligned}\frac{\partial}{\partial k}\{L_S(n,k) + ck\} &= \Psi_{n,S,c}(s) - s\Psi'_{n,S}(s)\\ &= \frac{2}{q}L^2 + O_A(L\ell),\end{aligned} \tag{A.9}$$

because $s = (2/q)L + O_A(\ell)$, (A.8) holds, and $\Psi_{n,S,c}(s) = O_A(\ell)$ throughout the corridor. This proves (3.7) and $\varepsilon^{\text{slope}}_{n,A} \to 0$ with one eventuality threshold for the complete phase. □

Appendix B. Rounding the selected profile and its first moment

Rounding and first-moment estimates in Proposition 5.3. The correction satisfies $\Delta k_2 + \Delta k_3 = -e_0$ and $2\Delta k_2 + 3\Delta k_3 = -e_1$, so it removes both constraint errors by changing $O(1)$ counts. To estimate the resulting loss in the counting exponent, put $\mathbf{k}^* = (k_{\text{co}}p_i^{(n)})_{i=2}^5$ and let $\Delta^{\text{tot}} = \mathbf{k} - \mathbf{k}^*$ be the total rounding and correction displacement. There is a phase-independent constant C_Δ such that $\|\Delta^{\text{tot}}\|_\infty \le C_\Delta$. Moreover, Δ^{tot} satisfies both homogeneous constraint equations, because $\mathbf{k}^*$ and $\mathbf{k}$ obey the same two constraints; thus it is tangent to the feasible affine plane.

For Taylor's formula we need the optimizer and the corrected profile to be joined inside the positive orthant. By (5.15), $k_i^* \ge c_p k_{\text{co}}$; hence, uniformly in $t \in [0,1]$ and in the phase,

$$k_i^* + t\Delta_i^{\text{tot}} \ge c_p k_{\text{co}} - C_\Delta \ge \frac{c_p}{2}k_{\text{co}}$$

for all sufficiently large n. In particular the correction preserves positivity. At the exact constrained optimizer the linear term in every tangent direction vanishes, while at a positive vector x the Hessian of the entropy term is $-\operatorname{diag}(1/x_i)$. Its operator norm along the rounding segment is therefore at most $2/(c_p k_{\text{co}})$. Taylor's formula with integral remainder gives

$$|L_{\mathbf{k}^*+\Delta^{\mathrm{tot}}} - L_{\mathbf{k}^*}| \le \frac{1}{2} \sup_{0\le t\le 1} \left\|\nabla^2 L_{\mathbf{k}^*+t\Delta^{\mathrm{tot}}}\right\|_{\mathrm{op}} \|\Delta^{\mathrm{tot}}\|_2^2 = O(1/k_{\mathrm{co}}).$$

The fixed displacement above the root provides the margin: $\Phi_n(r_4^{\mathrm{co}}) = 0$ and $16 \le k_{\mathrm{co}} - r_4^{\mathrm{co}} < 17$. By Equation (3.7), with $S = S_4$ and $c = \log 2$, there is a phase-independent $c_1 > 0$ such that $\Phi_n'(k) \ge c_1(\log n)^2$ throughout this interval for all sufficiently large n. Integration therefore gives

$$L_{S_4}(n, k_{\mathrm{co}}) + (\log 2)k_{\mathrm{co}} = \Phi_n(k_{\mathrm{co}}) \ge 16c_1(\log n)^2.$$

The tangent rounding loss above is $O(k_{\mathrm{co}}^{-1})$, and the five Stirling approximations contribute $O(\log n)$ in total. Both are smaller than the integrated margin; decreasing c_Z proves the assertion. □

Appendix C. An explicit tail bound for amplification

Lemma C.1 (Amplification by vertex exposure)**.** *There is an absolute constant $C > 0$ and a deterministic sequence $\varepsilon_n^{\mathrm{left}} \ge 0$ with $\varepsilon_n^{\mathrm{left}} \to 0$ such that the following holds. Suppose deterministic integers $k_n \ge 0$ and deterministic reals $\Lambda_n \ge 0$ satisfy, for all sufficiently large n,*

$$\mathbb{P}(\zeta(G_n) \le k_n) \ge e^{-\Lambda_n}. \tag{C.1}$$

For every deterministic choice $r = r(n) > 0$, with the same C and $\varepsilon_n^{\mathrm{left}}$ independent of k_n, Λ_n, r, the bound

$$\begin{aligned}\mathbb{P}\left(\zeta(G_n) > k_n + C\left(\frac{\sqrt{n\Lambda_n} + \sqrt{nr}}{\log n} + n^{1/3} + 1\right)\right) \\ \le e^{-r} + \varepsilon_n^{\mathrm{left}}.\end{aligned} \tag{C.2}$$

holds at every sufficiently large n for which (C.1) holds.

Proof. Let $\mathcal{G}_n$ be the simultaneous event in Lemma 10.1, and set $\varepsilon_n^{\mathrm{left}} := \mathbb{P}(\mathcal{G}_n^c) = o(1)$. This sequence is independent of k_n, Λ_n, and r.

Let S_{k_n} be the maximum cocolorable-set size defined in (10.4). The vertex-exposure argument in the proof of Lemma 10.2 gives (10.5), with $k = k_n$ and $\Lambda = \Lambda_n$.

The lower tail with radius $\sqrt{(n-1)r/2}$ gives, outside an event of probability at most e^{-r},

$$n - S_{k_n} \le \sqrt{(n-1)\Lambda_n/2} + \sqrt{(n-1)r/2}. \tag{C.3}$$

Choose a maximizing set and its complement V_{left} as in (10.6).

The set V_{left} was selected from the graph, but $\mathcal{G}_n$ controls it along with every other set. On this event and the event in (C.3), apply (10.3) with $|V_{\mathrm{left}}| = n - S_{k_n}$ in (10.6). Enlarging the absolute constant C gives the additional parts in (C.2). The two exceptional events have total probability at most $e^{-r} + \varepsilon_n^{\mathrm{left}}$. The bound holds at every sufficiently large n satisfying (C.1), with C and $\varepsilon_n^{\mathrm{left}}$ independent of the three deterministic parameters. □

Appendix D. Uniform Stirling expansion and adjacent-size ratios

Proof of the phase-expansion lemma in Section 2. Because $\alpha = O(L)$ uniformly in the phase, one has $\alpha < n/2$ for all sufficiently large n. For $0 \le x \le 1/2$, $-2x \le \log(1-x) \le -x$. Hence

$$\log(n)_\alpha = \alpha L + \sum_{j=0}^{\alpha-1} \log\left(1 - \frac{j}{n}\right) = \alpha L + R_{n,\alpha}, \qquad |R_{n,\alpha}| \le \frac{\alpha^2}{n}.$$

Using the uniform Stirling remainder

$$\log(\alpha!) = \alpha\log\alpha - \alpha + \frac{1}{2}\log(2\pi\alpha) + O(1/\alpha)$$

in

$$\log \mu_\alpha = \log(n)_\alpha - \log(\alpha!) - q\binom{\alpha}{2}$$

gives

$$\log \mu_\alpha = \alpha \left(L - \log \alpha + 1 - \frac{q}{2}(\alpha - 1) \right) - \frac{1}{2} \log(2\pi\alpha) + O(1/L). \tag{2.5}$$

All implied constants and remainder bounds here are uniform in δ.

Write

$$\alpha = \frac{2L}{q}(1 + x_n), \qquad x_n := \frac{-\ell + C + qb/2}{L}.$$

Uniformly for the actual range $0 < b \le 1$ (and on its closure), one has $|x_n| \le 1/2$ for all sufficiently large n. The Taylor formula $\log(1 + x) = x + O(x^2)$ therefore gives

$$\log \alpha = \ell + q - \log q + \frac{-\ell + C + qb/2}{L} + O\left(\frac{\ell^2}{L^2} \right).$$

Since

$$\frac{q}{2}(\alpha - 1) = L - \ell + C - \frac{q\delta}{2},$$

we obtain the uniform expansion

$$L - \log \alpha + 1 - \frac{q}{2}(\alpha - 1) = \frac{q\delta}{2} + \frac{\ell - C - qb/2}{L} + O\left(\frac{\ell^2}{L^2} \right). \tag{2.6}$$

Multiplying by $\alpha = 2(L - \ell + C)/q + b$ and collecting terms yields

$$\begin{aligned} \alpha \left(L - \log \alpha + 1 - \frac{q}{2}(\alpha - 1) \right) = \delta L + \left(\frac{2}{q} - \delta \right) \ell \\ + \delta C + \frac{q}{2}\delta(1 - \delta) - \frac{2C}{q} - (1 - \delta) + O\left(\frac{1 + \ell^2}{L} \right). \end{aligned}$$

Also

$$-\frac{1}{2} \log(2\pi\alpha) = -\frac{1}{2}\ell - \frac{1}{2} \log(2\pi) - \frac{q}{2} + \frac{1}{2} \log q + O(\ell/L).$$

Thus (2.2) holds with

$$\begin{aligned} K(\delta) := \delta C + \frac{q}{2}\delta(1 - \delta) - \frac{2C}{q} - (1 - \delta) \\ - \frac{1}{2} \log(2\pi) - \frac{q}{2} + \frac{1}{2} \log q. \end{aligned} \tag{2.7}$$

This function is continuous, hence bounded, on $[0, 1]$. The combined remainder in (2.5)–(2.7) has absolute value at most $C'(1 + \ell^2)/L$ for an absolute C'. Taking $C_{\mathrm{ph}} \ge C'$ gives $|E_n^{\mathrm{ph}}| \le \varepsilon_n^{\mathrm{ph}}$ in (2.2), where E_n^{ph} denotes the combined remainder just obtained.

For integers $1 \le s < n$, the adjacent-size ratios are exact:

$$\frac{\mu_{s+1}}{\mu_s} = \frac{n - s}{s + 1} 2^{-s}, \qquad \frac{\mu_{s-1}}{\mu_s} = \frac{s}{n - s + 1} 2^{s-1}. \tag{2.8}$$

The displayed formula for α gives the exact phase representation

$$2^\alpha = \exp(q\alpha) = \exp(2C + qb) \frac{n^2}{L^2}.$$

Since $0 < b \le 1$, the multiplicative factor $\exp(2C + qb)$ is bounded above and below by positive absolute constants. Therefore, for every fixed M, uniformly over integers s with $|s - \alpha| \le M$,

$$\frac{\mu_{s+1}}{\mu_s} = \Theta(L/n), \qquad \frac{\mu_{s-1}}{\mu_s} = \Theta(n/L).$$

In particular,

$$\mu_{\alpha+2} = \mu_\alpha \, \Theta(L^2/n^2), \qquad \mu_{\alpha-2} = \mu_\alpha \, \Theta(n^2/L^2). \tag{2.8a}$$

Taking logarithms and using (2.2) proves the first line of (2.3). Since $\delta \le 1$, its leading term is at most $-L$, and the coefficient of ℓ is uniformly bounded; this proves the second line.

For the lower adjacent size, (2.2) and (2.8a) give

$$\log \mu_{\alpha-2} = 2L + \left(\frac{2}{q} - \frac{5}{2} \right) \ell + \delta(L - \ell) + K(\delta) + O(1).$$

For large n, $L - \ell > 0$, while K is bounded below. Exponentiating gives (2.4), with one absolute constant c and one eventuality threshold valid for the complete phase. □

Appendix E. Tilt brackets and tails for the entropy bound

Proof of Lemma 5.2. Fix $\delta \in [0, 1]$, put $T_0 = 1 + 2/\log 2 - \delta$, and abbreviate

$$\lambda_4 := \lambda_{S_4}(T_0).$$

We will evaluate the unrestricted dual function at this four-size tilt. To control the omitted weights, first bracket the tilt. At $\lambda = 2\log 2$, set $j = i - 2$, which maps S_4 onto $\{0, 1, 2, 3\}$. Substitution of $i = j + 2$ gives

$$\begin{aligned} w_{j+2}(2\log 2) &= \exp\left(2\log 2(j+2) - \frac{\log 2}{2}(j+2)^2\right) \\ &= \exp\left(2\log 2 - \frac{\log 2}{2} j^2\right) = 4\, 2^{-j^2/2}. \end{aligned}$$

Consequently, reindexing the two sums in the definition of the mean gives

$$\begin{aligned} Z_{S_4}(2\log 2) &= \sum_{i=2}^{5} w_i(2\log 2) = 4 \sum_{j=0}^{3} 2^{-j^2/2}, \\ \sum_{i=2}^{5} i\, w_i(2\log 2) &= 4 \sum_{j=0}^{3} (j+2) 2^{-j^2/2} \\ &= 8 \sum_{j=0}^{3} 2^{-j^2/2} + 4 \sum_{j=0}^{3} j 2^{-j^2/2}. \end{aligned}$$

Dividing the second equality by the first gives

$$\begin{aligned} M_{S_4}(2\log 2) &= 2 + \frac{\sum_{j=0}^{3} j 2^{-j^2/2}}{\sum_{j=0}^{3} 2^{-j^2/2}} \\ &< 2 + \frac{\sum_{j\geq 0} j 2^{-j^2/2}}{\sum_{j\geq 0} 2^{-j^2/2}} < 2 + \frac{4}{5} < \frac{2}{\log 2} \leq T_0. \end{aligned} \tag{5.4}$$

We verify the first strict inequality in (5.4) directly. Put

$$Z_0 = \sum_{j=0}^{3} 2^{-j^2/2}, \quad N_0 = \sum_{j=0}^{3} j 2^{-j^2/2}, \quad Z_1 = \sum_{j=4}^{\infty} 2^{-j^2/2}, \quad N_1 = \sum_{j=4}^{\infty} j 2^{-j^2/2}.$$

Here $N_0/Z_0 \leq 3$, while $N_1/Z_1 \geq 4$. Since $Z_0, Z_1 > 0$,

$$\frac{N_0 + N_1}{Z_0 + Z_1} - \frac{N_0}{Z_0} = \frac{Z_0 N_1 - N_0 Z_1}{Z_0(Z_0 + Z_1)} > 0.$$

Thus adjoining the terms whose new indices satisfy $j \geq 4$ strictly increases the weighted mean. For the next inequality, put $a_j = j 2^{-j^2/2}$. For $j \geq 4$,

$$\frac{a_{j+1}}{a_j} = \frac{j+1}{j} 2^{-(2j+1)/2} \leq \frac{1}{16},$$

and $a_4 = 1/64$. Hence

$$\sum_{j\geq 4} j 2^{-j^2/2} \leq \frac{1/64}{1 - 1/16} = \frac{1}{60}.$$

Writing $\vartheta = 2^{-1/2}$, retaining $j = 0, 1, 2$ in the denominator and $j = 0, 1, 2, 3$ in the numerator gives

$$\frac{\sum_{j\geq 0} j 2^{-j^2/2}}{\sum_{j\geq 0} 2^{-j^2/2}} < \frac{19\vartheta/16 + 1/2 + 1/60}{5/4 + \vartheta} < \frac{4}{5},$$

where the last inequality follows from $\vartheta < 71/100$. Finally, $2/3 < \log 2 < 5/7$ gives $2 + 4/5 < 2/\log 2$.

At $\lambda = 9\log 2/2$, put $t = 2^{-1/8}$. The identity

$$\frac{9\log 2}{2} i - \frac{\log 2}{2} i^2 = \frac{\log 2}{8}\{81 - (2i-9)^2\}$$

gives

$$w_i(9\log 2/2) = 2^{81/8} t^{(2i-9)^2}.$$

Thus the four weights, in the order $i = 2, 3, 4, 5$, are proportional to

$$t^{25}, \quad t^9, \quad t, \quad t.$$

The numerator of $M_{S_4}(9\log 2/2) - 4$ is therefore

$$t - t^9 - 2t^{25} = t(1 - t^8 - 2t^{24}) = t/4 > 0, \tag{5.5}$$

and the denominator is positive. Hence this mean is greater than four, whereas $T_0 \le 1 + 2/\log 2 < 4$. Since M_{S_4} is strictly increasing, the two mean comparisons give

$$2\log 2 < \lambda_4 < 9\log 2/2. \tag{5.6}$$

For a tilt λ, define the omitted low- and high-deficit ratios

$$R_{\mathrm{low}}(\lambda) := \frac{\sum_{i=-1}^{1} w_i(\lambda)}{Z_{S_4}(\lambda)}, \qquad R_{\mathrm{high}}(\lambda) := \frac{\sum_{i \ge 6} w_i(\lambda)}{Z_{S_4}(\lambda)}.$$

Let $M_A(\lambda)$ denote the mean of i under the weights $w_i(\lambda)$ restricted to an index set A. Differentiation gives

$$\frac{d}{d\lambda} \log R_{\mathrm{low}}(\lambda) = M_{\{-1,0,1\}}(\lambda) - M_{S_4}(\lambda) \le 1 - 2 = -1,$$
$$\frac{d}{d\lambda} \log R_{\mathrm{high}}(\lambda) = M_{\{6,7,\dots\}}(\lambda) - M_{S_4}(\lambda) \ge 6 - 5 = 1.$$

Thus R_{low} decreases and R_{high} increases on $\mathbb{R}$. We can bound them at the endpoints of the tilt interval, splitting at its intermediate reference tilt. The required values follow from the direct calculations and Gaussian-tail bounds below:

$$R_{\mathrm{low}}(2\log 2) < \frac{51}{100}, \quad R_{\mathrm{high}}(3\log 2) < \frac{1}{50}, \quad R_{\mathrm{low}}(3\log 2) < \frac{3}{25}, \quad R_{\mathrm{high}}(9\log 2/2) < \frac{1}{4}. \tag{5.7}$$

At $2\log 2$, with $\vartheta = 2^{-1/2}$,

$$R_{\mathrm{low}}(2\log 2) = \frac{\vartheta + 1/4 + \vartheta/16}{1 + \vartheta + 1/4 + \vartheta/16} < 51/100.$$

At $\lambda = 3\log 2$,

$$w_i(3\log 2) = \exp\left((\log 2)\left(3i - \frac{i^2}{2}\right)\right) = 2^{9/2} 2^{-(i-3)^2/2}.$$

Set $\vartheta = 2^{-1/2}$. After canceling the common factor $2^{9/2}$, the weight at index i is $\vartheta^{(i-3)^2}$. Thus the four S_4-weights, for $i = 2, 3, 4, 5$, are

$$\vartheta, \quad 1, \quad \vartheta, \quad \vartheta^4,$$

and

$$2^{-9/2} Z_{S_4}(3\log 2) = 1 + 2\vartheta + \vartheta^4 = \frac{5}{4} + 2\vartheta.$$

For the omitted low indices $i = -1, 0, 1$, the weights are

$$(\vartheta^{16}, \vartheta^9, \vartheta^4) = \left(\frac{1}{256}, \frac{\vartheta}{16}, \frac{1}{4}\right).$$

For the omitted high indices $i \ge 6$, they are

$$\vartheta^9 + \vartheta^{16} + \sum_{r \ge 5} \vartheta^{r^2}.$$

The first term in the remaining sum is $\vartheta^{25} = \vartheta/4096$, and the ratio between successive terms is at most $\vartheta^{11} < 1/32$. Therefore

$$\sum_{r\geq 5} \vartheta^{r^2} \leq \frac{\vartheta^{25}}{1-1/32} < \frac{1}{3968}.$$

It follows that

$$R_{\text{low}}(3\log 2) = \frac{1/256 + \vartheta/16 + 1/4}{5/4 + 2\vartheta} < \frac{3}{25},$$

and

$$R_{\text{high}}(3\log 2) \leq \frac{\vartheta/16 + 1/256 + 1/3968}{5/4 + 2\vartheta} < \frac{1}{50},$$

where the final inequalities follow from $7/10 < \vartheta < 71/100$.

Finally, at $\lambda = 9\log 2/2$,

$$\frac{w_i(9\log 2/2)}{2^{81/8}} = t^{(2i-9)^2}.$$

After dividing by the common factor t of the $i = 4, 5$ weights, the high indices $i = 6, 7, 8$ contribute t^8, t^{24}, t^{48}. From $i = 8$ onward the ratio between successive terms is at most $1/16$, so the tail beginning with t^{48} is less than $1/60$. The normalized S_4 denominator is $2 + t^8 + t^{24}$. Hence $R_{\text{high}}(9\log 2/2) < 1/4$ follows from

$$3(t^8 + t^{24}) + 4/60 < 2, \tag{5.8}$$

because $t^8 = 1/2$ and $t^{24} = 1/8$.

If $\lambda_4 \leq 3\log 2$, monotonicity and (5.7) give $R_{\text{low}}(\lambda_4) + R_{\text{high}}(\lambda_4) < 53/100$; if $\lambda_4 \geq 3\log 2$, they give $R_{\text{low}}(\lambda_4) + R_{\text{high}}(\lambda_4) < 37/100$. By the dual representation above,

$$\mathcal{F}_{S_4}(T_0) = \log Z_{S_4}(\lambda_4) - \lambda_4 T_0,$$

whereas evaluating the S_+ dual function at the same parameter gives

$$\mathcal{F}_{S_+}(T_0) \leq \log Z_{S_+}(\lambda_4) - \lambda_4 T_0.$$

Subtracting these two relations yields

$$\begin{aligned} D_4(\delta) &\leq \log \frac{Z_{S_+}(\lambda_4)}{Z_{S_4}(\lambda_4)} \\ &= \log\bigl(1 + R_{\text{low}}(\lambda_4) + R_{\text{high}}(\lambda_4)\bigr) < \log(153/100). \end{aligned}$$

Since S_4 is a subset of S_+, the loss is nonnegative. Subtracting from $\log 2$ completes the proof. □

Appendix F. Three-range estimates for common subprofiles

Proof of Lemma 7.1. Set

$$\eta := \frac{\log\log n}{32\log n}.$$

For all sufficiently large n we have $\eta < 31/32$. Divide the selected mass into $m \leq \eta n$, $m > \eta n$ with $n - m > n/32$, and $n - m \leq n/32$. These ranges are disjoint. The first contains the term that will contribute one to the sum; we show that all other terms in it, and both remaining ranges, have vanishing total contribution. We use (7.4) in the first range, a rate estimate in the second, and (7.6) in the third.

Small selected mass.

Let

$$\xi_i = \frac{k_i^2}{2\mu_{u_i}(n)}, \qquad \Xi_{\text{empty}} = \sum_i \xi_i. \tag{7.8}$$

Equations (2.4), (2.8), and $k_i = \Theta(n/\log n)$ imply, with one phase-independent constant,

$$\Xi_{\text{empty}} = O\Bigl((\log n)^{-(2/\log 2 - 1/2)}\Bigr). \tag{7.9}$$

For $m \leq \eta n$, we compare every denominator in the recurrence with the corresponding first moment on the full ambient vertex set.

Fix a final subprofile ℓ with selected mass $m \le \eta n$, and expose its blocks in any order. At an intermediate selected mass $m' \le m$, the exact ratio of first moments is

$$\frac{\mu_{u_i}(n)}{\mu_{u_i}(n-m')} = \frac{(n)_{u_i}}{(n-m')_{u_i}} = \prod_{r=0}^{u_i-1} \frac{n-r}{n-m'-r}.$$

Uniformly in the phase, $u_i = O(\log n)$ whereas $\eta n \gg \log n$. Hence, for all sufficiently large n, $m' + r \le 2\eta n$ throughout this product, and therefore

$$\frac{\mu_{u_i}(n)}{\mu_{u_i}(n-m')} \le (1-2\eta)^{-u_i}.$$

The phase expansion and $\log 2 > 2/3$ give $u_i \le 3 \log n$ eventually. Since $\eta \le 1/4$ and $-\log(1-2\eta) \le 4\eta$,

$$(1-2\eta)^{-u_i} \le \exp(4\eta u_i) \le \exp\left(\frac{3}{8}\log\log n\right) = (\log n)^{3/8}. \tag{7.10}$$

Every threshold in this estimate is independent of the phase.

Suppose the current partial profile is a and the next selected block has type i. The exact recurrence (7.4), the inequality $(k_i - a_i)^2 \le k_i^2$, and (7.10) give

$$\frac{D(a+e_i)}{D(a)} \le \frac{(\log n)^{3/8}\xi_i}{a_i+1}.$$

Multiplication along any exposure order gives the order-independent bound

$$D(\ell) \le \prod_i \frac{\left((\log n)^{3/8}\xi_i\right)^{\ell_i}}{\ell_i!}.$$

All summands are nonnegative, so enlarging from the selected-mass range to the whole nonnegative four-dimensional lattice and expanding the exponential series yields

$$1 \le \sum_{m \le \eta n} D(\ell) \le \exp\left((\log n)^{3/8}\Xi_{\text{empty}}\right). \tag{7.11}$$

The lower bound is the empty marked subprofile.

The exponent tends to zero with a fixed power margin. By (7.9),

$$(\log n)^{3/8}\Xi_{\text{empty}} = O\left((\log n)^{-(2/q-7/8)}\right),$$

and $q < 7/10$ gives

$$\frac{2}{q} - \frac{7}{8} > \frac{20}{7} - \frac{7}{8} = \frac{111}{56} > 0.$$

Thus, for one deterministic sequence $\varepsilon_n^{\text{empty}} \to 0$ independent of the phase,

$$1 \le \sum_{m \le \eta n} D(\ell) \le e^{\varepsilon_n^{\text{empty}}} = 1 + o(1).$$

Intermediate selected mass.

Write

$$\begin{gathered} p_i = k_i/k_{\text{co}}, \quad y_i = \ell_i/k_{\text{co}}, \quad Y = \sum_i y_i, \quad I = \sum_i i y_i, \\ z_i = p_i - y_i, \quad R = \sum_i z_i = 1 - Y, \quad T := \sum_i i p_i = \alpha - \frac{n}{k_{\text{co}}}, \quad I_r = \sum_i i z_i = T - I. \end{gathered} \tag{7.12}$$

The residual vertex fraction is

$$\rho = \frac{n-m}{n} = R + \frac{I - TY}{\alpha - T}. \tag{7.13}$$

A uniform estimate for the marked weights.

Before approximating any factorial, equations (7.2) and (7.3) give the exact identity

$$\log D(\ell) = 2\sum_i \{\log(k_i!) - \log(\ell_i!) - \log((k_i-\ell_i)!)\} \\ - \ell_\bullet \log 2 - \log(n!) + \log((n-m)!) + \sum_i \ell_i \log(u_i!) + \sum_i \log(\ell_i!) + (\log 2) b_\ell.$$

Some selected or remaining type counts may vanish. We therefore use Stirling's estimate in a form valid also at zero: for an absolute constant $C_{\mathrm{S}} > 0$,

$$\log(t!) = t\log t - t + \sigma(t), \qquad |\sigma(t)| \le C_{\mathrm{S}} \log(t+2) \qquad (t \in \mathbb{N}_0),$$

with the convention $0 \log 0 = 0$. Apply it only to the ambient and profile-multiplicity factorials $n!$, $(n-m)!$, $k_i!$, $\ell_i!$, and $(k_i - \ell_i)!$; the coefficient-weighted class-size factorials $u_i!$ remain exact. Only a fixed number of factorials are approximated, and every argument is at most n. Therefore the total remainder is $O(\log n)$ uniformly, including when some ℓ_i or $k_i - \ell_i$ is zero.

Set $K = k_{\mathrm{co}}$, so $k_i = K p_i$, $\ell_i = K y_i$, and $k_i - \ell_i = K z_i$. The profile-multiplicity factorials then contribute

$$2\sum_i \log(k_i!) - \sum_i \log(\ell_i!) - 2\sum_i \log((k_i-\ell_i)!) \\ = K\left[2\sum_i p_i \log p_i - \sum_i y_i \log y_i - 2\sum_i z_i \log z_i - Y + Y\log K\right] + O(\log n).$$

Since $m = K(\alpha Y - I) = K\sum_i u_i y_i$, $n = K(\alpha - T)$, and $n - m = n\rho$, the two ambient factorials contribute

$$-\log(n!) + \log((n-m)!) = n\rho\log\rho + K\sum_i y_i\big(u_i - u_i \log n\big) + O(\log n).$$

Finally, $\ell_\bullet = K\sum_i y_i$ and $b_\ell = K\sum_i y_i \binom{u_i}{2}$, so the class-size and sign terms are exactly

$$K\sum_i y_i \left[\log(u_i!) + (\log 2)\binom{u_i}{2} - \log 2\right].$$

Adding these three displayed contributions and distributing $Y \log K = \sum_i y_i \log K$ gives

$$\log D(\ell) = n\rho\log\rho \\ + k_{\mathrm{co}} \sum_i \{2p_i \log p_i - 2(p_i - y_i)\log(p_i - y_i) - y_i \log y_i - y_i + y_i E_i\} + O(\log n), \tag{7.14}$$

where

$$E_i = \log k_{\mathrm{co}} + \log(u_i!) + u_i - u_i \log n + (\log 2)\binom{u_i}{2} - \log 2. \tag{7.15}$$

We next bound the entropy term in (7.14), allowing any coordinate of the selected subprofile to reach an endpoint. By (5.15) and the bounded rounding displacement, $p_i \ge c_p > 0$ for the exact integer profile, after decreasing c_p if needed. There is therefore an absolute C_{H} such that, for $c_p \le p \le 1$ and $0 \le y \le p$,

$$2p\log p - 2(p-y)\log(p-y) - y\log y - y \le C_{\mathrm{H}} y \log\left(\frac{\mathrm{e}}{y}\right). \tag{7.15a}$$

Here and below, $y\log(\mathrm{e}/y)$ takes its continuous value 0 at $y = 0$; the analogous convention applies to $Y\log(4\mathrm{e}/Y)$. Indeed, when $y \le p/2$, the mean-value theorem bounds $p\log p - (p-y)\log(p-y)$ by $C_{\mathrm{H}} y$, while $-y\log y$ supplies the only singular term. When $y \ge p/2$, one has $y \ge c_p/2$, and the assertion follows by compactness. Finally,

$$\sum_i y_i \log\left(\frac{\mathrm{e}}{y_i}\right) \le Y \log\left(\frac{4\mathrm{e}}{Y}\right),$$

by the entropy bound on four coordinates. The resulting sum is $O(Y\log(\mathrm{e}/Y))$ uniformly.

Since $u_{i+1} = u_i - 1$, exact subtraction in (7.15) gives

$$E_{i+1} - E_i = \log n - \log(\alpha - i) - (\log 2)(\alpha - i) + \log 2 - 1 \\ = -\frac{\log 2}{2}\alpha + O(1), \tag{7.16}$$

uniformly for $2 \le i \le 4$. For the second line, use $(\log 2)\alpha/2 = \log n - \log\log n + O(1)$ and $\log(\alpha - i) = \log\log n + O(1)$.

The coefficients have a large affine variation, but their profile-weighted mean is bounded above by the complete first moment. Let $\bar{E} = \sum_i p_i E_i$. Applying Stirling to that moment gives

$$-\frac{1}{k_{\mathrm{co}}} \log \mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big] = \sum_i p_i \log p_i - 1 + \bar{E} + o(1). \tag{7.17}$$

The phase-uniform growth estimate (5.19) is eventually positive. Solving (7.17) for $\bar{E}$ and using $-\sum_i p_i \log p_i \le \log 4$ therefore gives

$$\bar{E} = 1 - \sum_i p_i \log p_i - \frac{1}{k_{\mathrm{co}}} \log \mathbb{E} Z_{\mathbf{k}}^{\mathrm{sgn}} + o(1) \le 1 + \log 4 + o(1).$$

In particular, $\bar{E} \le C_E$ for one phase-independent constant. To expose the affine structure in (7.16), set

$$F_i := E_i + \frac{(\log 2)\alpha}{2} i.$$

Then $F_{i+1} - F_i = O(1)$. Because the support has only four consecutive indices,

$$F_i = \sum_j p_j F_j + O(1) = \bar{E} + \frac{(\log 2)\alpha}{2} T + O(1).$$

Consequently

$$E_i = \bar{E} - \frac{(\log 2)\alpha}{2}(i - T) + O(1),$$

and therefore

$$\sum_i y_i E_i \le \frac{(\log 2)\alpha}{2}(TY - I) + O(Y). \tag{7.18}$$

Only the upper bound on $\bar{E}$ is used here.

The ambient term uses the remaining vertex fraction ρ, whereas the coefficient bound uses the remaining class fraction R. We put them on the same scale. Since $2 \le i \le 5$ and $2 \le T \le 5$,

$$|I - TY| = \left| \sum_i (i - T) y_i \right| \le 3Y,$$

so (7.13) gives

$$|\rho - R| \le \frac{3Y}{\alpha - T} = O(Y/\alpha).$$

If $\rho \ge 1/32$, then $Y \le 1$ and the preceding bound implies $R \ge 1/64$ for all sufficiently large n. Hence $x \mapsto x \log x$ is uniformly Lipschitz on the interval between R and ρ. Moreover, $-R \log R \le 1 - R = Y$. Choose a fixed Lipschitz constant $C_{\mathrm{lip}} > 0$ for $x \mapsto x \log x$ on $[1/64, 1]$. Then

$$\begin{aligned} &|n\rho \log \rho - k_{\mathrm{co}} \alpha R \log R| \\ &\quad \le k_{\mathrm{co}}(\alpha - T) C_{\mathrm{lip}} |\rho - R| + k_{\mathrm{co}} T |R \log R| \le (3C_{\mathrm{lip}} + 5) k_{\mathrm{co}} Y, \end{aligned}$$

which proves

$$n\rho \log \rho = k_{\mathrm{co}} \alpha R \log R + O(k_{\mathrm{co}} Y). \tag{7.19}$$

Under the standing condition $\rho \ge 1/32$, combining (7.14), (7.15a), (7.18), and (7.19) yields one absolute constant C_{diag} such that

$$\log D(\ell) \le k_{\mathrm{co}} \alpha \Phi_T(z) + C_{\mathrm{diag}} k_{\mathrm{co}} Y \log\left(\frac{\mathrm{e}}{Y}\right) + C_{\mathrm{diag}} \log n, \tag{7.20}$$

where

$$\Phi_T(z) = R \log R + \frac{\log 2}{2}(I_r - TR). \tag{7.21}$$

All constants and eventuality thresholds in this extraction are independent of the phase.

A negative bound for the leading rate.

The inequalities used in the main proof of Lemma 7.1 are

$$I_r - TR \le (5 - T)R, \qquad I_r - TR \le (T - 2)(1 - R). \tag{7.22}$$

Their weighted sum gives

$$I_r - TR \le 3R(1-R). \tag{7.22a}$$

For $\rho \ge 1/32$, we have $R \ge 1/64$ eventually; (3.6) and (5.13) also give $T < 4$. The logarithm inequalities in that proof imply

$$\Phi_T(z) \le -\frac{Y}{5000}. \tag{7.23}$$

We retain the error estimates below, since Y can tend to zero.

Summing the intermediate profiles.

Assume

$$m > \eta n, \qquad n - m > n/32. \tag{7.24}$$

Then $\rho > 1/32$. Moreover,

$$\frac{m}{n} = \frac{\alpha Y - I}{\alpha - T} \le \frac{\alpha Y}{\alpha - T}$$

and $m/n > \eta$ imply that, for all sufficiently large n,

$$Y > \eta \frac{\alpha - T}{\alpha} \ge \frac{\eta}{2}.$$

Define the deterministic error ratio

$$\varepsilon_n^{\mathrm{diag}} := \frac{C_{\mathrm{diag}} \log(2\mathrm{e}/\eta)}{\alpha} + \frac{2C_{\mathrm{diag}} \log n}{k_{\mathrm{co}} \alpha \eta}. \tag{7.24a}$$

Since

$$\alpha = \Theta(\log n), \qquad k_{\mathrm{co}} = \Theta(n/\log n), \qquad \eta = \Theta(\log\log n/\log n),$$

one has $\varepsilon_n^{\mathrm{diag}} \to 0$, uniformly in the phase. Because $Y \ge \eta/2$, equations (7.20) and (7.23) give

$$\log D(\ell) \le -k_{\mathrm{co}} \alpha Y \left(\frac{1}{5000} - \varepsilon_n^{\mathrm{diag}} \right).$$

For all sufficiently large n, the parenthesis is at least $1/10000$. Moreover, $\alpha\eta = \Theta(\log\log n)$. Hence there is a phase-independent $c > 0$ such that

$$D(\ell) \le \exp\{-c k_{\mathrm{co}} \log\log n\}. \tag{7.25}$$

Fix phase-independent constants $0 < c_K < C_K$ such that, eventually,

$$c_K \frac{n}{\log n} \le k_{\mathrm{co}} \le C_K \frac{n}{\log n}.$$

Define the phase-independent envelope

$$\varepsilon_n^{\mathrm{central}} := \left(C_K \frac{n}{\log n} + 1 \right)^4 \exp\left\{ -c c_K \frac{n}{\log n} \log\log n \right\}.$$

There are at most $(k_{\mathrm{co}} + 1)^4$ subprofiles, so the total central-range contribution is at most $\varepsilon_n^{\mathrm{central}}$. Since $n \log\log n/\log n \gg \log n$, this deterministic sequence tends to zero.

Nearly full selected mass.

Put $h = \mathbf{k} - \ell$ and write

$$v(h) := \sum_i u_i h_i = n - m.$$

The full corner is the range

$$v(h) \le \frac{n}{32}. \tag{7.26}$$

We use the exact reverse recurrence (7.6) without replacing any ambient falling factorial.

The residual first moments are uniformly small. For every integer $0 \le w \le n$ and every one of the four block sizes, the ratio is zero when $w < u_i$; otherwise

$$\frac{\mu_{u_i}(w)}{\mu_{u_i}(n)} = \frac{(w)_{u_i}}{(n)_{u_i}} = \prod_{r=0}^{u_i - 1} \frac{w - r}{n - r} \le \left(\frac{w}{n} \right)^{u_i}.$$

Equations (2.2) and (2.8) give $\mu_{u_i}(n) \le n^{6+o(1)}$ uniformly in the phase. Moreover, $u_i = (2/\log 2 + o(1)) \log n$, and $\log 32 = 5 \log 2$. Consequently, uniformly for $w \le n/32$,

$$\mu_{u_i}(w) \le n^{6+o(1)} 32^{-u_i} = n^{-4+o(1)}. \tag{7.26a}$$

In particular, there is one phase-independent threshold after which

$$\mu_{u_i}(w) \le n^{-3} \qquad (w \le n/32,\ 2 \le i \le 5). \tag{7.26b}$$

We use the weaker exponent -3 so that the next recurrence bound has a single threshold valid for every phase.

Fix a final residual profile h satisfying (7.26), and expose its blocks in any order. Let a be the current residual profile and $v(a) = \sum_i u_i a_i$. If the next block has type i, then $a_i < h_i$ and therefore

$$v(a) + u_i \le v(h) \le n/32.$$

The exact recurrence (7.6), $k_i - a_i \le k_i \le n$, and (7.26b) now give

$$\frac{B(a + e_i)}{B(a)} = \frac{2(k_i - a_i)\mu_{u_i}(v(a) + u_i)}{(a_i + 1)^2} \le 2n^{-2} < 1$$

for all sufficiently large n. By convention $M_{\mathbf{0}}(0) = 1$, and hence $B(0) = 1$. Thus B never increases along the path from 0 to h, and

$$B(h) \le B(0) = 1. \tag{7.27}$$

By the exact complementary identity (7.5) and the phase-uniform first-moment growth estimate (5.19),

$$D(\mathbf{k} - h) = \frac{B(h)}{\mathbb{E}\big[Z_{\mathbf{k}}^{\mathrm{sgn}}\big]} \le \exp\{-c_Z(\log n)^2\}. \tag{7.27a}$$

There are at most $(k_{\mathrm{co}} + 1)^4$ residual profiles. Therefore

$$\sum_{v(h) \le n/32} D(\mathbf{k} - h) \le \exp\{4 \log(k_{\mathrm{co}} + 1) - c_Z(\log n)^2\} = o(1)$$

uniformly in the phase. More explicitly, it is bounded by the phase-independent sequence

$$\varepsilon_n^{\mathrm{full}} := \left(C_K \frac{n}{\log n} + 1\right)^4 e^{-c_Z(\log n)^2} \longrightarrow 0.$$

Summing the three ranges.

For sufficiently large n, one has $\eta < 31/32$. The coordinate box of all partial subprofiles is then the disjoint union of

$$\begin{aligned} \mathcal{P}_n^{\mathrm{empty}} &:= \{\ell : m \le \eta n\}, \\ \mathcal{P}_n^{\mathrm{central}} &:= \{\ell : m > \eta n,\ n - m > n/32\}, \\ \mathcal{P}_n^{\mathrm{full}} &:= \{\ell : n - m \le n/32\}. \end{aligned}$$

Indeed, if $\ell \in \mathcal{P}_n^{\mathrm{full}}$, then $m \ge 31n/32 > \eta n$, so the full corner cannot meet the empty corner; after excluding those two cases, the two strict inequalities defining $\mathcal{P}_n^{\mathrm{central}}$ are automatic. Thus no boundary term is counted twice and no subprofile is omitted.

The three estimates proved above have the form

$$\begin{aligned} 1 &\le \sum_{\ell \in \mathcal{P}_n^{\mathrm{empty}}} D(\ell) \le e^{\varepsilon_n^{\mathrm{empty}}}, \\ 0 &\le \sum_{\ell \in \mathcal{P}_n^{\mathrm{central}}} D(\ell) \le \varepsilon_n^{\mathrm{central}}, \\ 0 &\le \sum_{\ell \in \mathcal{P}_n^{\mathrm{full}}} D(\ell) \le \varepsilon_n^{\mathrm{full}}, \end{aligned}$$

where every error sequence is deterministic, phase-independent, and tends to zero. The lower bound in the first line is the term $D(0) = 1$. Adding the three disjoint sums gives

$$1 \le \sum_{0 \le \ell_i \le k_i} D(\ell) \le e^{\varepsilon_n^{\mathrm{empty}}} + \varepsilon_n^{\mathrm{central}} + \varepsilon_n^{\mathrm{full}} = 1 + o(1),$$

which proves (7.7) with one eventuality threshold for the complete phase. □

Appendix G. Quadratic bounds for residual cell activities

Proof of Lemma 9.4. Fix a cell and write $\theta = \theta_{ab}$. If $\theta = 0$, then $q_{ab} = 0$, so the claim holds. Assume $\theta > 0$. If $R < 3$, the higher-reward sum is empty, so $q_{ab} = \theta^2/2$ and the claim is immediate. We may therefore assume $R \geq 3$. Since $0 \leq \Delta_x \leq g(x)$, it is enough to bound

$$a_x := g(x)\frac{\theta^x}{x!}, \qquad x \geq 3.$$

For $x \geq 3$,

$$\frac{a_{x+1}}{a_x} = \frac{2^x\theta}{x+1}, \qquad \frac{a_{x+2}/a_{x+1}}{a_{x+1}/a_x} = 2\frac{x+1}{x+2} > 1.$$

Thus (a_x) is log-convex, and its maximum on $[3, R]$ occurs at an endpoint. Consequently

$$\lambda_{ab} \leq R(a_3 + a_R).$$

Now $a_3 = (2/3)\theta^3$. Since $R\theta = O(U^3 2^{-U/3})$, one has $Ra_3 \leq \theta^2$ for all sufficiently large U.

For the other endpoint, using $R = \lfloor U/2 \rfloor$ and $\theta \leq \mathrm{e}U^2 2^{-U/3}$ gives

$$\log_2 \theta \leq \log_2 \mathrm{e} + 2\log_2 U - \frac{U}{3}.$$

For sufficiently large U, one has $R \geq 3$ and $g(R) = 2^{\binom{R}{2}-1}$. Hence

$$\begin{aligned}
\log_2\left(\frac{a_R}{\theta^2}\right) &= \binom{R}{2} - 1 + (R-2)\log_2\theta - \log_2(R!) \\
&\leq \left(\frac{U^2}{8} + O(U)\right) - \left(\frac{U^2}{6} + O(U)\right) + O(U\log U) \\
&= -\frac{U^2}{24} + O(U\log U).
\end{aligned}$$

The growth of $g(R)$ contributes the first quadratic term, and the small intensity in θ^{R-2} contributes the second. Their difference is negative; the factorial further lowers the bound. Thus $Ra_R \leq \theta^2$ for all sufficiently large U. Combining the two endpoint estimates gives $\lambda_{ab} \leq 2\theta^2$, and hence $q_{ab} \leq (5/2)\theta^2$.

Choose an integer U_0 beyond which the preceding two endpoint bounds hold. For $1 \leq U < U_0$, define

$$q_U(\theta) := \frac{\theta^2}{2} + \sum_{x=3}^{\lfloor U/2 \rfloor} \Delta_x \frac{\theta^x}{x!}, \qquad \Theta_U := \mathrm{e}U^2 2^{-U/3},$$

where an empty sum is zero. The quotient $q_U(\theta)/\theta^2$ extends continuously to $\theta = 0$ with value $1/2$. Therefore

$$C_{\mathrm{small}} := \max_{1 \leq U < U_0} \max_{0 \leq \theta \leq \Theta_U} \frac{q_U(\theta)}{\theta^2} < \infty.$$

Taking $C_0 = \max\{5/2, C_{\mathrm{small}}\}$ proves the lemma for every U. □

Appendix H. The Lean theorem and reproducibility records

A companion Lean 4 development kernel-checks the explicit full-sequence uniform consequence of the main theorem (Petkov, 2026). Its formal coefficient is $((\log 2)^2/4)\log(200/153)$. The top-level declaration is `Erdos625.erdos625`; its type is `Erdos625.Erdos625Statement`. The phase-resolved refinement involving $A_4(\delta_n)$ is proved in this manuscript and is not claimed as part of the Lean theorem. The released proof uses tail amplification. The mean-positive-excess bound in Lemma 10.2 is proved here but is not asserted as a separately released Lean theorem. The cited source revision pins Lean and Mathlib at version 4.31.0 and contains the modular sources and generated single-file closure. The cited replay revision contains the trust-boundary audit and a clean-environment replay. Reproduction commands, checksums, and transcripts are recorded in the sharp-constant replay record.

References

Tudor Achim, Alex Best, Alberto Bietti, Kevin Der, Mathïs Fédérico, Sergei Gukov, Daniel Halpern-Leistner, Kirsten Henningsgard, Yury Kudryashov, Alexander Meiburg, Martin Michelsen, Riley Patterson, Eric Rodriguez, Laura Scharff, Vikram Shanker, Vladmir Sicca, Hari Sowrirajan, Aidan Swope, Matyas Tamas, Vlad Tenev, Jonathan Thomm, Harold Williams, and Lawrence Wu. Aristotle: IMO-level automated theorem proving, 2025. URL https://doi.org/10.48550/arXiv.2510.01346.

Thomas F. Bloom. Erdős problem #625. https://www.erdosproblems.com/625, 2026. Accessed 16 July 2026.

Béla Bollobás. The chromatic number of random graphs. *Combinatorica*, 8(1):49–55, 1988. doi: 10.1007/BF02122551. URL https://doi.org/10.1007/BF02122551.

Béla Bollobás and Paul Erdős. Cliques in random graphs. *Mathematical Proceedings of the Cambridge Philosophical Society*, 80 (3):419–427, 1976. doi: 10.1017/S0305004100053056. URL https://doi.org/10.1017/S0305004100053056.

Béla Bollobás and Andrew Thomason. Generalized chromatic numbers of random graphs. *Random Structures & Algorithms*, 6 (2–3):353–356, 1995. doi: 10.1002/rsa.3240060222. URL https://doi.org/10.1002/rsa.3240060222.

Paul Erdős and John Gimbel. Some problems and results in cochromatic theory. In John Gimbel, John W. Kennedy, and Louis V. Quintas, editors, *Quo Vadis, Graph Theory?*, volume 55 of *Annals of Discrete Mathematics*, pages 261–264. North-Holland/Elsevier, 1993. doi: 10.1016/S0167-5060(08)70393-5. URL https://doi.org/10.1016/S0167-5060(08)70393-5.

John Gimbel. Some of my favorite coloring problems for graphs and digraphs. In Ralucca Gera, Stephen T. Hedetniemi, and Craig Larson, editors, *Graph Theory: Favorite Conjectures and Open Problems – 1*, Problem Books in Mathematics, pages 95–108. Springer, Cham, 2016. ISBN 978-3-319-31940-7. doi: 10.1007/978-3-319-31940-7_7. URL https://doi.org/10.1007/978-3-319-31940-7_7.

Geoffrey R. Grimmett and Colin J. H. McDiarmid. On colouring random graphs. *Mathematical Proceedings of the Cambridge Philosophical Society*, 77(2):313–324, 1975. doi: 10.1017/S0305004100051124. URL https://doi.org/10.1017/S0305004100051124.

Annika Heckel. The chromatic number of dense random graphs. *Random Structures & Algorithms*, 53(1):140–182, 2018. doi: 10.1002/rsa.20757. URL https://doi.org/10.1002/rsa.20757.

Annika Heckel. On a question of Erdős and Gimbel on the cochromatic number. *The Electronic Journal of Combinatorics*, 31(4): P4.72, 2024. doi: 10.37236/13346. URL https://doi.org/10.37236/13346.

Annika Heckel. The difference between the chromatic and the cochromatic number of a random graph, 2025. URL https://arxiv.org/abs/2409.17614. Version 2, revised 19 February 2025; first submitted in 2024.

Annika Heckel and Konstantinos Panagiotou. Colouring random graphs: Tame colourings, 2023. URL https://arxiv.org/abs/2306.07253. arXiv:2306.07253.

Colin McDiarmid. On the method of bounded differences. In Johannes Siemons, editor, *Surveys in Combinatorics, 1989*, volume 141 of *London Mathematical Society Lecture Note Series*, pages 148–188. Cambridge University Press, 1989. doi: 10.1017/CBO9781107359949.008. URL https://doi.org/10.1017/CBO9781107359949.008.

Colin McDiarmid. On the chromatic number of random graphs. *Random Structures & Algorithms*, 1(4):435–442, 1990. doi: 10.1002/rsa.3240010404. URL https://doi.org/10.1002/rsa.3240010404.

Colin McDiarmid. Concentration. In Michel Habib, Colin McDiarmid, Jorge Ramirez-Alfonsin, and Bruce Reed, editors, *Probabilistic Methods for Algorithmic Discrete Mathematics*, volume 16 of *Algorithms and Combinatorics*, pages 195–248. Springer, Berlin, Heidelberg, 1998. doi: 10.1007/978-3-662-12788-9_6. URL https://cgm.cs.mcgill.ca/~breed/conc/colin.pdf. Theorem 3.1; author preprint, pp. 10–11.

Konstantinos Panagiotou and Angelika Steger. A note on the chromatic number of a dense random graph. *Discrete Mathematics*, 309(10):3420–3423, 2009. doi: 10.1016/j.disc.2008.09.019. URL https://doi.org/10.1016/j.disc.2008.09.019.

Samuil Petkov. Erdős Problem 625: Lean 4 formalization and reproducibility record. Software companion at exact Git revisions, 2026. URL https://github.com/SamPetkov/Erdos/commit/31bbe00c529a996bdb61b880120d71240172d18f. Formal source revision 824e4b609466d2e26b216a76ecf103184dac2663; replay and archive revision 31bbe00c529a996bdb61b880120d71240172d18f; Lean and Mathlib v4.31.0.

Edward R. Scheinerman. Generalized chromatic numbers of random graphs. *SIAM Journal on Discrete Mathematics*, 5(1):74–80, 1992. doi: 10.1137/0405006. URL https://doi.org/10.1137/0405006.

Alex Scott. On the concentration of the chromatic number of random graphs, 2017. URL https://arxiv.org/abs/0806.0178. Version 2; first submitted in 2008.

Raphael Steiner. On the difference between the chromatic and cochromatic number. *SIAM Journal on Discrete Mathematics*, 39 (4):2268–2274, 2025. doi: 10.1137/24M1715180. URL https://doi.org/10.1137/24M1715180. Published online 17 November 2025; arXiv:2408.02400.

Yaming Yu. Relative log-concavity and a pair of triangle inequalities. *Bernoulli*, 16(2):459–470, 2010. doi: 10.3150/09-BEJ216. URL https://doi.org/10.3150/09-BEJ216.

École normale supérieure, Université PSL, Paris, France; samuil.petkov@ens.psl.eu